

Virtual I/O Device (VIRTIO) Version 1.0

Working Draft

03 March 2016

Specification URIs

This version:

<http://docs.oasis-open.org/virtio/virtio/v1.0/wd10/tex/> (Authoritative)
<http://docs.oasis-open.org/virtio/virtio/v1.0/wd10/virtio-v1.0-wd10.pdf>
<http://docs.oasis-open.org/virtio/virtio/v1.0/wd10/virtio-v1.0-wd10.html>

Previous version:

<http://docs.oasis-open.org/virtio/virtio/v1.0/cs03/tex/> (Authoritative)
<http://docs.oasis-open.org/virtio/virtio/v1.0/cs03/virtio-v1.0-cs03.pdf>
<http://docs.oasis-open.org/virtio/virtio/v1.0/cs03/virtio-v1.0-cs03.html>

Latest version:

<http://docs.oasis-open.org/virtio/virtio/v1.0/virtio-v1.0.pdf>
<http://docs.oasis-open.org/virtio/virtio/v1.0/virtio-v1.0.html>

Technical Committee:

OASIS Virtual I/O Device (VIRTIO) TC

Chairs:

Michael S. Tsirkin (mst@redhat.com), Red Hat

Editors:

Michael S. Tsirkin (mst@redhat.com), Red Hat
Cornelia Huck (cohuck@redhat.com), Red Hat
Pawel Moll (pawel.moll@arm.com), ARM

Additional artifacts:

This prose specification is one component of a Work Product that also includes:

- Example Driver Listing:
<http://docs.oasis-open.org/virtio/virtio/v1.0/wd10/listings/>

Related work:

This specification replaces or supersedes:

- Virtio PCI Card Specification Version 0.9.5:
<http://ozlabs.org/~rusty/virtio-spec/virtio-0.9.5.pdf>

Abstract:

This document describes the specifications of the “virtio” family of devices. These devices are found in virtual environments, yet by design they look like physical devices to the guest within the virtual machine - and this document treats them as such. This similarity allows the guest to use standard drivers and discovery mechanisms.

The purpose of virtio and this specification is that virtual environments and guests should have a straightforward, efficient, standard and extensible mechanism for virtual devices, rather than boutique per-environment or per-OS mechanisms.

Status:

This document was last revised or approved by the Virtual I/O Device (VIRTIO) TC on the above date. The level of approval is also listed above. Check the “Latest version” location noted above for possible later revisions of this document. Any other numbered Versions and other technical work produced by the Technical Committee (TC) are listed at https://www.oasis-open.org/committees/tc_home.php?wg_abbrev=virtio#technical.

Technical Committee members should send comments on this specification to the Technical Committee’s email list. Others should send comments to the Technical Committee by using the “Send A Comment” button on the Technical Committee’s web page at <https://www.oasis-open.org/committees/virtio/>.

For information on whether any patents have been disclosed that may be essential to implementing this specification, and any offers of patent licensing terms, please refer to the Intellectual Property Rights section of the Technical Committee web page (<https://www.oasis-open.org/committees/virtio/ipr.php>).

Citation format:

When referencing this specification the following citation format should be used:

[VIRTIO-v1.0]

Virtual I/O Device (VIRTIO) Version 1.0. Edited by Rusty Russell, Michael S. Tsirkin, Cornelia Huck, and Pawel Moll. 03 March 2016. OASIS Working Draft. <http://docs.oasis-open.org/virtio/virtio/v1.0/wd10/virtio-v1.0-wd10.html>. Latest version: <http://docs.oasis-open.org/virtio/virtio/v1.0/virtio-v1.0.html>.

Notices

Copyright © OASIS Open 2015. All Rights Reserved.

All capitalized terms in the following text have the meanings assigned to them in the OASIS Intellectual Property Rights Policy (the "OASIS IPR Policy"). The full [Policy](#) may be found at the OASIS website.

This document and translations of it may be copied and furnished to others, and derivative works that comment on or otherwise explain it or assist in its implementation may be prepared, copied, published, and distributed, in whole or in part, without restriction of any kind, provided that the above copyright notice and this section are included on all such copies and derivative works. However, this document itself may not be modified in any way, including by removing the copyright notice or references to OASIS, except as needed for the purpose of developing any document or deliverable produced by an OASIS Technical Committee (in which case the rules applicable to copyrights, as set forth in the OASIS IPR Policy, must be followed) or as required to translate it into languages other than English.

The limited permissions granted above are perpetual and will not be revoked by OASIS or its successors or assigns.

This document and the information contained herein is provided on an "AS IS" basis and OASIS DISCLAIMS ALL WARRANTIES, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY WARRANTY THAT THE USE OF THE INFORMATION HEREIN WILL NOT INFRINGE ANY OWNERSHIP RIGHTS OR ANY IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE.

OASIS requests that any OASIS Party or any other party that believes it has patent claims that would necessarily be infringed by implementations of this OASIS Committee Specification or OASIS Standard, to notify OASIS TC Administrator and provide an indication of its willingness to grant patent licenses to such patent claims in a manner consistent with the IPR Mode of the OASIS Technical Committee that produced this specification.

OASIS invites any party to contact the OASIS TC Administrator if it is aware of a claim of ownership of any patent claims that would necessarily be infringed by implementations of this specification by a patent holder that is not willing to provide a license to such patent claims in a manner consistent with the IPR Mode of the OASIS Technical Committee that produced this specification. OASIS may include such claims on its website, but disclaims any obligation to do so.

OASIS takes no position regarding the validity or scope of any intellectual property or other rights that might be claimed to pertain to the implementation or use of the technology described in this document or the extent to which any license under such rights might or might not be available; neither does it represent that it has made any effort to identify any such rights. Information on OASIS' procedures with respect to rights in any document or deliverable produced by an OASIS Technical Committee can be found on the OASIS website. Copies of claims of rights made available for publication and any assurances of licenses to be made available, or the result of an attempt made to obtain a general license or permission for the use of such proprietary rights by implementers or users of this OASIS Committee Specification or OASIS Standard, can be obtained from the OASIS TC Administrator. OASIS makes no representation that any information or list of intellectual property rights will at any time be complete, or that any claims in such list are, in fact, Essential Claims.

The name "OASIS" is a trademark of [OASIS](#), the owner and developer of this specification, and should be used only to refer to the organization and its official outputs. OASIS welcomes reference to, and implementation and use of, specifications, while reserving the right to enforce its marks against misleading uses. Please see <https://www.oasis-open.org/policies-guidelines/trademark> for above guidance.

Table of Contents

1	Introduction	11
1.1	Normative References	11
1.2	Non-Normative References	12
1.3	Terminology	12
1.3.1	Legacy Interface: Terminology	12
1.3.2	Transition from earlier specification drafts	12
1.4	Structure Specifications	13
2	Basic Facilities of a Virtio Device	14
2.1	Device Status Field	14
2.1.1	Driver Requirements: Device Status Field	14
2.1.2	Device Requirements: Device Status Field	15
2.2	Feature Bits	15
2.2.1	Driver Requirements: Feature Bits	15
2.2.2	Device Requirements: Feature Bits	15
2.2.3	Legacy Interface: A Note on Feature Bits	15
2.3	Device Configuration Space	16
2.3.1	Driver Requirements: Device Configuration Space	16
2.3.2	Device Requirements: Device Configuration Space	16
2.3.3	Legacy Interface: A Note on Device Configuration Space endian-ness	16
2.3.4	Legacy Interface: Device Configuration Space	17
2.4	Virtqueues	17
2.4.1	Driver-notifications	17
2.5	Split Virtqueues	18
2.5.1	Driver Requirements: Virtqueues	18
2.5.2	Legacy Interfaces: A Note on Virtqueue Layout	19
2.5.3	Legacy Interfaces: A Note on Virtqueue Endianness	19
2.5.4	Message Framing	19
2.5.4.1	Device Requirements: Message Framing	19
2.5.4.2	Driver Requirements: Message Framing	20
2.5.4.3	Legacy Interface: Message Framing	20
2.5.5	The Virtqueue Descriptor Table	20
2.5.5.1	Device Requirements: The Virtqueue Descriptor Table	20
2.5.5.2	Driver Requirements: The Virtqueue Descriptor Table	21
2.5.5.3	Indirect Descriptors	21
2.5.5.3.1	Driver Requirements: Indirect Descriptors	21
2.5.5.3.2	Device Requirements: Indirect Descriptors	21
2.5.6	The Virtqueue Available Ring	22
2.5.6.1	Driver Requirements: The Virtqueue Available Ring	22
2.5.7	Virtqueue Interrupt Suppression	22
2.5.7.1	Driver Requirements: Virtqueue Interrupt Suppression	22
2.5.7.2	Device Requirements: Virtqueue Interrupt Suppression	22
2.5.8	The Virtqueue Used Ring	23
2.5.8.1	Legacy Interface: The Virtqueue Used Ring	23
2.5.8.2	Device Requirements: The Virtqueue Used Ring	23
2.5.8.3	Driver Requirements: The Virtqueue Used Ring	24
2.5.9	Virtqueue Notification Suppression	24
2.5.9.1	Driver Requirements: Virtqueue Notification Suppression	24

2.5.9.2	Device Requirements: Virtqueue Notification Suppression	24
2.5.10	Helpers for Operating Virtqueues	25
2.5.11	Virtqueue Operation	25
2.5.12	Supplying Buffers to The Device	25
2.5.12.1	Placing Buffers Into The Descriptor Table	25
2.5.12.2	Updating The Available Ring	26
2.5.12.3	Updating <i>idx</i>	26
2.5.12.3.1	Driver Requirements: Updating <i>idx</i>	26
2.5.12.4	Notifying The Device	26
2.5.12.4.1	Driver Requirements: Notifying The Device	26
2.5.13	Receiving Used Buffers From The Device	27
2.6	Packed Virtqueues	27
2.6.1	Driver and Device Ring Wrap Counters	28
2.6.2	Polling of available and used descriptors	28
2.6.3	Write Flag	29
2.6.4	Element Address and Length	29
2.6.5	Scatter-Gather Support	29
2.6.6	Next Flag: Descriptor Chaining	29
2.6.7	Indirect Flag: Scatter-Gather Support	30
2.6.8	In-order use of descriptors	30
2.6.9	Multi-buffer requests	31
2.6.10	Driver and Device Event Suppression	31
2.6.10.1	Structure Size and Alignment	31
2.6.11	Driver Requirements: Virtqueues	32
2.6.12	Device Requirements: Virtqueues	32
2.6.13	The Virtqueue Descriptor Format	32
2.6.14	Event Suppression Structure Format	32
2.6.15	Device Requirements: The Virtqueue Descriptor Table	32
2.6.16	Driver Requirements: The Virtqueue Descriptor Table	33
2.6.17	Driver Requirements: Scatter-Gather Support	33
2.6.18	Device Requirements: Scatter-Gather Support	33
2.6.19	Driver Requirements: Indirect Descriptors	33
2.6.20	Virtqueue Operation	33
2.6.21	Supplying Buffers to The Device	33
2.6.21.1	Placing Available Buffers Into The Descriptor Ring	34
2.6.21.1.1	Driver Requirements: Updating flags	34
2.6.21.2	Notifying The Device	34
2.6.21.3	Implementation Example	35
2.6.21.3.1	Driver Requirements: Notifying The Device	35
2.6.22	Receiving Used Buffers From The Device	35
2.6.22.1	Driver-notifications	36
3	General Initialization And Device Operation	37
3.1	Device Initialization	37
3.1.1	Driver Requirements: Device Initialization	37
3.1.2	Legacy Interface: Device Initialization	37
3.2	Device Operation	38
3.2.1	Notification of Device Configuration Changes	38
3.3	Device Cleanup	38
3.3.1	Driver Requirements: Device Cleanup	38
4	Virtio Transport Options	39
4.1	Virtio Over PCI Bus	39
4.1.1	Device Requirements: Virtio Over PCI Bus	39
4.1.2	PCI Device Discovery	39
4.1.2.1	Device Requirements: PCI Device Discovery	39
4.1.2.2	Driver Requirements: PCI Device Discovery	40

4.1.2.3	Legacy Interfaces: A Note on PCI Device Discovery	40
4.1.3	PCI Device Layout	40
4.1.3.1	Driver Requirements: PCI Device Layout	40
4.1.3.2	Device Requirements: PCI Device Layout	40
4.1.4	Virtio Structure PCI Capabilities	40
4.1.4.1	Driver Requirements: Virtio Structure PCI Capabilities	42
4.1.4.2	Device Requirements: Virtio Structure PCI Capabilities	42
4.1.4.3	Common configuration structure layout	42
4.1.4.3.1	Device Requirements: Common configuration structure layout	43
4.1.4.3.2	Driver Requirements: Common configuration structure layout	44
4.1.4.4	Notification structure layout	44
4.1.4.4.1	Device Requirements: Notification capability	44
4.1.4.5	ISR status capability	45
4.1.4.5.1	Device Requirements: ISR status capability	45
4.1.4.5.2	Driver Requirements: ISR status capability	45
4.1.4.6	Device-specific configuration	45
4.1.4.6.1	Device Requirements: Device-specific configuration	45
4.1.4.7	PCI configuration access capability	46
4.1.4.7.1	Device Requirements: PCI configuration access capability	46
4.1.4.7.2	Driver Requirements: PCI configuration access capability	46
4.1.4.8	Legacy Interfaces: A Note on PCI Device Layout	46
4.1.4.9	Non-transitional Device With Legacy Driver: A Note on PCI Device Layout	47
4.1.5	PCI-specific Initialization And Device Operation	48
4.1.5.1	Device Initialization	48
4.1.5.1.1	Virtio Device Configuration Layout Detection	48
4.1.5.1.2	MSI-X Vector Configuration	48
4.1.5.1.3	Virtqueue Configuration	49
4.1.5.2	Notifying The Device	50
4.1.5.3	Virtqueue Interrupts From The Device	50
4.1.5.3.1	Device Requirements: Virtqueue Interrupts From The Device	50
4.1.5.4	Notification of Device Configuration Changes	50
4.1.5.4.1	Device Requirements: Notification of Device Configuration Changes	50
4.1.5.4.2	Driver Requirements: Notification of Device Configuration Changes	50
4.1.5.5	Driver Handling Interrupts	51
4.2	Virtio Over MMIO	51
4.2.1	MMIO Device Discovery	51
4.2.2	MMIO Device Register Layout	51
4.2.2.1	Device Requirements: MMIO Device Register Layout	54
4.2.2.2	Driver Requirements: MMIO Device Register Layout	54
4.2.3	MMIO-specific Initialization And Device Operation	55
4.2.3.1	Device Initialization	55
4.2.3.1.1	Driver Requirements: Device Initialization	55
4.2.3.2	Virtqueue Configuration	55
4.2.3.3	Notifying The Device	55
4.2.3.4	Notifications From The Device	55
4.2.3.4.1	Driver Requirements: Notifications From The Device	55
4.2.4	Legacy interface	55
4.3	Virtio Over Channel I/O	58
4.3.1	Basic Concepts	58
4.3.1.1	Device Requirements: Basic Concepts	59
4.3.1.2	Driver Requirements: Basic Concepts	59
4.3.2	Device Initialization	59
4.3.2.1	Setting the Virtio Revision	59
4.3.2.1.1	Device Requirements: Setting the Virtio Revision	59
4.3.2.1.2	Driver Requirements: Setting the Virtio Revision	60
4.3.2.1.3	Legacy Interfaces: A Note on Setting the Virtio Revision	60
4.3.2.2	Configuring a Virtqueue	60

4.3.2.2.1	Device Requirements: Configuring a Virtqueue	61
4.3.2.2.2	Legacy Interface: A Note on Configuring a Virtqueue	61
4.3.2.3	Communicating Status Information	61
4.3.2.3.1	Driver Requirements: Communicating Status Information	61
4.3.2.3.2	Device Requirements: Communicating Status Information	61
4.3.2.4	Handling Device Features	61
4.3.2.5	Device Configuration	62
4.3.2.6	Setting Up Indicators	62
4.3.2.6.1	Setting Up Classic Queue Indicators	62
4.3.2.6.2	Setting Up Configuration Change Indicators	62
4.3.2.6.3	Setting Up Two-Stage Queue Indicators	63
4.3.2.6.4	Legacy Interfaces: A Note on Setting Up Indicators	63
4.3.3	Device Operation	63
4.3.3.1	Host->Guest Notification	63
4.3.3.1.1	Notification via Classic I/O Interrupts	63
4.3.3.1.2	Notification via Adapter I/O Interrupts	64
4.3.3.1.3	Legacy Interfaces: A Note on Host->Guest Notification	64
4.3.3.2	Guest->Host Notification	64
4.3.3.2.1	Device Requirements: Guest->Host Notification	65
4.3.3.2.2	Driver Requirements: Guest->Host Notification	65
4.3.3.3	Resetting Devices	65
5	Device Types	66
5.1	Network Device	66
5.1.1	Device ID	67
5.1.2	Virtqueues	67
5.1.3	Feature bits	67
5.1.3.1	Feature bit requirements	68
5.1.3.2	Legacy Interface: Feature bits	68
5.1.4	Device configuration layout	68
5.1.4.1	Device Requirements: Device configuration layout	69
5.1.4.2	Driver Requirements: Device configuration layout	69
5.1.4.3	Legacy Interface: Device configuration layout	69
5.1.5	Device Initialization	69
5.1.6	Device Operation	70
5.1.6.1	Legacy Interface: Device Operation	70
5.1.6.2	Packet Transmission	71
5.1.6.2.1	Driver Requirements: Packet Transmission	71
5.1.6.2.2	Device Requirements: Packet Transmission	72
5.1.6.2.3	Packet Transmission Interrupt	72
5.1.6.3	Setting Up Receive Buffers	72
5.1.6.3.1	Driver Requirements: Setting Up Receive Buffers	73
5.1.6.3.2	Device Requirements: Setting Up Receive Buffers	73
5.1.6.4	Processing of Incoming Packets	73
5.1.6.4.1	Device Requirements: Processing of Incoming Packets	74
5.1.6.4.2	Driver Requirements: Processing of Incoming Packets	74
5.1.6.5	Control Virtqueue	75
5.1.6.5.1	Packet Receive Filtering	75
5.1.6.5.2	Setting MAC Address Filtering	76
5.1.6.5.3	VLAN Filtering	77
5.1.6.5.4	Gratuitous Packet Sending	77
5.1.6.5.5	Automatic receive steering in multiqueue mode	78
5.1.6.5.6	Offloads State Configuration	79
5.1.6.6	Legacy Interface: Framing Requirements	80
5.2	Block Device	80
5.2.1	Device ID	80
5.2.2	Virtqueues	80

5.2.3	Feature bits	80
5.2.3.1	Legacy Interface: Feature bits	81
5.2.4	Device configuration layout	81
5.2.4.1	Legacy Interface: Device configuration layout	81
5.2.5	Device Initialization	81
5.2.5.1	Driver Requirements: Device Initialization	82
5.2.5.2	Device Requirements: Device Initialization	82
5.2.5.3	Legacy Interface: Device Initialization	82
5.2.6	Device Operation	82
5.2.6.1	Driver Requirements: Device Operation	83
5.2.6.2	Device Requirements: Device Operation	83
5.2.6.3	Legacy Interface: Device Operation	83
5.2.6.4	Legacy Interface: Framing Requirements	85
5.3	Console Device	85
5.3.1	Device ID	85
5.3.2	Virtqueues	85
5.3.3	Feature bits	85
5.3.4	Device configuration layout	86
5.3.4.1	Legacy Interface: Device configuration layout	86
5.3.5	Device Initialization	86
5.3.5.1	Device Requirements: Device Initialization	86
5.3.6	Device Operation	86
5.3.6.1	Driver Requirements: Device Operation	87
5.3.6.2	Multiport Device Operation	87
5.3.6.2.1	Device Requirements: Multiport Device Operation	88
5.3.6.2.2	Driver Requirements: Multiport Device Operation	88
5.3.6.3	Legacy Interface: Device Operation	88
5.3.6.4	Legacy Interface: Framing Requirements	88
5.4	Entropy Device	88
5.4.1	Device ID	88
5.4.2	Virtqueues	88
5.4.3	Feature bits	88
5.4.4	Device configuration layout	89
5.4.5	Device Initialization	89
5.4.6	Device Operation	89
5.4.6.1	Driver Requirements: Device Operation	89
5.4.6.2	Device Requirements: Device Operation	89
5.5	Traditional Memory Balloon Device	89
5.5.1	Device ID	89
5.5.2	Virtqueues	89
5.5.3	Feature bits	90
5.5.3.1	Driver Requirements: Feature bits	90
5.5.3.2	Device Requirements: Feature bits	90
5.5.4	Device configuration layout	90
5.5.5	Device Initialization	90
5.5.6	Device Operation	91
5.5.6.1	Driver Requirements: Device Operation	91
5.5.6.2	Device Requirements: Device Operation	92
5.5.6.2.1	Legacy Interface: Device Operation	92
5.5.6.3	Memory Statistics	92
5.5.6.3.1	Driver Requirements: Memory Statistics	93
5.5.6.3.2	Device Requirements: Memory Statistics	93
5.5.6.3.3	Legacy Interface: Memory Statistics	93
5.5.6.4	Memory Statistics Tags	93
5.6	SCSI Host Device	94
5.6.1	Device ID	94
5.6.2	Virtqueues	94

5.6.3	Feature bits	94
5.6.4	Device configuration layout	94
5.6.4.1	Driver Requirements: Device configuration layout	95
5.6.4.2	Device Requirements: Device configuration layout	95
5.6.4.3	Legacy Interface: Device configuration layout	95
5.6.5	Device Requirements: Device Initialization	95
5.6.6	Device Operation	95
5.6.6.0.1	Legacy Interface: Device Operation	96
5.6.6.1	Device Operation: Request Queues	96
5.6.6.1.1	Device Requirements: Device Operation: Request Queues	97
5.6.6.1.2	Driver Requirements: Device Operation: Request Queues	98
5.6.6.1.3	Legacy Interface: Device Operation: Request Queues	98
5.6.6.2	Device Operation: controlq	98
5.6.6.2.1	Legacy Interface: Device Operation: controlq	100
5.6.6.3	Device Operation: eventq	100
5.6.6.3.1	Driver Requirements: Device Operation: eventq	102
5.6.6.3.2	Device Requirements: Device Operation: eventq	102
5.6.6.3.3	Legacy Interface: Device Operation: eventq	102
5.6.6.4	Legacy Interface: Framing Requirements	102
6	Reserved Feature Bits	103
6.1	Driver Requirements: Reserved Feature Bits	103
6.2	Device Requirements: Reserved Feature Bits	103
6.3	Legacy Interface: Reserved Feature Bits	104
7	Conformance	105
7.1	Conformance Targets	105
7.2	Driver Conformance	105
7.2.1	PCI Driver Conformance	106
7.2.2	MMIO Driver Conformance	106
7.2.3	Channel I/O Driver Conformance	106
7.2.4	Network Driver Conformance	106
7.2.5	Block Driver Conformance	107
7.2.6	Console Driver Conformance	107
7.2.7	Entropy Driver Conformance	107
7.2.8	Traditional Memory Balloon Driver Conformance	107
7.2.9	SCSI Host Driver Conformance	107
7.3	Device Conformance	107
7.3.1	PCI Device Conformance	108
7.3.2	MMIO Device Conformance	108
7.3.3	Channel I/O Device Conformance	108
7.3.4	Network Device Conformance	109
7.3.5	Block Device Conformance	109
7.3.6	Console Device Conformance	109
7.3.7	Entropy Device Conformance	109
7.3.8	Traditional Memory Balloon Device Conformance	109
7.3.9	SCSI Host Device Conformance	109
7.4	Legacy Interface: Transitional Device and Transitional Driver Conformance	110
A	virtio_queue.h	112
B	Creating New Device Types	114
B.1	How Many Virtqueues?	114
B.2	What Device Configuration Space Layout?	114
B.3	What Device Number?	114
B.4	How many MSI-X vectors? (for PCI)	114
B.5	Device Improvements	115

C Acknowledgements	116
D Revision History	117

1 Introduction

This document describes the specifications of the “virtio” family of devices. These devices are found in virtual environments, yet by design they look like physical devices to the guest within the virtual machine - and this document treats them as such. This similarity allows the guest to use standard drivers and discovery mechanisms.

The purpose of virtio and this specification is that virtual environments and guests should have a straightforward, efficient, standard and extensible mechanism for virtual devices, rather than boutique per-environment or per-OS mechanisms.

Straightforward: Virtio devices use normal bus mechanisms of interrupts and DMA which should be familiar to any device driver author. There is no exotic page-flipping or COW mechanism: it’s just a normal device.¹

Efficient: Virtio devices consist of rings of descriptors for both input and output, which are neatly laid out to avoid cache effects from both driver and device writing to the same cache lines.

Standard: Virtio makes no assumptions about the environment in which it operates, beyond supporting the bus to which device is attached. In this specification, virtio devices are implemented over MMIO, Channel I/O and PCI bus transports², earlier drafts have been implemented on other buses not included here.

Extensible: Virtio devices contain feature bits which are acknowledged by the guest operating system during device setup. This allows forwards and backwards compatibility: the device offers all the features it knows about, and the driver acknowledges those it understands and wishes to use.

1.1 Normative References

- | | |
|--------------------------|---|
| [RFC2119] | Bradner S., “Key words for use in RFCs to Indicate Requirement Levels”, BCP 14, RFC 2119, March 1997.
http://www.ietf.org/rfc/rfc2119.txt |
| [S390 PoP] | z/Architecture Principles of Operation, IBM Publication SA22-7832,
http://publibfi.boulder.ibm.com/epubs/pdf/dz9zr009.pdf , and any future revisions |
| [S390 Common I/O] | ESA/390 Common I/O-Device and Self-Description, IBM Publication SA22-7204,
http://publibfp.dhe.ibm.com/cgi-bin/bookmgr/BOOKS/dz9ar501/CCONTENTS ,
and any future revisions |
| [PCI] | Conventional PCI Specifications,
http://www.pcisig.com/specifications/conventional/ , PCI-SIG |
| [PCIe] | PCI Express Specifications
http://www.pcisig.com/specifications/pciexpress/ , PCI-SIG |
| [IEEE 802] | IEEE Standard for Local and Metropolitan Area Networks: Overview and Architecture,
http://standards.ieee.org/about/get/802/802.html , IEEE |

¹This lack of page-sharing implies that the implementation of the device (e.g. the hypervisor or host) needs full access to the guest memory. Communication with untrusted parties (i.e. inter-guest communication) requires copying.

²The Linux implementation further separates the virtio transport code from the specific virtio drivers: these drivers are shared between different transports.

[SAM]	SCSI Architectural Model, http://www.t10.org/cgi-bin/ac.pl?t=f&f=sam4r05.pdf
[SCSI MMC]	SCSI Multimedia Commands, http://www.t10.org/cgi-bin/ac.pl?t=f&f=mmc6r00.pdf

1.2 Non-Normative References

[Virtio PCI Draft]	Virtio PCI Draft Specification http://ozlabs.org/~rusty/virtio-spec/virtio-0.9.5.pdf
---------------------------	---

1.3 Terminology

The key words “MUST”, “MUST NOT”, “REQUIRED”, “SHALL”, “SHALL NOT”, “SHOULD”, “SHOULD NOT”, “RECOMMENDED”, “MAY”, and “OPTIONAL” in this document are to be interpreted as described in [\[RFC2119\]](#).

1.3.1 Legacy Interface: Terminology

Earlier drafts of this specification (i.e. revisions before 1.0, see e.g. [\[Virtio PCI Draft\]](#)) defined a similar, but different interface between the driver and the device. Since these are widely deployed, this specification accommodates OPTIONAL features to simplify transition from these earlier draft interfaces.

Specifically devices and drivers MAY support:

Legacy Interface is an interface specified by an earlier draft of this specification (before 1.0)

Legacy Device is a device implemented before this specification was released, and implementing a legacy interface on the host side

Legacy Driver is a driver implemented before this specification was released, and implementing a legacy interface on the guest side

Legacy devices and legacy drivers are not compliant with this specification.

To simplify transition from these earlier draft interfaces, a device MAY implement:

Transitional Device a device supporting both drivers conforming to this specification, and allowing legacy drivers.

Similarly, a driver MAY implement:

Transitional Driver a driver supporting both devices conforming to this specification, and legacy devices.

Note: Legacy interfaces are not required; ie. don't implement them unless you have a need for backwards compatibility!

Devices or drivers with no legacy compatibility are referred to as non-transitional devices and drivers, respectively.

1.3.2 Transition from earlier specification drafts

For devices and drivers already implementing the legacy interface, some changes will have to be made to support this specification.

In this case, it might be beneficial for the reader to focus on sections tagged "Legacy Interface" in the section title. These highlight the changes made since the earlier drafts.

1.4 Structure Specifications

Many device and driver in-memory structure layouts are documented using the C struct syntax. All structures are assumed to be without additional padding. To stress this, cases where common C compilers are known to insert extra padding within structures are tagged using the GNU C `__attribute__((packed))` syntax.

For the integer data types used in the structure definitions, the following conventions are used:

u8, u16, u32, u64 An unsigned integer of the specified length in bits.

le16, le32, le64 An unsigned integer of the specified length in bits, in little-endian byte order.

be16, be32, be64 An unsigned integer of the specified length in bits, in big-endian byte order.

Some of the fields to be defined in this specification don't start or don't end on a byte boundary. Such fields are called bit-fields. A set of bit-fields is always a sub-division of an integer typed field.

Bit-fields within integer fields are always listed in order, from the least significant to the most significant bit. The bit-fields are considered unsigned integers of the specified width with the next in significance relationship of the bits preserved.

For example:

```
struct S {  
    be16 {  
        A : 15;  
        B : 1;  
    } x;  
    be16 y;  
};
```

documents the value A stored in the low 15 bit of x and the value B stored in the high bit of x, the 16-bit integer x in turn stored using the big-endian byte order at the beginning of the structure S, and being followed immediately by an unsigned integer y stored in big-endian byte order at an offset of 2 bytes (16 bits) from the beginning of the structure.

Note that this notation somewhat resembles the C bitfield syntax but should not be naively converted to a bitfield notation for portable code: it matches the way bitfields are packed by C compilers on little-endian architectures but not the way bitfields are packed by C compilers on big-endian architectures.

Assuming that CPU_TO_BE16 converts a 16-bit integer from a native CPU to the big-endian byte order, the following is the equivalent portable C code to generate a value to be stored into x:

```
CPU_TO_BE16(B << 15 | A)
```

2 Basic Facilities of a Virtio Device

A virtio device is discovered and identified by a bus-specific method (see the bus specific sections: [4.1 Virtio Over PCI Bus](#), [4.2 Virtio Over MMIO](#) and [4.3 Virtio Over Channel I/O](#)). Each device consists of the following parts:

- Device status field
- Feature bits
- Device Configuration space
- One or more virtqueues

2.1 Device Status Field

During device initialization by a driver, the driver follows the sequence of steps specified in [3.1](#).

The *device status* field provides a simple low-level indication of the completed steps of this sequence. It's most useful to imagine it hooked up to traffic lights on the console indicating the status of each device. The following bits are defined (listed below in the order in which they would be typically set):

ACKNOWLEDGE (1) Indicates that the guest OS has found the device and recognized it as a valid virtio device.

DRIVER (2) Indicates that the guest OS knows how to drive the device.

Note: There could be a significant (or infinite) delay before setting this bit. For example, under Linux, drivers can be loadable modules.

FAILED (128) Indicates that something went wrong in the guest, and it has given up on the device. This could be an internal error, or the driver didn't like the device for some reason, or even a fatal error during device operation.

FEATURES_OK (8) Indicates that the driver has acknowledged all the features it understands, and feature negotiation is complete.

DRIVER_OK (4) Indicates that the driver is set up and ready to drive the device.

DEVICE_NEEDS_RESET (64) Indicates that the device has experienced an error from which it can't recover.

2.1.1 Driver Requirements: Device Status Field

The driver **MUST** update *device status*, setting bits to indicate the completed steps of the driver initialization sequence specified in [3.1](#). The driver **MUST NOT** clear a *device status* bit. If the driver sets the FAILED bit, the driver **MUST** later reset the device before attempting to re-initialize.

The driver **SHOULD NOT** rely on completion of operations of a device if **DEVICE_NEEDS_RESET** is set.

Note: For example, the driver can't assume requests in flight will be completed if **DEVICE_NEEDS_RESET** is set, nor can it assume that they have not been completed. A good implementation will try to recover by issuing a reset.

2.1.2 Device Requirements: Device Status Field

The device **MUST** initialize *device status* to 0 upon reset.

The device **MUST NOT** consume buffers or notify the driver before `DRIVER_OK`.

The device **SHOULD** set `DEVICE_NEEDS_RESET` when it enters an error state that a reset is needed. If `DRIVER_OK` is set, after it sets `DEVICE_NEEDS_RESET`, the device **MUST** send a device configuration change notification to the driver.

2.2 Feature Bits

Each virtio device offers all the features it understands. During device initialization, the driver reads this and tells the device the subset that it accepts. The only way to renegotiate is to reset the device.

This allows for forwards and backwards compatibility: if the device is enhanced with a new feature bit, older drivers will not write that feature bit back to the device. Similarly, if a driver is enhanced with a feature that the device doesn't support, it sees the new feature is not offered.

Feature bits are allocated as follows:

0 to 23 Feature bits for the specific device type

24 to 33 Feature bits reserved for extensions to the queue and feature negotiation mechanisms

34 and above Feature bits reserved for future extensions.

Note: For example, feature bit 0 for a network device (i.e. Device ID 1) indicates that the device supports checksumming of packets.

In particular, new fields in the device configuration space are indicated by offering a new feature bit.

2.2.1 Driver Requirements: Feature Bits

The driver **MUST NOT** accept a feature which the device did not offer, and **MUST NOT** accept a feature which requires another feature which was not accepted.

The driver **SHOULD** go into backwards compatibility mode if the device does not offer a feature it understands, otherwise **MUST** set the `FAILED device status` bit and cease initialization.

2.2.2 Device Requirements: Feature Bits

The device **MUST NOT** offer a feature which requires another feature which was not offered. The device **SHOULD** accept any valid subset of features the driver accepts, otherwise it **MUST** fail to set the `FEATURES_OK device status` bit when the driver writes it.

2.2.3 Legacy Interface: A Note on Feature Bits

Transitional Drivers **MUST** detect Legacy Devices by detecting that the feature bit `VIRTIO_F_VERSION_1` is not offered. Transitional devices **MUST** detect Legacy drivers by detecting that `VIRTIO_F_VERSION_1` has not been acknowledged by the driver.

In this case device is used through the legacy interface.

Legacy interface support is **OPTIONAL**. Thus, both transitional and non-transitional devices and drivers are compliant with this specification.

Requirements pertaining to transitional devices and drivers is contained in sections named 'Legacy Interface' like this one.

When device is used through the legacy interface, transitional devices and transitional drivers **MUST** operate according to the requirements documented within these legacy interface sections. Specification text within these sections generally does not apply to non-transitional devices.

2.3 Device Configuration Space

Device configuration space is generally used for rarely-changing or initialization-time parameters. Where configuration fields are optional, their existence is indicated by feature bits: Future versions of this specification will likely extend the device configuration space by adding extra fields at the tail.

Note: The device configuration space uses the little-endian format for multi-byte fields.

Each transport also provides a generation count for the device configuration space, which will change whenever there is a possibility that two accesses to the device configuration space can see different versions of that space.

2.3.1 Driver Requirements: Device Configuration Space

Drivers **MUST NOT** assume reads from fields greater than 32 bits wide are atomic, nor are reads from multiple fields: drivers **SHOULD** read device configuration space fields like so:

```
u32 before, after;
do {
    before = get_config_generation(device);
    // read config entry/entries.
    after = get_config_generation(device);
} while (after != before);
```

For optional configuration space fields, the driver **MUST** check that the corresponding feature is offered before accessing that part of the configuration space.

Note: See section 3.1 for details on feature negotiation.

Drivers **MUST NOT** limit structure size and device configuration space size. Instead, drivers **SHOULD** only check that device configuration space is *large enough* to contain the fields necessary for device operation.

Note: For example, if the specification states that device configuration space 'includes a single 8-bit field' drivers should understand this to mean that the device configuration space might also include an arbitrary amount of tail padding, and accept any device configuration space size equal to or greater than the specified 8-bit size.

2.3.2 Device Requirements: Device Configuration Space

The device **MUST** allow reading of any device-specific configuration field before **FEATURES_OK** is set by the driver. This includes fields which are conditional on feature bits, as long as those feature bits are offered by the device.

2.3.3 Legacy Interface: A Note on Device Configuration Space endian-ness

Note that for legacy interfaces, device configuration space is generally the guest's native endian, rather than PCI's little-endian. The correct endian-ness is documented for each device.

2.3.4 Legacy Interface: Device Configuration Space

Legacy devices did not have a configuration generation field, thus are susceptible to race conditions if configuration is updated. This affects the block *capacity* (see 5.2.4) and network *mac* (see 5.1.4) fields; when using the legacy interface, drivers SHOULD read these fields multiple times until two reads generate a consistent result.

2.4 Virtqueues

The mechanism for bulk data transport on virtio devices is pretentiously called a virtqueue. Each device can have zero or more virtqueues¹.

Driver makes requests available to device by adding an available buffer to the queue - i.e. adding a buffer describing the request to a virtqueue, and optionally triggering a driver event - i.e. sending a notification to the device.

Device executes the requests and - when complete - adds a used buffer to the queue - i.e. lets the driver know by marking the buffer as used. Device can then trigger a device event - i.e. send an interrupt to the driver.

Device reports the number of bytes it has written to memory for each buffer it uses. This is referred to as "used length".

Device is not generally required to use buffers in the same order in which they have been made available by the driver.

Some devices always use descriptors in the same order in which they have been made available. These devices can offer the VIRTIO_F_IN_ORDER feature. If negotiated, this knowledge might allow optimizations or simplify driver and/or device code.

Each virtqueue can consist of up to 3 parts:

- Descriptor Area - used for describing buffers
- Driver Area - extra data supplied by driver to the device
- Device Area - extra data supplied by device to driver

Note: Note that previous versions of this spec used different names for these parts (following 2.5):

- Descriptor Table - for the Descriptor Area
- Available Ring - for the Driver Area
- Used Ring - for the Device Area

Two formats are supported: Split Virtqueues (see 2.5 Split Virtqueues) and Packed Virtqueues (see 2.6 Packed Virtqueues).

Every driver and device supports either the Packed or the Split Virtqueue format, or both.

2.4.1 Driver notifications

~~Driver is sometimes required to notify the device after making changes to the virtqueue.~~

~~When VIRTIO_F_NOTIFICATION_DATA has not been negotiated, this notification involves sending the virtqueue number to the device (depending on the transport).~~

~~However, some devices benefit from ability to find out the number of available descriptors in the ring, and whether to send interrupts to drivers without accessing virtqueue in memory: for efficiency or as a debugging aid.~~

¹For example, the simplest network device has one virtqueue for transmit and one for receive.

To help with these optimizations, when VIRTIO_F_NOTIFICATION_DATA has been negotiated, driver notifications to the device include the following information:–

VQ-numberOffsetWithin the ring where the next available ring entry will be written. Without VIRTIO_F_RING_PACKED this refers to the 15 least significant bits of the available index. With VIRTIO_F_RING_PACKED this refers to the offset (in units of descriptor entries) within the descriptor ring where the next available descriptor will be written. Wrap CounterWith VIRTIO_F_RING_PACKED this is the wrap counter referring to the next available descriptor. Without VIRTIO_F_RING_PACKED this is the most significant bit of the available index.–

Note that driver can trigger multiple notifications even without making any more changes to the ring. When VIRTIO_F_NOTIFICATION_DATA has been negotiated, these notifications would then have identical Offset and Wrap Counter values.–

2.5 Split Virtqueues

The split virtqueue format was the only format supported by the version 1.0 (and earlier) of this standard.

The split virtqueue format separates the virtqueue into several parts, where each part is write-able by either the driver or the device, but not both. Multiple parts and/or locations within a part need to be updated when making a buffer available and when marking it as used.

Each queue has a 16-bit queue size parameter, which sets the number of entries and implies the total size of the queue.

Each virtqueue consists of three parts:

- Descriptor Table - occupies the Descriptor Area
- Available Ring - occupies the Driver Area
- Used Ring - occupies the Device Area

where each part is physically-contiguous in guest memory, and has different alignment requirements.

The memory alignment and size requirements, in bytes, of each part of the virtqueue are summarized in the following table:

Virtqueue Part	Alignment	Size
Descriptor Table	16	16*(Queue Size)
Available Ring	2	6 + 2*(Queue Size)
Used Ring	4	6 + 8*(Queue Size)

The Alignment column gives the minimum alignment for each part of the virtqueue.

The Size column gives the total number of bytes for each part of the virtqueue.

Queue Size corresponds to the maximum number of buffers in the virtqueue². Queue Size value is always a power of 2. The maximum Queue Size value is 32768. This value is specified in a bus-specific way.

When the driver wants to send a buffer to the device, it fills in a slot in the descriptor table (or chains several together), and writes the descriptor index into the available ring. It then notifies the device. When the device has finished a buffer, it writes the descriptor index into the used ring, and sends an interrupt.

2.5.1 Driver Requirements: Virtqueues

The driver MUST ensure that the physical address of the first byte of each virtqueue part is a multiple of the specified alignment value in the above table.

²For example, if Queue Size is 4 then at most 4 buffers can be queued at any given time.

2.5.2 Legacy Interfaces: A Note on Virtqueue Layout

For Legacy Interfaces, several additional restrictions are placed on the virtqueue layout:

Each virtqueue occupies two or more physically-contiguous pages (usually defined as 4096 bytes, but depending on the transport; henceforth referred to as Queue Align) and consists of three parts:

Descriptor Table	Available Ring (...padding...)	Used Ring
------------------	--------------------------------	-----------

The bus-specific Queue Size field controls the total number of bytes for the virtqueue. When using the legacy interface, the transitional driver **MUST** retrieve the Queue Size field from the device and **MUST** allocate the total number of bytes for the virtqueue according to the following formula (Queue Align given in qalign and Queue Size given in qsz):

```
#define ALIGN(x) (((x) + qalign) & ~qalign)
static inline unsigned virtq_size(unsigned int qsz)
{
    return ALIGN(sizeof(struct virtq_desc)*qsz + sizeof(u16)*(3 + qsz))
        + ALIGN(sizeof(u16)*3 + sizeof(struct virtq_used_elem)*qsz);
}
```

This wastes some space with padding. When using the legacy interface, both transitional devices and drivers **MUST** use the following virtqueue layout structure to locate elements of the virtqueue:

```
struct virtq {
    // The actual descriptors (16 bytes each)
    struct virtq_desc desc[ Queue Size ];

    // A ring of available descriptor heads with free-running index.
    struct virtq_avail avail;

    // Padding to the next Queue Align boundary.
    u8 pad[ Padding ];

    // A ring of used descriptor heads with free-running index.
    struct virtq_used used;
};
```

2.5.3 Legacy Interfaces: A Note on Virtqueue Endianness

Note that when using the legacy interface, transitional devices and drivers **MUST** use the native endian of the guest as the endian of fields and in the virtqueue. This is opposed to little-endian for non-legacy interface as specified by this standard. It is assumed that the host is already aware of the guest endian.

2.5.4 Message Framing

The framing of messages with descriptors is independent of the contents of the buffers. For example, a network transmit buffer consists of a 12 byte header followed by the network packet. This could be most simply placed in the descriptor table as a 12 byte output descriptor followed by a 1514 byte output descriptor, but it could also consist of a single 1526 byte output descriptor in the case where the header and packet are adjacent, or even three or more descriptors (possibly with loss of efficiency in that case).

Note that, some device implementations have large-but-reasonable restrictions on total descriptor size (such as based on IOV_MAX in the host OS). This has not been a problem in practice: little sympathy will be given to drivers which create unreasonably-sized descriptors such as by dividing a network packet into 1500 single-byte descriptors!

2.5.4.1 Device Requirements: Message Framing

The device **MUST NOT** make assumptions about the particular arrangement of descriptors. The device **MAY** have a reasonable limit of descriptors it will allow in a chain.

2.5.4.2 Driver Requirements: Message Framing

The driver **MUST** place any device-writable descriptor elements after any device-readable descriptor elements.

The driver **SHOULD NOT** use an excessive number of descriptors to describe a buffer.

2.5.4.3 Legacy Interface: Message Framing

Regrettably, initial driver implementations used simple layouts, and devices came to rely on it, despite this specification wording. In addition, the specification for virtio_blk SCSI commands required intuiting field lengths from frame boundaries (see [5.2.6.3 Legacy Interface: Device Operation](#))

Thus when using the legacy interface, the VIRTIO_F_ANY_LAYOUT feature indicates to both the device and the driver that no assumptions were made about framing. Requirements for transitional drivers when this is not negotiated are included in each device section.

2.5.5 The Virtqueue Descriptor Table

The descriptor table refers to the buffers the driver is using for the device. *addr* is a physical address, and the buffers can be chained via *next*. Each descriptor describes a buffer which is read-only for the device (“device-readable”) or write-only for the device (“device-writable”), but a chain of descriptors can contain both device-readable and device-writable buffers.

The actual contents of the memory offered to the device depends on the device type. Most common is to begin the data with a header (containing little-endian fields) for the device to read, and postfix it with a status tailer for the device to write.

```
struct virtq_desc {
    /* Address (guest-physical). */
    le64 addr;
    /* Length. */
    le32 len;

    /* This marks a buffer as continuing via the next field. */
#define VIRTQ_DESC_F_NEXT 1
    /* This marks a buffer as device write-only (otherwise device read-only). */
#define VIRTQ_DESC_F_WRITE 2
    /* This means the buffer contains a list of buffer descriptors. */
#define VIRTQ_DESC_F_INDIRECT 4
    /* The flags as indicated above. */
    le16 flags;
    /* Next field if flags & NEXT */
    le16 next;
};
```

The number of descriptors in the table is defined by the queue size for this virtqueue: this is the maximum possible descriptor chain length.

If VIRTIO_F_IN_ORDER has been negotiated, driver uses descriptors in ring order: starting from offset 0 in the table, and wrapping around at the end of the table.

Note: The legacy [\[Virtio PCI Draft\]](#) referred to this structure as `vring_desc`, and the constants as `VRING_DESC_F_NEXT`, etc, but the layout and values were identical.

2.5.5.1 Device Requirements: The Virtqueue Descriptor Table

A device **MUST NOT** write to a device-readable buffer, and a device **SHOULD NOT** read a device-writable buffer (it **MAY** do so for debugging or diagnostic purposes).

2.5.5.2 Driver Requirements: The Virtqueue Descriptor Table

Drivers MUST NOT add a descriptor chain over than 2^{32} bytes long in total; this implies that loops in the descriptor chain are forbidden!

If `VIRTIO_F_IN_ORDER` has been negotiated, and when making a descriptor with `VRING_DESC_F_NEXT` set in *flags* at offset *x* in the table available to the device, driver MUST set *next* to 0 for the last descriptor in the table (where $x = \text{queue_size} - 1$) and to $x + 1$ for the rest of the descriptors.

2.5.5.3 Indirect Descriptors

Some devices benefit by concurrently dispatching a large number of large requests. The `VIRTIO_F_INDIRECT_DESC` feature allows this (see [A virtio_queue.h](#)). To increase ring capacity the driver can store a table of indirect descriptors anywhere in memory, and insert a descriptor in main virtqueue (with *flags*&`VIRTQ_DESC_F_INDIRECT` on) that refers to memory buffer containing this indirect descriptor table; *addr* and *len* refer to the indirect table address and length in bytes, respectively.

The indirect table layout structure looks like this (*len* is the length of the descriptor that refers to this table, which is a variable, so this code won't compile):

```
struct indirect_descriptor_table {
    /* The actual descriptors (16 bytes each) */
    struct virtq_desc desc[len / 16];
};
```

The first indirect descriptor is located at start of the indirect descriptor table (index 0), additional indirect descriptors are chained by *next*. An indirect descriptor without a valid *next* (with *flags*&`VIRTQ_DESC_F_NEXT` off) signals the end of the descriptor. A single indirect descriptor table can include both device-readable and device-writable descriptors.

If `VIRTIO_F_IN_ORDER` has been negotiated, indirect descriptors use sequential indices, in-order: index 0 followed by index 1 followed by index 2, etc.

2.5.5.3.1 Driver Requirements: Indirect Descriptors

The driver MUST NOT set the `VIRTQ_DESC_F_INDIRECT` flag unless the `VIRTIO_F_INDIRECT_DESC` feature was negotiated. The driver MUST NOT set the `VIRTQ_DESC_F_INDIRECT` flag within an indirect descriptor (ie. only one table per descriptor).

A driver MUST NOT create a descriptor chain longer than the Queue Size of the device.

A driver MUST NOT set both `VIRTQ_DESC_F_INDIRECT` and `VIRTQ_DESC_F_NEXT` in *flags*.

If `VIRTIO_F_IN_ORDER` has been negotiated, indirect descriptors MUST appear sequentially, with *next* taking the value of 1 for the 1st descriptor, 2 for the 2nd one, etc.

2.5.5.3.2 Device Requirements: Indirect Descriptors

The device MUST ignore the write-only flag (*flags*&`VIRTQ_DESC_F_WRITE`) in the descriptor that refers to an indirect table.

The device MUST handle the case of zero or more normal chained descriptors followed by a single descriptor with *flags*&`VIRTQ_DESC_F_INDIRECT`.

Note: While unusual (most implementations either create a chain solely using non-indirect descriptors, or use a single indirect element), such a layout is valid.

2.5.6 The Virtqueue Available Ring

```
struct virtq_avail {
#define VIRTQ_AVAIL_F_NO_INTERRUPT    1
    le16 flags;
    le16 idx;
    le16 ring[ /* Queue Size */ ];
    le16 used_event; /* Only if VIRTIO_F_EVENT_IDX */
};
```

The driver uses the available ring to offer buffers to the device: each ring entry refers to the head of a descriptor chain. It is only written by the driver and read by the device.

idx field indicates where the driver would put the next descriptor entry in the ring (modulo the queue size). This starts at 0, and increases.

Note: The legacy [\[Virtio PCI Draft\]](#) referred to this structure as `vring_avail`, and the constant as `VRING_AVAIL_F_NO_INTERRUPT`, but the layout and value were identical.

2.5.6.1 Driver Requirements: The Virtqueue Available Ring

A driver MUST NOT decrement the available *idx* on a virtqueue (ie. there is no way to “unexpose” buffers).

2.5.7 Virtqueue Interrupt Suppression

If the `VIRTIO_F_EVENT_IDX` feature bit is not negotiated, the *flags* field in the available ring offers a crude mechanism for the driver to inform the device that it doesn't want interrupts when buffers are used. Otherwise *used_event* is a more performant alternative where the driver specifies how far the device can progress before interrupting.

Neither of these interrupt suppression methods are reliable, as they are not synchronized with the device, but they serve as useful optimizations.

2.5.7.1 Driver Requirements: Virtqueue Interrupt Suppression

If the `VIRTIO_F_EVENT_IDX` feature bit is not negotiated:

- The driver MUST set *flags* to 0 or 1.
- The driver MAY set *flags* to 1 to advise the device that interrupts are not needed.

Otherwise, if the `VIRTIO_F_EVENT_IDX` feature bit is negotiated:

- The driver MUST set *flags* to 0.
- The driver MAY use *used_event* to advise the device that interrupts are unnecessary until the device writes entry with an index specified by *used_event* into the used ring (equivalently, until *idx* in the used ring will reach the value *used_event* + 1).

The driver MUST handle spurious interrupts from the device.

2.5.7.2 Device Requirements: Virtqueue Interrupt Suppression

If the `VIRTIO_F_EVENT_IDX` feature bit is not negotiated:

- The device MUST ignore the *used_event* value.
- After the device writes a descriptor index into the used ring:
 - If *flags* is 1, the device SHOULD NOT send an interrupt.
 - If *flags* is 0, the device MUST send an interrupt.

Otherwise, if the VIRTIO_F_EVENT_IDX feature bit is negotiated:

- The device MUST ignore the lower bit of *flags*.
- After the device writes a descriptor index into the used ring:
 - If the *idx* field in the used ring (which determined where that descriptor index was placed) was equal to *used_event*, the device MUST send an interrupt.
 - Otherwise the device SHOULD NOT send an interrupt.

Note: For example, if *used_event* is 0, then a device using VIRTIO_F_EVENT_IDX would interrupt after the first buffer is used (and again after the 65536th buffer, etc).

2.5.8 The Virtqueue Used Ring

```
struct virtq_used {
#define VIRTQ_USED_F_NO_NOTIFY 1
    le16 flags;
    le16 idx;
    struct virtq_used_elem ring[ /* Queue Size */;
    le16 avail_event; /* Only if VIRTIO_F_EVENT_IDX */
};

/* le32 is used here for ids for padding reasons. */
struct virtq_used_elem {
    /* Index of start of used descriptor chain. */
    le32 id;
    /* Total length of the descriptor chain which was used (written to) */
    le32 len;
};
```

The used ring is where the device returns buffers once it is done with them: it is only written to by the device, and read by the driver.

Each entry in the ring is a pair: *id* indicates the head entry of the descriptor chain describing the buffer (this matches an entry placed in the available ring by the guest earlier), and *len* the total of bytes written into the buffer.

Note: *len* is particularly useful for drivers using untrusted buffers: if a driver does not know exactly how much has been written by the device, the driver would have to zero the buffer in advance to ensure no data leakage occurs.

For example, a network driver may hand a received buffer directly to an unprivileged userspace application. If the network device has not overwritten the bytes which were in that buffer, this could leak the contents of freed memory from other processes to the application.

idx field indicates where the driver would put the next descriptor entry in the ring (modulo the queue size). This starts at 0, and increases.

Note: The legacy [Virtio PCI Draft] referred to these structures as *vring_used* and *vring_used_elem*, and the constant as *VRING_USED_F_NO_NOTIFY*, but the layout and value were identical.

2.5.8.1 Legacy Interface: The Virtqueue Used Ring

Historically, many drivers ignored the *len* value, as a result, many devices set *len* incorrectly. Thus, when using the legacy interface, it is generally a good idea to ignore the *len* value in used ring entries if possible. Specific known issues are listed per device type.

2.5.8.2 Device Requirements: The Virtqueue Used Ring

The device MUST set *len* prior to updating the used *idx*.

The device **MUST** write at least *len* bytes to descriptor, beginning at the first device-writable buffer, prior to updating the used *idx*.

The device **MAY** write more than *len* bytes to descriptor.

Note: There are potential error cases where a device might not know what parts of the buffers have been written. This is why *len* is permitted to be an underestimate: that's preferable to the driver believing that uninitialized memory has been overwritten when it has not.

2.5.8.3 Driver Requirements: The Virtqueue Used Ring

The driver **MUST NOT** make assumptions about data in device-writable buffers beyond the first *len* bytes, and **SHOULD** ignore this data.

2.5.9 Virtqueue Notification Suppression

The device can suppress notifications in a manner analogous to the way drivers can suppress interrupts as detailed in section 2.5.7. The device manipulates *flags* or *avail_event* in the used ring the same way the driver manipulates *flags* or *used_event* in the available ring.

2.5.9.1 Driver Requirements: Virtqueue Notification Suppression

The driver **MUST** initialize *flags* in the used ring to 0 when allocating the used ring.

If the VIRTIO_F_EVENT_IDX feature bit is not negotiated:

- The driver **MUST** ignore the *avail_event* value.
- After the driver writes a descriptor index into the available ring:
 - If *flags* is 1, the driver **SHOULD NOT** send a notification.
 - If *flags* is 0, the driver **MUST** send a notification.

Otherwise, if the VIRTIO_F_EVENT_IDX feature bit is negotiated:

- The driver **MUST** ignore the lower bit of *flags*.
- After the driver writes a descriptor index into the available ring:
 - If the *idx* field in the available ring (which determined where that descriptor index was placed) was equal to *avail_event*, the driver **MUST** send a notification.
 - Otherwise the driver **SHOULD NOT** send a notification.

2.5.9.2 Device Requirements: Virtqueue Notification Suppression

If the VIRTIO_F_EVENT_IDX feature bit is not negotiated:

- The device **MUST** set *flags* to 0 or 1.
- The device **MAY** set *flags* to 1 to advise the driver that notifications are not needed.

Otherwise, if the VIRTIO_F_EVENT_IDX feature bit is negotiated:

- The device **MUST** set *flags* to 0.
- The device **MAY** use *avail_event* to advise the driver that notifications are unnecessary until the driver writes entry with an index specified by *avail_event* into the available ring (equivalently, until *idx* in the available ring will reach the value *avail_event* + 1).

The device **MUST** handle spurious notifications from the driver.

2.5.10 Helpers for Operating Virtqueues

The Linux Kernel Source code contains the definitions above and helper routines in a more usable form, in `include/uapi/linux/virtio_ring.h`. This was explicitly licensed by IBM and Red Hat under the (3-clause) BSD license so that it can be freely used by all other projects, and is reproduced (with slight variation) in [A virtio_queue.h](#).

2.5.11 Virtqueue Operation

There are two parts to virtqueue operation: supplying new available buffers to the device, and processing used buffers from the device.

Note: As an example, the simplest virtio network device has two virtqueues: the transmit virtqueue and the receive virtqueue. The driver adds outgoing (device-readable) packets to the transmit virtqueue, and then frees them after they are used. Similarly, incoming (device-writable) buffers are added to the receive virtqueue, and processed after they are used.

What follows is the requirements of each of these two parts when using the split virtqueue format in more detail.

2.5.12 Supplying Buffers to The Device

The driver offers buffers to one of the device's virtqueues as follows:

1. The driver places the buffer into free descriptor(s) in the descriptor table, chaining as necessary (see [2.5.5 The Virtqueue Descriptor Table](#)).
2. The driver places the index of the head of the descriptor chain into the next ring entry of the available ring.
3. Steps 1 and 2 MAY be performed repeatedly if batching is possible.
4. The driver performs suitable a memory barrier to ensure the device sees the updated descriptor table and available ring before the next step.
5. The available *idx* is increased by the number of descriptor chain heads added to the available ring.
6. The driver performs a suitable memory barrier to ensure that it updates the *idx* field before checking for notification suppression.
7. If notifications are not suppressed, the driver notifies the device of the new available buffers.

Note that the above code does not take precautions against the available ring buffer wrapping around: this is not possible since the ring buffer is the same size as the descriptor table, so step (1) will prevent such a condition.

In addition, the maximum queue size is 32768 (the highest power of 2 which fits in 16 bits), so the 16-bit *idx* value can always distinguish between a full and empty buffer.

What follows is the requirements of each stage in more detail.

2.5.12.1 Placing Buffers Into The Descriptor Table

A buffer consists of zero or more device-readable physically-contiguous elements followed by zero or more physically-contiguous device-writable elements (each has at least one element). This algorithm maps it into the descriptor table to form a descriptor chain:

for each buffer element, b:

1. Get the next free descriptor table entry, d
2. Set *d.addr* to the physical address of the start of b

3. Set *d.len* to the length of *b*.
4. If *b* is device-writable, set *d.flags* to `VIRTQ_DESC_F_WRITE`, otherwise 0.
5. If there is a buffer element after this:
 - (a) Set *d.next* to the index of the next free descriptor element.
 - (b) Set the `VIRTQ_DESC_F_NEXT` bit in *d.flags*.

In practice, *d.next* is usually used to chain free descriptors, and a separate count kept to check there are enough free descriptors before beginning the mappings.

2.5.12.2 Updating The Available Ring

The descriptor chain head is the first *d* in the algorithm above, ie. the index of the descriptor table entry referring to the first part of the buffer. A naive driver implementation MAY do the following (with the appropriate conversion to-and-from little-endian assumed):

```
avail->ring[avail->idx
```

However, in general the driver MAY add many descriptor chains before it updates *idx* (at which point they become visible to the device), so it is common to keep a counter of how many the driver has added:

```
avail->ring[(avail->idx + added++)
```

2.5.12.3 Updating *idx*

idx always increments, and wraps naturally at 65536:

```
avail->idx += added;
```

Once available *idx* is updated by the driver, this exposes the descriptor and its contents. The device MAY access the descriptor chains the driver created and the memory they refer to immediately.

2.5.12.3.1 Driver Requirements: Updating *idx*

The driver MUST perform a suitable memory barrier before the *idx* update, to ensure the device sees the most up-to-date copy.

2.5.12.4 Notifying The Device

The actual method of device notification is bus-specific, but generally it can be expensive. So the device MAY suppress such notifications if it doesn't need them, as detailed in section [2.5.9](#).

The driver has to be careful to expose the new *idx* value before checking if notifications are suppressed.

2.5.12.4.1 Driver Requirements: Notifying The Device

The driver MUST perform a suitable memory barrier before reading *flags* or *avail_event*, to avoid missing a notification.

2.5.13 Receiving Used Buffers From The Device

Once the device has used buffers referred to by a descriptor (read from or written to them, or parts of both, depending on the nature of the virtqueue and the device), it interrupts the driver as detailed in section 2.5.7.

Note: For optimal performance, a driver MAY disable interrupts while processing the used ring, but beware the problem of missing interrupts between emptying the ring and reenabling interrupts. This is usually handled by re-checking for more used buffers after interrupts are re-enabled:

```
virtq_disable_interrupts(vq);

for (;;) {
    if (vq->last_seen_used != le16_to_cpu(virtq->used.idx)) {
        virtq_enable_interrupts(vq);
        mb();

        if (vq->last_seen_used != le16_to_cpu(virtq->used.idx))
            break;

        virtq_disable_interrupts(vq);
    }

    struct virtq_used_elem *e = virtq.used->ring[vq->last_seen_used]
        ↪ process_buffer(e);
    vq->last_seen_used++;
}
```

2.6 Packed Virtqueues

Packed virtqueues is an alternative compact virtqueue layout using read-write memory, that is memory that is both read and written by both host and guest.

Use of packed virtqueues is negotiated by the VIRTIO_F_RING_PACKED feature bit.

Packed virtqueues support up to 2^{15} entries each.

With current transports, virtqueues are located in guest memory allocated by driver. Each packed virtqueue consists of three parts:

- Descriptor Ring - occupies the Descriptor Area
- Driver Event Suppression - occupies the Driver Area
- Device Event Suppression - occupies the Device Area

Where Descriptor Ring in turn consists of descriptors, and where each descriptor can contain the following parts:

- Buffer ID
- Element Address
- Element Length
- Flags

A buffer consists of zero or more device-readable physically-contiguous elements followed by zero or more physically-contiguous device-writable elements (each buffer has at least one element).

When the driver wants to send such a buffer to the device, it writes at least one available descriptor describing elements of the buffer into the Descriptor Ring. The descriptor(s) are associated with a buffer by means of a Buffer ID stored within the descriptor.

Driver then notifies the device. When the device has finished processing the buffer, it writes a used device descriptor including the Buffer ID into the Descriptor Ring (overwriting a driver descriptor previously made available), and sends an interrupt.

Descriptor Ring is used in a circular manner: driver writes descriptors into the ring in order. After reaching end of ring, the next descriptor is placed at head of the ring. Once ring is full of driver descriptors, driver stops sending new requests and waits for device to start processing descriptors and to write out some used descriptors before making new driver descriptors available.

Similarly, device reads descriptors from the ring in order and detects that a driver descriptor has been made available. As processing of descriptors is completed used descriptors are written by the device back into the ring.

Note: after reading driver descriptors and starting their processing in order, device might complete their processing out of order. Used device descriptors are written in the order in which their processing is complete.

Device Event Suppression data structure is write-only by the device. It includes information for reducing the number of device events - i.e. driver notifications to device.

Driver Event Suppression data structure is read-only by the device. It includes information for reducing the number of driver events - i.e. device interrupts to driver.

2.6.1 Driver and Device Ring Wrap Counters

Each of the driver and the device are expected to maintain, internally, a single-bit ring wrap counter initialized to 1.

The counter maintained by the driver is called the Driver Ring Wrap Counter. Driver changes the value of this counter each time it makes available the last descriptor in the ring (after making the last descriptor available).

The counter maintained by the device is called the Device Ring Wrap Counter. Device changes the value of this counter each time it uses the last descriptor in the ring (after marking the last descriptor used).

It is easy to see that the Driver Ring Wrap Counter in the driver matches the Device Ring Wrap Counter in the device when both are processing the same descriptor, or when all available descriptors have been used.

To mark a descriptor as available and used, both driver and device use the following two flags:

```
#define VIRTQ_DESC_F_AVAIL    (1 << 7)
#define VIRTQ_DESC_F_USED    (1 << 15)
```

To mark a descriptor as available, driver sets the VIRTQ_DESC_F_AVAIL bit in Flags to match the internal Driver Ring Wrap Counter. It also sets the VIRTQ_DESC_F_USED bit to match the *inverse* value (i.e. to not match the internal Driver Ring Wrap Counter).

To mark a descriptor as used, device sets the VIRTQ_DESC_F_USED bit in Flags to match the internal Device Ring Wrap Counter. It also sets the VIRTQ_DESC_F_AVAIL bit to match the *same* value.

Thus VIRTQ_DESC_F_AVAIL and VIRTQ_DESC_F_USED bits are different for an available descriptor and equal for a used descriptor.

Note that this observation is mostly useful for sanity-checking as these are necessary but not sufficient conditions - for example, all descriptors are zero-initialized. To detect used and available descriptors it is possible for drivers and devices to keep track of the last observed value of VIRTQ_DESC_F_USED/VIRTQ_DESC_F_AVAIL. Other techniques to detect VIRTQ_DESC_F_AVAIL/VIRTQ_DESC_F_USED bit changes might also be possible.

2.6.2 Polling of available and used descriptors

Writes of device and driver descriptors can generally be reordered, but each side (driver and device) are only required to poll (or test) a single location in memory: next device descriptor after the one they processed previously, in circular order.

Sometimes device needs to only write out a single used descriptor after processing a batch of multiple available descriptors. As described in more detail below, this can happen when using descriptor chaining

or with in-order use of descriptors. In this case, device writes out a used descriptor with buffer id of the last descriptor in the group. After processing the used descriptor, both device and driver then skip forward in the ring the number of the remaining descriptors in the group until processing (reading for the driver and writing for the device) the next used descriptor.

2.6.3 Write Flag

In an available descriptor, `VIRTQ_DESC_F_WRITE` bit within Flags is used to mark a descriptor as corresponding to a write-only or read-only element of a buffer.

```
/* This marks a descriptor as device write-only (otherwise device read-only). */
#define VIRTQ_DESC_F_WRITE      2
```

In a used descriptor, this bit is used to specify whether any data has been written by the device into any parts of the buffer.

2.6.4 Element Address and Length

In an available descriptor, Element Address corresponds to the physical address of the buffer element. The length of the element assumed to be physically contiguous is stored in Element Length.

In a used descriptor, Element Address is unused. Element Length specifies the length of the buffer that has been initialized (written to) by the device.

Element length is reserved for used descriptors without the `VIRTQ_DESC_F_WRITE` flag, and is ignored by drivers.

2.6.5 Scatter-Gather Support

Some drivers need an ability to supply a list of multiple buffer elements (also known as a scatter/gather list) with a request. Two features support this: descriptor chaining and indirect descriptors.

If neither feature is in use by the driver, each buffer is physically-contiguous, either read-only or write-only and is described completely by a single descriptor.

While unusual (most implementations either create all lists solely using non-indirect descriptors, or always use a single indirect element), if both features have been negotiated, mixing direct and indirect descriptors in a ring is valid, as long as each list only contains descriptors of a given type.

Scatter/gather lists only apply to available descriptors. A single used descriptor corresponds to the whole list.

The device limits the number of descriptors in a list through a transport-specific and/or device-specific value. If not limited, the maximum number of descriptors in a list is the virt queue size.

2.6.6 Next Flag: Descriptor Chaining

The packed ring format allows driver to supply a scatter/gather list to the device by using multiple descriptors, and setting the `VIRTQ_DESC_F_NEXT` in Flags for all but the last available descriptor.

```
/* This marks a buffer as continuing. */
#define VIRTQ_DESC_F_NEXT      1
```

Buffer ID is included in the last descriptor in the list.

The driver always makes the first descriptor in the list available after the rest of the list has been written out into the ring. This guarantees that the device will never observe a partial scatter/gather list in the ring.

Note: all flags, including VIRTQ_DESC_F_AVAIL, VIRTQ_DESC_F_USED, VIRTQ_DESC_F_WRITE must be set/cleared correctly in all descriptors in the list, not just the first one.

Device only writes out a single used descriptor for the whole list. It then skips forward according to the number of descriptors in the list. Driver needs to keep track of the size of the list corresponding to each buffer ID, to be able to skip to where the next used descriptor is written by the device.

For example, if descriptors are used in the same order in which they are made available, this will result in the used descriptor overwriting the first available descriptor in the list, the used descriptor for the next list overwriting the first available descriptor in the next list, etc.

VIRTQ_DESC_F_NEXT is reserved in used descriptors, and should be ignored by drivers.

2.6.7 Indirect Flag: Scatter-Gather Support

Some devices benefit by concurrently dispatching a large number of large requests. The VIRTIO_F_INDIRECT_DESC feature allows this. To increase ring capacity the driver can store a (read-only by the device) table of indirect descriptors anywhere in memory, and insert a descriptor in main virtqueue (with *Flags* bit VIRTQ_DESC_F_INDIRECT on) that refers to a buffer element containing this indirect descriptor table; *addr* and *len* refer to the indirect table address and length in bytes, respectively.

```
/* This means the element contains a table of descriptors. */
#define VIRTQ_DESC_F_INDIRECT 4
```

The indirect table layout structure looks like this (*len* is the Buffer Length of the descriptor that refers to this table, which is a variable):

```
struct pvirtq_indirect_descriptor_table {
    /* The actual descriptor structures (struct pvirtq_desc each) */
    struct pvirtq_desc desc[len / sizeof(struct pvirtq_desc)];
};
```

The first descriptor is located at start of the indirect descriptor table, additional indirect descriptors come immediately afterwards. *Flags* bit VIRTQ_DESC_F_WRITE is the only valid flag for descriptors in the indirect table. Others are reserved and are ignored by the device. Buffer ID is also reserved and is ignored by the device.

In Descriptors with VIRTQ_DESC_F_INDIRECT set VIRTQ_DESC_F_WRITE is reserved and is ignored by the device.

2.6.8 In-order use of descriptors

Some devices always use descriptors in the same order in which they have been made available. These devices can offer the VIRTIO_F_IN_ORDER feature. If negotiated, this knowledge allows devices to notify the use of a batch of buffers to the driver by only writing out a single used descriptor with the Buffer ID corresponding to the last descriptor in the batch.

Device then skips forward in the ring according to the size of the batch. Driver needs to look up the used Buffer ID and calculate the batch size to be able to advance to where the next used descriptor will be written by the device.

This will result in the used descriptor overwriting the first available descriptor in the batch, the used descriptor for the next batch overwriting the first available descriptor in the next batch, etc.

The skipped buffers (for which no used descriptor was written) are assumed to have been used (read or written) by the device completely.

2.6.9 Multi-buffer requests

Some devices combine multiple buffers as part of processing of a single request. These devices always mark the descriptor corresponding to the first buffer in the request used after the rest of the descriptors (corresponding to rest of the buffers) in the request - which follow the first descriptor in ring order - has been marked used and written out into the ring. This guarantees that the driver will never observe a partial request in the ring.

2.6.10 Driver and Device Event Suppression

In many systems driver and device notifications involve significant overhead. To mitigate this overhead, each virtqueue includes two identical structures used for controlling notifications between device and driver.

Driver Event Suppression structure is read-only by the device and controls the events sent by the device to the driver (e.g. interrupts).

Device Event Suppression structure is read-only by the driver and controls the events sent by the driver to the device (e.g. IO).

Each of these Event Suppression structures controls both Descriptor Ring events and structure events, and each includes the following fields:

Descriptor Ring Change Event Flags Takes values:

```
/* Enable events */
#define RING_EVENT_FLAGS_ENABLE 0x0
/* Disable events */
#define RING_EVENT_FLAGS_DISABLE 0x1
/*
 * Enable events for a specific descriptor
 * (as specified by Descriptor Ring Change Event Offset/Wrap Counter).
 * Only valid if VIRTIO_F_RING_EVENT_IDX has been negotiated.
 */
#define RING_EVENT_FLAGS_DESC 0x2
/* The value 0x3 is reserved */
```

Descriptor Ring Change Event Offset If Event Flags set to descriptor specific event: offset within the ring (in units of descriptor size). Event will only trigger when this descriptor is made available/used respectively.

Descriptor Ring Change Event Wrap Counter If Event Flags set to descriptor specific event: offset within the ring (in units of descriptor size). Event will only trigger when Ring Wrap Counter matches this value and a descriptor is made available/used respectively.

After writing out some descriptors, both device and driver are expected to consult the relevant structure to find out whether interrupt/notification should be sent.

2.6.10.1 Structure Size and Alignment

Each part of the virtqueue is physically-contiguous in guest memory, and has different alignment requirements.

The memory alignment and size requirements, in bytes, of each part of the virtqueue are summarized in the following table:

Virtqueue Part	Alignment	Size
Descriptor Ring	16	16*(Queue Size)
Device Event Suppression	4	4
Driver Event Suppression	4	4

The Alignment column gives the minimum alignment for each part of the virtqueue.

The Size column gives the total number of bytes for each part of the virtqueue.

Queue Size corresponds to the maximum number of descriptors in the virtqueue³. Queue Size value does not have to be a power of 2 unless enforced by the transport.

2.6.11 Driver Requirements: Virtqueues

The driver MUST ensure that the physical address of the first byte of each virtqueue part is a multiple of the specified alignment value in the above table.

2.6.12 Device Requirements: Virtqueues

The device MUST start processing driver descriptors in the order in which they appear in the ring. The device MUST start writing device descriptors into the ring in the order in which they complete. Device MAY reorder descriptor writes once they are started.

2.6.13 The Virtqueue Descriptor Format

The available descriptor refers to the buffers the driver is sending to the device. *addr* is a physical address, and the descriptor is identified with a buffer using the *id* field.

```
struct pvirtq_desc {
    /* Buffer Address. */
    le64 addr;
    /* Buffer Length. */
    le32 len;
    /* Buffer ID. */
    le16 id;
    /* The flags depending on descriptor type. */
    le16 flags;
};
```

The descriptor ring is zero-initialized.

2.6.14 Event Suppression Structure Format

The following structure is used to reduce the number of notifications sent between driver and device.

```
struct pvirtq_event_suppress {
    le16 {
        desc_event_off : 15; /* Descriptor Ring Change Event Offset */
        desc_event_wrap : 1; /* Descriptor Ring Change Event Wrap Counter */
    } desc; /* If desc_event_flags set to RING_EVENT_FLAGS_DESC */
    le16 {
        desc_event_flags : 2; /* Descriptor Ring Change Event Flags */
        reserved : 14; /* Reserved, set to 0 */
    } flags;
};
```

2.6.15 Device Requirements: The Virtqueue Descriptor Table

A device MUST NOT write to a device-readable buffer, and a device SHOULD NOT read a device-writable buffer. A device MUST NOT use a descriptor unless it observes VIRTQ_DESC_F_AVAIL bit in its *flags* being changed (e.g. as compared to the initial zero value). A device MUST NOT change a descriptor after changing it's VIRTQ_DESC_F_USED bit in its *flags*.

³For example, if Queue Size is 4 then at most 4 buffers can be queued at any given time.

2.6.16 Driver Requirements: The Virtqueue Descriptor Table

A driver MUST NOT change a descriptor unless it observes `VIRTQ_DESC_F_USED` bit in its *flags* being changed. A driver MUST NOT change a descriptor after changing `VIRTQ_DESC_F_AVAIL` bit in its *flags*. When notifying the device, driver MUST set *next_off* and *next_wrap* to match the next descriptor not yet made available to the device. A driver MAY send multiple notifications without making any new descriptors available to the device.

2.6.17 Driver Requirements: Scatter-Gather Support

A driver MUST NOT create a descriptor list longer than allowed by the device.

A driver MUST NOT create a descriptor list longer than the Queue Size.

This implies that loops in the descriptor list are forbidden!

The driver MUST place any device-writable descriptor elements after any device-readable descriptor elements.

A driver MUST NOT depend on the device to use more descriptors to be able to write out all descriptors in a list. A driver MUST make sure there's enough space in the ring for the whole list before making the first descriptor in the list available to the device.

A driver MUST NOT make the first descriptor in the list available before all subsequent descriptors comprising the list are made available.

2.6.18 Device Requirements: Scatter-Gather Support

The device MUST use descriptors in a list chained by the `VIRTQ_DESC_F_NEXT` flag in the same order that they were made available by the driver.

The device MAY limit the number of buffers it will allow in a list.

2.6.19 Driver Requirements: Indirect Descriptors

The driver MUST NOT set the `DESC_F_INDIRECT` flag unless the `VIRTIO_F_INDIRECT_DESC` feature was negotiated. The driver MUST NOT set any flags except `DESC_F_WRITE` within an indirect descriptor.

A driver MUST NOT create a descriptor chain longer than allowed by the device.

A driver MUST NOT write direct descriptors with `DESC_F_INDIRECT` set in a scatter-gather list linked by `VIRTQ_DESC_F_NEXT`. *flags*.

2.6.20 Virtqueue Operation

There are two parts to virtqueue operation: supplying new available buffers to the device, and processing used buffers from the device.

What follows is the requirements of each of these two parts when using the packed virtqueue format in more detail.

2.6.21 Supplying Buffers to The Device

The driver offers buffers to one of the device's virtqueues as follows:

1. The driver places the buffer into free descriptor in the Descriptor Ring.

2. The driver performs a suitable memory barrier to ensure that it updates the descriptor(s) before checking for notification suppression.
3. If notifications are not suppressed, the driver notifies the device of the new available buffers.

What follows is the requirements of each stage in more detail.

2.6.21.1 Placing Available Buffers Into The Descriptor Ring

For each buffer element, *b*:

1. Get the next descriptor table entry, *d*
2. Get the next free buffer id value
3. Set *d.addr* to the physical address of the start of *b*
4. Set *d.len* to the length of *b*.
5. Set *d.id* to the buffer id
6. Calculate the flags as follows:
 - (a) If *b* is device-writable, set the VIRTQ_DESC_F_WRITE bit to 1, otherwise 0
 - (b) Set VIRTQ_DESC_F_AVAIL bit to the current value of the Driver Ring Wrap Counter
 - (c) Set VIRTQ_DESC_F_USED bit to inverse value
7. Perform a memory barrier to ensure that the descriptor has been initialized
8. Set *d.flags* to the calculated flags value
9. If *d* is the last descriptor in the ring, toggle the Driver Ring Wrap Counter
10. Otherwise, increment *d* to point at the next descriptor

This makes a single descriptor buffer available. However, in general the driver MAY make use of a batch of descriptors as part of a single request. In that case, it defers updating the descriptor flags for the first descriptor (and the previous memory barrier) until after the rest of the descriptors have been initialized.

Once the descriptor *flags* is updated by the driver, this exposes the descriptor and its contents. The device MAY access the descriptor and any following descriptors the driver created and the memory they refer to immediately.

2.6.21.1.1 Driver Requirements: Updating flags

The driver MUST perform a suitable memory barrier before the *flags* update, to ensure the device sees the most up-to-date copy.

2.6.21.2 Notifying The Device

The actual method of device notification is bus-specific, but generally it can be expensive. So the device MAY suppress such notifications if it doesn't need them, using the Driver Event Suppression structure as detailed in section [2.6.14](#).

The driver has to be careful to expose the new *flags* value before checking if notifications are suppressed.

2.6.21.3 Implementation Example

Below is an example driver code. It does not attempt to reduce the number of device interrupts, neither does it support the `VIRTIO_F_RING_EVENT_IDX` feature.

```
/* Note: vq->avail_wrap_count is initialized to 1 */
/* Note: vq->sgs is an array same size as the ring */

id = alloc_id(vq);

first = vq->next_avail;
sgs = 0;
for (each buffer element b) {
    sgs++;

    vq->ids[vq->next_avail] = -1;
    vq->desc[vq->next_avail].address = get_addr(b);
    vq->desc[vq->next_avail].len = get_len(b);

    avail = vq->avail_wrap_count ? VIRTQ_DESC_F_AVAIL : 0;
    used = !vq->avail_wrap_count ? VIRTQ_DESC_F_USED : 0;
    f = get_flags(b) | avail | used;
    if (b is not the last buffer element) {
        f |= VIRTQ_DESC_F_NEXT;
    }

    /* Don't mark the 1st descriptor available until all of them are ready. */
    if (vq->next_avail == first) {
        flags = f;
    } else {
        vq->desc[vq->next_avail].flags = f;
    }

    last = vq->next_avail;

    vq->next_avail++;

    if (vq->next_avail >= vq->size) {
        vq->next_avail = 0;
        vq->avail_wrap_count ^= 1;
    }
}
vq->sgs[id] = sgs;
/* ID included in the last descriptor in the list */
vq->desc[last].id = id;
write_memory_barrier();
vq->desc[first].flags = flags;

memory_barrier();

if (vq->device_event.flags != RING_EVENT_FLAGS_DISABLE) {
    notify_device(vq);
}
```

2.6.21.3.1 Driver Requirements: Notifying The Device

The driver **MUST** perform a suitable memory barrier before reading the Driver Event Suppression structure, to avoid missing a notification.

2.6.22 Receiving Used Buffers From The Device

Once the device has used buffers referred to by a descriptor (read from or written to them, or parts of both, depending on the nature of the virtqueue and the device), it interrupts the driver as detailed in section [2.6.14](#).

Note: For optimal performance, a driver MAY disable interrupts while processing the used buffers, but beware the problem of missing interrupts between emptying the ring and reenabling interrupts. This is usually handled by re-checking for more used buffers after interrupts are re-enabled:

```
/* Note: vq->used_wrap_count is initialized to 1 */
vq->driver_event.flags = RING_EVENT_FLAGS_DISABLE;

for (;;) {
    struct pvirtq_desc *d = vq->desc[vq->next_used];

    flags = d->flags;
    bool used = flags & VIRTQ_DESC_F_USED;

    if (used != vq->used_wrap_count) {
        vq->driver_event.flags = RING_EVENT_FLAGS_ENABLE;
        memory_barrier();

        flags = d->flags;
        bool used = flags & VIRTQ_DESC_F_USED;
        if (used != vq->used_wrap_count) {
            break;
        }

        vq->driver_event.flags = RING_EVENT_FLAGS_DISABLE;
    }

    read_memory_barrier();

    /* skip descriptors until the next buffer */
    id = d->id;
    assert(id < vq->size);
    sgs = vq->sgs[id];
    vq->next_used += sgs;
    if (vq->next_used >= vq->size) {
        vq->next_used -= vq->size;
        vq->used_wrap_count ^= 1;
    }

    free_id(vq, id);

    process_buffer(d);
}
```

2.6.22.1 Driver notifications

~~Whenever not suppressed by Device Event Suppression, driver is required to notify the device after making changes to the virtqueue.~~

~~Some devices benefit from ability to find out the number of available descriptors in the ring, and whether to send interrupts to drivers without accessing virtqueue in memory: for efficiency or as a debugging aid.~~

~~To help with these optimizations, driver notifications to the device include the following information:~~

~~VQ number Offset (in units of descriptor size) within the ring where the next available descriptor will be written Wrap Counter referring to the next available descriptor~~

~~Note that driver can trigger multiple notifications even without making any more changes to the ring. These would then have identical Offset and Wrap Counter values.~~

3 General Initialization And Device Operation

We start with an overview of device initialization, then expand on the details of the device and how each step is preformed. This section is best read along with the bus-specific section which describes how to communicate with the specific device.

3.1 Device Initialization

3.1.1 Driver Requirements: Device Initialization

The driver **MUST** follow this sequence to initialize a device:

1. Reset the device.
2. Set the ACKNOWLEDGE status bit: the guest OS has notice the device.
3. Set the DRIVER status bit: the guest OS knows how to drive the device.
4. Read device feature bits, and write the subset of feature bits understood by the OS and driver to the device. During this step the driver **MAY** read (but **MUST NOT** write) the device-specific configuration fields to check that it can support the device before accepting it.
5. Set the FEATURES_OK status bit. The driver **MUST NOT** accept new feature bits after this step.
6. Re-read *device status* to ensure the FEATURES_OK bit is still set: otherwise, the device does not support our subset of features and the device is unusable.
7. Perform device-specific setup, including discovery of virtqueues for the device, optional per-bus setup, reading and possibly writing the device's virtio configuration space, and population of virtqueues.
8. Set the DRIVER_OK status bit. At this point the device is "live".

If any of these steps go irrecoverably wrong, the driver **SHOULD** set the FAILED status bit to indicate that it has given up on the device (it can reset the device later to restart if desired). The driver **MUST NOT** continue initialization in that case.

The driver **MUST NOT** notify the device before setting DRIVER_OK.

3.1.2 Legacy Interface: Device Initialization

Legacy devices did not support the FEATURES_OK status bit, and thus did not have a graceful way for the device to indicate unsupported feature combinations. They also did not provide a clear mechanism to end feature negotiation, which meant that devices finalized features on first-use, and no features could be introduced which radically changed the initial operation of the device.

Legacy driver implementations often used the device before setting the DRIVER_OK bit, and sometimes even before writing the feature bits to the device.

The result was the steps 5 and 6 were omitted, and steps 4, 7 and 8 were conflated.

Therefore, when using the legacy interface:

- The transitional driver **MUST** execute the initialization sequence as described in 3.1 but omitting the steps 5 and 6.

- The transitional device **MUST** support the driver writing device configuration fields before the step [4](#).
- The transitional device **MUST** support the driver using the device before the step [8](#).

3.2 Device Operation

When operating the device, each field in the device configuration space can be changed by either the driver or the device.

Whenever such a configuration change is triggered by the device, driver is notified. This makes it possible for drivers to cache device configuration, avoiding expensive configuration reads unless notified.

3.2.1 Notification of Device Configuration Changes

For devices where the device-specific configuration information can be changed, an interrupt is delivered when a device-specific configuration change occurs.

In addition, this interrupt is triggered by the device setting `DEVICE_NEEDS_RESET` (see [2.1.2](#)).

3.3 Device Cleanup

Once the driver has set the `DRIVER_OK` status bit, all the configured virtqueue of the device are considered live. None of the virtqueues of a device are live once the device has been reset.

3.3.1 Driver Requirements: Device Cleanup

A driver **MUST NOT** alter virtqueue entries for exposed buffers - i.e. buffers which have been made available to the device (and not been used by the device) of a live virtqueue.

Thus a driver **MUST** ensure a virtqueue isn't live (by device reset) before removing exposed buffers.

4 Virtio Transport Options

Virtio can use various different buses, thus the standard is split into virtio general and bus-specific sections.

4.1 Virtio Over PCI Bus

Virtio devices are commonly implemented as PCI devices.

A Virtio device can be implemented as any kind of PCI device: a Conventional PCI device or a PCI Express device. To assure designs meet the latest level requirements, see the PCI-SIG home page at <http://www.pcisig.com> for any approved changes.

4.1.1 Device Requirements: Virtio Over PCI Bus

A Virtio device using Virtio Over PCI Bus MUST expose to guest an interface that meets the specification requirements of the appropriate PCI specification: [PCI] and [PCIe] respectively.

4.1.2 PCI Device Discovery

Any PCI device with PCI Vendor ID 0x1AF4, and PCI Device ID 0x1000 through 0x107F inclusive is a virtio device. The actual value within this range indicates which virtio device is supported by the device. The PCI Device ID is calculated by adding 0x1040 to the Virtio Device ID, as indicated in section 5. Additionally, devices MAY utilize a Transitional PCI Device ID range, 0x1000 to 0x103F depending on the device type.

4.1.2.1 Device Requirements: PCI Device Discovery

Devices MUST have the PCI Vendor ID 0x1AF4. Devices MUST either have the PCI Device ID calculated by adding 0x1040 to the Virtio Device ID, as indicated in section 5 or have the Transitional PCI Device ID depending on the device type, as follows:

Transitional PCI Device ID	Virtio Device
0x1000	network card
0x1001	block device
0x1002	memory ballooning (traditional)
0x1003	console
0x1004	SCSI host
0x1005	entropy source
0x1009	9P transport

For example, the network card device with the Virtio Device ID 1 has the PCI Device ID 0x1041 or the Transitional PCI Device ID 0x1000.

The PCI Subsystem Vendor ID and the PCI Subsystem Device ID MAY reflect the PCI Vendor and Device ID of the environment (for informational purposes by the driver).

Non-transitional devices SHOULD have a PCI Device ID in the range 0x1040 to 0x107f. Non-transitional devices SHOULD have a PCI Revision ID of 1 or higher. Non-transitional devices SHOULD have a PCI Subsystem Device ID of 0x40 or higher.

This is to reduce the chance of a legacy driver attempting to drive the device.

4.1.2.2 Driver Requirements: PCI Device Discovery

Drivers MUST match devices with the PCI Vendor ID 0x1AF4 and the PCI Device ID in the range 0x1040 to 0x107f, calculated by adding 0x1040 to the Virtio Device ID, as indicated in section 5. Drivers for device types listed in section 4.1.2 MUST match devices with the PCI Vendor ID 0x1AF4 and the Transitional PCI Device ID indicated in section 4.1.2.

Drivers MUST match any PCI Revision ID value. Drivers MAY match any PCI Subsystem Vendor ID and any PCI Subsystem Device ID value.

4.1.2.3 Legacy Interfaces: A Note on PCI Device Discovery

Transitional devices MUST have a PCI Revision ID of 0. Transitional devices MUST have the PCI Subsystem Device ID matching the Virtio Device ID, as indicated in section 5. Transitional devices MUST have the Transitional PCI Device ID in the range 0x1000 to 0x103f.

This is to match legacy drivers.

4.1.3 PCI Device Layout

The device is configured via I/O and/or memory regions (though see 4.1.4.7 for access via the PCI configuration space), as specified by Virtio Structure PCI Capabilities.

Fields of different sizes are present in the device configuration regions. All 64-bit, 32-bit and 16-bit fields are little-endian. 64-bit fields are to be treated as two 32-bit fields, with low 32 bit part followed by the high 32 bit part.

4.1.3.1 Driver Requirements: PCI Device Layout

For device configuration access, the driver MUST use 8-bit wide accesses for 8-bit wide fields, 16-bit wide and aligned accesses for 16-bit wide fields and 32-bit wide and aligned accesses for 32-bit and 64-bit wide fields. For 64-bit fields, the driver MAY access each of the high and low 32-bit parts of the field independently.

4.1.3.2 Device Requirements: PCI Device Layout

For 64-bit device configuration fields, the device MUST allow driver independent access to high and low 32-bit parts of the field.

4.1.4 Virtio Structure PCI Capabilities

The virtio device configuration layout includes several structures:

- Common configuration
- Notifications
- ISR Status
- Device-specific configuration (optional)
- PCI configuration access

Each structure can be mapped by a Base Address register (BAR) belonging to the function, or accessed via the special VIRTIO_PCI_CAP_PCI_CFG field in the PCI configuration space.

The location of each structure is specified using a vendor-specific PCI capability located on the capability list in PCI configuration space of the device. This virtio structure capability uses little-endian format; all fields are read-only for the driver unless stated otherwise:

```
struct virtio_pci_cap {
    u8 cap_vndr;    /* Generic PCI field: PCI_CAP_ID_VNDR */
    u8 cap_next;    /* Generic PCI field: next ptr. */
    u8 cap_len;     /* Generic PCI field: capability length */
    u8 cfg_type;    /* Identifies the structure. */
    u8 bar;         /* Where to find it. */
    u8 padding[3];  /* Pad to full dword. */
    le32 offset;    /* Offset within bar. */
    le32 length;    /* Length of the structure, in bytes. */
};
```

This structure can be followed by extra data, depending on *cfg_type*, as documented below.

The fields are interpreted as follows:

cap_vndr 0x09; Identifies a vendor-specific capability.

cap_next Link to next capability in the capability list in the PCI configuration space.

cap_len Length of this capability structure, including the whole of struct virtio_pci_cap, and extra data if any. This length MAY include padding, or fields unused by the driver.

cfg_type identifies the structure, according to the following table:

```
/* Common configuration */
#define VIRTIO_PCI_CAP_COMMON_CFG      1
/* Notifications */
#define VIRTIO_PCI_CAP_NOTIFY_CFG      2
/* ISR Status */
#define VIRTIO_PCI_CAP_ISR_CFG         3
/* Device specific configuration */
#define VIRTIO_PCI_CAP_DEVICE_CFG      4
/* PCI configuration access */
#define VIRTIO_PCI_CAP_PCI_CFG         5
```

Any other value is reserved for future use.

Each structure is detailed individually below.

The device MAY offer more than one structure of any type - this makes it possible for the device to expose multiple interfaces to drivers. The order of the capabilities in the capability list specifies the order of preference suggested by the device.

Note: For example, on some hypervisors, notifications using IO accesses are faster than memory accesses. In this case, the device would expose two capabilities with *cfg_type* set to VIRTIO_PCI_CAP_NOTIFY_CFG: the first one addressing an I/O BAR, the second one addressing a memory BAR. In this example, the driver would use the I/O BAR if I/O resources are available, and fall back on memory BAR when I/O resources are unavailable.

bar values 0x0 to 0x5 specify a Base Address register (BAR) belonging to the function located beginning at 10h in PCI Configuration Space and used to map the structure into Memory or I/O Space. The BAR is permitted to be either 32-bit or 64-bit, it can map Memory Space or I/O Space.

Any other value is reserved for future use.

offset indicates where the structure begins relative to the base address associated with the BAR. The alignment requirements of *offset* are indicated in each structure-specific section below.

length indicates the length of the structure.

length MAY include padding, or fields unused by the driver, or future extensions.

Note: For example, a future device might present a large structure size of several MBytes. As current devices never utilize structures larger than 4KBytes in size, driver MAY limit the mapped structure size to e.g. 4KBytes (thus ignoring parts of structure after the first 4KBytes) to allow forward compatibility with such devices without loss of functionality and without wasting resources.

4.1.4.1 Driver Requirements: Virtio Structure PCI Capabilities

The driver MUST ignore any vendor-specific capability structure which has a reserved *cfg_type* value.

The driver SHOULD use the first instance of each virtio structure type they can support.

The driver MUST accept a *cap_len* value which is larger than specified here.

The driver MUST ignore any vendor-specific capability structure which has a reserved *bar* value.

The drivers SHOULD only map part of configuration structure large enough for device operation. The drivers MUST handle an unexpectedly large *length*, but MAY check that *length* is large enough for device operation.

The driver MUST NOT write into any field of the capability structure, with the exception of those with *cap_type* VIRTIO_PCI_CAP_PCI_CFG as detailed in 4.1.4.7.2.

4.1.4.2 Device Requirements: Virtio Structure PCI Capabilities

The device MUST include any extra data (from the beginning of the *cap_vndr* field through end of the extra data fields if any) in *cap_len*. The device MAY append extra data or padding to any structure beyond that.

If the device presents multiple structures of the same type, it SHOULD order them from optimal (first) to least-optimal (last).

4.1.4.3 Common configuration structure layout

The common configuration structure is found at the *bar* and *offset* within the VIRTIO_PCI_CAP_COMMON_-CFG capability; its layout is below.

```
struct virtio_pci_common_cfg {
    /* About the whole device. */
    le32 device_feature_select; /* read-write */
    le32 device_feature; /* read-only for driver */
    le32 driver_feature_select; /* read-write */
    le32 driver_feature; /* read-write */
    le16 msix_config; /* read-write */
    le16 num_queues; /* read-only for driver */
    u8 device_status; /* read-write */
    u8 config_generation; /* read-only for driver */

    /* About a specific virtqueue. */
    le16 queue_select; /* read-write */
    le16 queue_size; /* read-write, power of 2, or 0. */
    le16 queue_msix_vector; /* read-write */
    le16 queue_enable; /* read-write */
    le16 queue_notify_off; /* read-only for driver */
    le64 queue_desc; /* read-write */
    le64 queue_driver; /* read-write */
    le64 queue_device; /* read-write */
};
```

device_feature_select The driver uses this to select which feature bits *device_feature* shows. Value 0x0 selects Feature Bits 0 to 31, 0x1 selects Feature Bits 32 to 63, etc.

device_feature The device uses this to report which feature bits it is offering to the driver: the driver writes to *device_feature_select* to select which feature bits are presented.

driver_feature_select The driver uses this to select which feature bits *driver_feature* shows. Value 0x0 selects Feature Bits 0 to 31, 0x1 selects Feature Bits 32 to 63, etc.

driver_feature The driver writes this to accept feature bits offered by the device. Driver Feature Bits selected by *driver_feature_select*.

config_msix_vector The driver sets the Configuration Vector for MSI-X.

num_queues The device specifies the maximum number of virtqueues supported here.

device_status The driver writes the device status here (see 2.1). Writing 0 into this field resets the device.

config_generation Configuration atomicity value. The device changes this every time the configuration noticeably changes.

queue_select Queue Select. The driver selects which virtqueue the following fields refer to.

queue_size Queue Size. On reset, specifies the maximum queue size supported by the hypervisor. This can be modified by driver to reduce memory requirements. A 0 means the queue is unavailable.

queue_msix_vector The driver uses this to specify the queue vector for MSI-X.

queue_enable The driver uses this to selectively prevent the device from executing requests from this virtqueue. 1 - enabled; 0 - disabled.

queue_notify_off The driver reads this to calculate the offset from start of Notification structure at which this virtqueue is located.

Note: this is *not an offset in bytes*. See 4.1.4.4 below.

queue_desc The driver writes the physical address of Descriptor Area here. See section 2.4.

queue_driver The driver writes the physical address of Driver Area here. See section 2.4.

queue_device The driver writes the physical address of Device Area here. See section 2.4.

4.1.4.3.1 Device Requirements: Common configuration structure layout

offset MUST be 4-byte aligned.

The device MUST present at least one common configuration capability.

The device MUST present the feature bits it is offering in *device_feature*, starting at bit *device_feature_select* * 32 for any *device_feature_select* written by the driver.

Note: This means that it will present 0 for any *device_feature_select* other than 0 or 1, since no feature defined here exceeds 63.

The device MUST present any valid feature bits the driver has written in *driver_feature*, starting at bit *driver_feature_select* * 32 for any *driver_feature_select* written by the driver. Valid feature bits are those which are subset of the corresponding *device_feature* bits. The device MAY present invalid bits written by the driver.

Note: This means that a device can ignore writes for feature bits it never offers, and simply present 0 on reads. Or it can just mirror what the driver wrote (but it will still have to check them when the driver sets FEATURES_OK).

Note: A driver shouldn't write invalid bits anyway, as per 3.1.1, but this attempts to handle it.

The device MUST present a changed *config_generation* after the driver has read a device-specific configuration value which has changed since any part of the device-specific configuration was last read.

Note: As *config_generation* is an 8-bit value, simply incrementing it on every configuration change could violate this requirement due to wrap. Better would be to set an internal flag when it has changed, and if that flag is set when the driver reads from the device-specific configuration, increment *config_generation* and clear the flag.

The device MUST reset when 0 is written to *device_status*, and present a 0 in *device_status* once that is done.

The device MUST present a 0 in *queue_enable* on reset.

The device MUST present a 0 in *queue_size* if the virtqueue corresponding to the current *queue_select* is unavailable.

4.1.4.3.2 Driver Requirements: Common configuration structure layout

The driver MUST NOT write to *device_feature*, *num_queues*, *config_generation* or *queue_notify_off*.

The driver MUST NOT write a value which is not a power of 2 to *queue_size*.

The driver MUST configure the other virtqueue fields before enabling the virtqueue with *queue_enable*.

After writing 0 to *device_status*, the driver MUST wait for a read of *device_status* to return 0 before reinitializing the device.

The driver MUST NOT write a 0 to *queue_enable*.

4.1.4.4 Notification structure layout

The notification location is found using the VIRTIO_PCI_CAP_NOTIFY_CFG capability. This capability is immediately followed by an additional field, like so:

```
struct virtio_pci_notify_cap {
    struct virtio_pci_cap cap;
    le32 notify_off_multiplier; /* Multiplier for queue_notify_off. */
};
```

notify_off_multiplier is combined with the *queue_notify_off* to derive the Queue Notify address within a BAR for a virtqueue:

$$\text{cap.offset} + \text{queue_notify_off} * \text{notify_off_multiplier}$$

The *cap.offset* and *notify_off_multiplier* are taken from the notification capability structure above, and the *queue_notify_off* is taken from the common configuration structure.

Note: For example, if *notify_off_multiplier* is 0, the device uses the same Queue Notify address for all queues.

4.1.4.4.1 Device Requirements: Notification capability

The device MUST present at least one notification capability.

~~For devices not offering VIRTIO_F_NOTIFICATION_DATA:-~~

The *cap.offset* MUST be 2-byte aligned.

The device MUST either present *notify_off_multiplier* as an even power of 2, or present *notify_off_multiplier* as 0.

The value *cap.length* presented by the device MUST be at least 2 and MUST be large enough to support queue notification offsets for all supported queues in all possible configurations.

For all queues, the value *cap.length* presented by the device MUST satisfy:

$$\text{cap.length} \geq \text{queue_notify_off} * \text{notify_off_multiplier} + 2$$

~~For devices offering VIRTIO_F_NOTIFICATION_DATA:-~~

~~The device MUST either present *notify_off_multiplier* as a number that is a power of 2 that is also a multiple 4, or present *notify_off_multiplier* as 0.-~~

~~The *cap.offset* MUST be 4-byte aligned.-~~

The value *cap.length* presented by the device MUST be at least 4 and MUST be large enough to support queue notification offsets for all supported queues in all possible configurations.

For all queues, the value *cap.length* presented by the device MUST satisfy:

$$\text{cap.length} \geq \text{queue_notify_off} * \text{notify_off_multiplier} + 4$$

4.1.4.5 ISR status capability

The VIRTIO_PCI_CAP_ISR_CFG capability refers to at least a single byte, which contains the 8-bit ISR status field to be used for INT#x interrupt handling.

The *offset* for the *ISR status* has no alignment requirements.

The ISR bits allow the device to distinguish between device-specific configuration change interrupts and normal virtqueue interrupts:

Bits	0	1	2 to 31
Purpose	Queue Interrupt	Device Configuration Interrupt	Reserved

To avoid an extra access, simply reading this register resets it to 0 and causes the device to de-assert the interrupt.

In this way, driver read of ISR status causes the device to de-assert an interrupt.

See sections 4.1.5.3 and 4.1.5.4 for how this is used.

4.1.4.5.1 Device Requirements: ISR status capability

The device MUST present at least one VIRTIO_PCI_CAP_ISR_CFG capability.

The device MUST set the Device Configuration Interrupt bit in *ISR status* before sending a device configuration change notification to the driver.

If MSI-X capability is disabled, the device MUST set the Queue Interrupt bit in *ISR status* before sending a virtqueue notification to the driver.

If MSI-X capability is disabled, the device MUST set the Interrupt Status bit in the PCI Status register in the PCI Configuration Header of the device to the logical OR of all bits in *ISR status* of the device. The device then asserts/deasserts INT#x interrupts unless masked according to standard PCI rules [PCI].

The device MUST reset *ISR status* to 0 on driver read.

4.1.4.5.2 Driver Requirements: ISR status capability

If MSI-X capability is enabled, the driver SHOULD NOT access *ISR status* upon detecting a Queue Interrupt.

4.1.4.6 Device-specific configuration

The device MUST present at least one VIRTIO_PCI_CAP_DEVICE_CFG capability for any device type which has a device-specific configuration.

4.1.4.6.1 Device Requirements: Device-specific configuration

The *offset* for the device-specific configuration MUST be 4-byte aligned.

4.1.4.7 PCI configuration access capability

The VIRTIO_PCI_CAP_PCI_CFG capability creates an alternative (and likely suboptimal) access method to the common configuration, notification, ISR and device-specific configuration regions.

The capability is immediately followed by an additional field like so:

```
struct virtio_pci_cfg_cap {
    struct virtio_pci_cap cap;
    u8 pci_cfg_data[4]; /* Data for BAR access. */
};
```

The fields *cap.bar*, *cap.length*, *cap.offset* and *pci_cfg_data* are read-write (RW) for the driver.

To access a device region, the driver writes into the capability structure (ie. within the PCI configuration space) as follows:

- The driver sets the BAR to access by writing to *cap.bar*.
- The driver sets the size of the access by writing 1, 2 or 4 to *cap.length*.
- The driver sets the offset within the BAR by writing to *cap.offset*.

At that point, *pci_cfg_data* will provide a window of size *cap.length* into the given *cap.bar* at offset *cap.offset*.

4.1.4.7.1 Device Requirements: PCI configuration access capability

The device MUST present at least one VIRTIO_PCI_CAP_PCI_CFG capability.

Upon detecting driver write access to *pci_cfg_data*, the device MUST execute a write access at offset *cap.offset* at BAR selected by *cap.bar* using the first *cap.length* bytes from *pci_cfg_data*.

Upon detecting driver read access to *pci_cfg_data*, the device MUST execute a read access of length *cap.length* at offset *cap.offset* at BAR selected by *cap.bar* and store the first *cap.length* bytes in *pci_cfg_data*.

4.1.4.7.2 Driver Requirements: PCI configuration access capability

The driver MUST NOT write a *cap.offset* which is not a multiple of *cap.length* (ie. all accesses MUST be aligned).

The driver MUST NOT read or write *pci_cfg_data* unless *cap.bar*, *cap.length* and *cap.offset* address *cap.length* bytes within a BAR range specified by some other Virtio Structure PCI Capability of type other than VIRTIO_PCI_CAP_PCI_CFG.

4.1.4.8 Legacy Interfaces: A Note on PCI Device Layout

Transitional devices MUST present part of configuration registers in a legacy configuration structure in BAR0 in the first I/O region of the PCI device, as documented below. When using the legacy interface, transitional drivers MUST use the legacy configuration structure in BAR0 in the first I/O region of the PCI device, as documented below.

When using the legacy interface the driver MAY access the device-specific configuration region using any width accesses, and a transitional device MUST present driver with the same results as when accessed using the “natural” access method (i.e. 32-bit accesses for 32-bit fields, etc).

Note that this is possible because while the virtio common configuration structure is PCI (i.e. little) endian, when using the legacy interface the device-specific configuration region is encoded in the native endian of the guest (where such distinction is applicable).

When used through the legacy interface, the virtio common configuration structure looks as follows:

Bits	32	32	32	16	16	16	8	8
Read / Write	R	R+W	R+W	R	R+W	R+W	R+W	R
Purpose	Device Features bits 0:31	Driver Features bits 0:31	Queue Address	<i>queue_size</i>	<i>queue_select</i>	Queue Notify	Device Status	ISR Status

If MSI-X is enabled for the device, two additional fields immediately follow this header:

Bits	16	16
Read/Write	R+W	R+W
Purpose (MSI-X)	<i>config_msix_vector</i>	<i>queue_msix_vector</i>

Note: When MSI-X capability is enabled, device-specific configuration starts at byte offset 24 in virtio common configuration structure. When MSI-X capability is not enabled, device-specific configuration starts at byte offset 20 in virtio header. ie. once you enable MSI-X on the device, the other fields move. If you turn it off again, they move back!

Any device-specific configuration space immediately follows these general headers:

Bits	Device Specific	
Read / Write	Device Specific	...
Purpose	Device Specific	

When accessing the device-specific configuration space using the legacy interface, transitional drivers MUST access the device-specific configuration space at an offset immediately following the general headers.

When using the legacy interface, transitional devices MUST present the device-specific configuration space if any at an offset immediately following the general headers.

Note that only Feature Bits 0 to 31 are accessible through the Legacy Interface. When used through the Legacy Interface, Transitional Devices MUST assume that Feature Bits 32 to 63 are not acknowledged by Driver.

As legacy devices had no *config_generation* field, see [2.3.4 Legacy Interface: Device Configuration Space](#) for workarounds.

4.1.4.9 Non-transitional Device With Legacy Driver: A Note on PCI Device Layout

All known legacy drivers check either the PCI Revision or the Device and Vendor IDs, and thus won't attempt to drive a non-transitional device.

A buggy legacy driver might mistakenly attempt to drive a non-transitional device. If support for such drivers is required (as opposed to fixing the bug), the following would be the recommended way to detect and handle them.

Note: Such buggy drivers are not currently known to be used in production.

4.1.4.9.0.1 ~~Driver Requirements: Non-transitional Device With Legacy Driver~~ Device Requirements: Non-transitional Device With Legacy Driver

Non-transitional devices, on a platform where a legacy driver for a legacy device with the same ID (including PCI Revision, Device and Vendor IDs) is known to have previously existed, SHOULD take the following steps to cause the legacy driver to fail gracefully when it attempts to drive them:

1. Present an I/O BAR in BAR0, and

2. Respond to a single-byte zero write to offset 18 (corresponding to Device Status register in the legacy layout) of BAR0 by presenting zeroes on every BAR and ignoring writes.

4.1.5 PCI-specific Initialization And Device Operation

4.1.5.1 Device Initialization

This documents PCI-specific steps executed during Device Initialization.

4.1.5.1.1 Virtio Device Configuration Layout Detection

As a prerequisite to device initialization, the driver scans the PCI capability list, detecting virtio configuration layout using Virtio Structure PCI capabilities as detailed in [4.1.4](#)

4.1.5.1.1.1 Legacy Interface: A Note on Device Layout Detection

Legacy drivers skipped the Device Layout Detection step, assuming legacy device configuration space in BAR0 in I/O space unconditionally.

Legacy devices did not have the Virtio PCI Capability in their capability list.

Therefore:

Transitional devices MUST expose the Legacy Interface in I/O space in BAR0.

Transitional drivers MUST look for the Virtio PCI Capabilities on the capability list. If these are not present, driver MUST assume a legacy device, and use it through the legacy interface.

Non-transitional drivers MUST look for the Virtio PCI Capabilities on the capability list. If these are not present, driver MUST assume a legacy device, and fail gracefully.

4.1.5.1.2 MSI-X Vector Configuration

When MSI-X capability is present and enabled in the device (through standard PCI configuration space) *config_msix_vector* and *queue_msix_vector* are used to map configuration change and queue interrupts to MSI-X vectors. In this case, the ISR Status is unused.

Writing a valid MSI-X Table entry number, 0 to 0x7FFF, to *config_msix_vector/queue_msix_vector* maps interrupts triggered by the configuration change/selected queue events respectively to the corresponding MSI-X vector. To disable interrupts for an event type, the driver unmaps this event by writing a special NO_VECTOR value:

```
/* Vector value used to disable MSI for queue */
#define VIRTIO_MSI_NO_VECTOR          0xffff
```

Note that mapping an event to vector might require device to allocate internal device resources, and thus could fail.

4.1.5.1.2.1 Device Requirements: MSI-X Vector Configuration

A device that has an MSI-X capability SHOULD support at least 2 and at most 0x800 MSI-X vectors. Device MUST report the number of vectors supported in *Table Size* in the MSI-X Capability as specified in [\[PCI\]](#). The device SHOULD restrict the reported MSI-X Table Size field to a value that might benefit system performance.

Note: For example, a device which does not expect to send interrupts at a high rate might only specify 2 MSI-X vectors.

Device **MUST** support mapping any event type to any valid vector 0 to MSI-X *Table Size*. Device **MUST** support unmapping any event type.

The device **MUST** return vector mapped to a given event, (NO_VECTOR if unmapped) on read of *config_msix_vector/queue_msix_vector*. The device **MUST** have all queue and configuration change events are unmapped upon reset.

Devices **SHOULD NOT** cause mapping an event to vector to fail unless it is impossible for the device to satisfy the mapping request. Devices **MUST** report mapping failures by returning the NO_VECTOR value when the relevant *config_msix_vector/queue_msix_vector* field is read.

4.1.5.1.2.2 Driver Requirements: MSI-X Vector Configuration

Driver **MUST** support device with any MSI-X Table Size 0 to 0x7FF. Driver **MAY** fall back on using INT#x interrupts for a device which only supports one MSI-X vector (MSI-X Table Size = 0).

Driver **MAY** interpret the Table Size as a hint from the device for the suggested number of MSI-X vectors to use.

Driver **MUST NOT** attempt to map an event to a vector outside the MSI-X Table supported by the device, as reported by *Table Size* in the MSI-X Capability.

After mapping an event to vector, the driver **MUST** verify success by reading the Vector field value: on success, the previously written value is returned, and on failure, NO_VECTOR is returned. If a mapping failure is detected, the driver **MAY** retry mapping with fewer vectors, disable MSI-X or report device failure.

4.1.5.1.3 Virtqueue Configuration

As a device can have zero or more virtqueues for bulk data transport¹, the driver needs to configure them as part of the device-specific configuration.

The driver typically does this as follows, for each virtqueue a device has:

1. Write the virtqueue index (first queue is 0) to *queue_select*.
2. Read the virtqueue size from *queue_size*. This controls how big the virtqueue is (see [2.4 Virtqueues](#)). If this field is 0, the virtqueue does not exist.
3. Optionally, select a smaller virtqueue size and write it to *queue_size*.
4. Allocate and zero Descriptor Table, Available and Used rings for the virtqueue in contiguous physical memory.
5. Optionally, if MSI-X capability is present and enabled on the device, select a vector to use to request interrupts triggered by virtqueue events. Write the MSI-X Table entry number corresponding to this vector into *queue_msix_vector*. Read *queue_msix_vector*: on success, previously written value is returned; on failure, NO_VECTOR value is returned.

4.1.5.1.3.1 Legacy Interface: A Note on Virtqueue Configuration

When using the legacy interface, the queue layout follows [2.5.2 Legacy Interfaces: A Note on Virtqueue Layout](#) with an alignment of 4096. Driver writes the physical address, divided by 4096 to the Queue Address field². There was no mechanism to negotiate the queue size.

¹For example, the simplest network device has two virtqueues.

²The 4096 is based on the x86 page size, but it's also large enough to ensure that the separate parts of the virtqueue are on separate cache lines.

4.1.5.2 Notifying The Device

~~When VIRTIO_F_NOTIFICATION_DATA has not been negotiated, the~~ The driver notifies the device by writing the 16-bit virtqueue index of this virtqueue ~~(in little-endian byte order format)~~ to the Queue Notify address.

~~When VIRTIO_F_NOTIFICATION_DATA has been negotiated, the driver notifies the device by writing the following 32-bit value to the Queue Notify address:-~~

~~See ?? for the definition of the components.~~

See 4.1.4.4 for how to calculate ~~the Queue Notify this~~ address.

4.1.5.3 Virtqueue Interrupts From The Device

If an interrupt is necessary for a virtqueue, the device would typically act as follows:

- If MSI-X capability is disabled:
 1. Set the lower bit of the ISR Status field for the device.
 2. Send the appropriate PCI interrupt for the device.
- If MSI-X capability is enabled:
 1. If *queue_msix_vector* is not NO_VECTOR, request the appropriate MSI-X interrupt message for the device, *queue_msix_vector* sets the MSI-X Table entry number.

4.1.5.3.1 Device Requirements: Virtqueue Interrupts From The Device

If MSI-X capability is enabled and *queue_msix_vector* is NO_VECTOR for a virtqueue, the device MUST NOT deliver an interrupt for that virtqueue.

4.1.5.4 Notification of Device Configuration Changes

Some virtio PCI devices can change the device configuration state, as reflected in the device-specific configuration region of the device. In this case:

- If MSI-X capability is disabled:
 1. Set the second lower bit of the ISR Status field for the device.
 2. Send the appropriate PCI interrupt for the device.
- If MSI-X capability is enabled:
 1. If *config_msix_vector* is not NO_VECTOR, request the appropriate MSI-X interrupt message for the device, *config_msix_vector* sets the MSI-X Table entry number.

A single interrupt MAY indicate both that one or more virtqueue has been used and that the configuration space has changed.

4.1.5.4.1 Device Requirements: Notification of Device Configuration Changes

If MSI-X capability is enabled and *config_msix_vector* is NO_VECTOR, the device MUST NOT deliver an interrupt for device configuration space changes.

4.1.5.4.2 Driver Requirements: Notification of Device Configuration Changes

A driver MUST handle the case where the same interrupt is used to indicate both device configuration space change and one or more virtqueues being used.

4.1.5.5 Driver Handling Interrupts

The driver interrupt handler would typically:

- If MSI-X capability is disabled:
 - Read the ISR Status field, which will reset it to zero.
 - If the lower bit is set: look through all virtqueues for the device, to see if any progress has been made by the device which requires servicing.
 - If the second lower bit is set: re-examine the configuration space to see what changed.
- If MSI-X capability is enabled:
 - Look through all virtqueues mapped to that MSI-X vector for the device, to see if any progress has been made by the device which requires servicing.
 - If the MSI-X vector is equal to *config_msix_vector*, re-examine the configuration space to see what changed.

4.2 Virtio Over MMIO

Virtual environments without PCI support (a common situation in embedded devices models) might use simple memory mapped device (“virtio-mmio”) instead of the PCI device.

The memory mapped virtio device behaviour is based on the PCI device specification. Therefore most operations including device initialization, queues configuration and buffer transfers are nearly identical. Existing differences are described in the following sections.

4.2.1 MMIO Device Discovery

Unlike PCI, MMIO provides no generic device discovery mechanism. For each device, the guest OS will need to know the location of the registers and interrupt(s) used. The suggested binding for systems using flattened device trees is shown in this example:

```
// EXAMPLE: virtio_block device taking 512 bytes at 0x1e000, interrupt 42.
virtio_block@1e000 {
    compatible = "virtio,mmio";
    reg = <0x1e000 0x200>;
    interrupts = <42>;
}
```

4.2.2 MMIO Device Register Layout

MMIO virtio devices provide a set of memory mapped control registers followed by a device-specific configuration space, described in the table 4.1.

All register values are organized as Little Endian.

Table 4.1: MMIO Device Register Layout

<i>Name</i>	Function
Offset from base	Description
Direction	
<i>MagicValue</i>	Magic value
0x000	0x74726976 (a Little Endian equivalent of the “virt” string).
R	

<i>Name</i> Offset from the base Direction	Function Description
<i>Version</i> 0x004 R	Device version number 0x2. Note: Legacy devices (see 4.2.4 Legacy interface) used 0x1.
<i>DeviceID</i> 0x008 R	Virtio Subsystem Device ID See 5 Device Types for possible values. Value zero (0x0) is used to define a system memory map with placeholder devices at static, well known addresses, assigning functions to them depending on user's needs.
<i>VendorID</i> 0x00c R	Virtio Subsystem Vendor ID
<i>DeviceFeatures</i> 0x010 R	Flags representing features the device supports Reading from this register returns 32 consecutive flag bits, the least significant bit depending on the last value written to <i>DeviceFeaturesSel</i> . Access to this register returns bits $DeviceFeaturesSel * 32$ to $(DeviceFeaturesSel * 32) + 31$, eg. feature bits 0 to 31 if <i>DeviceFeaturesSel</i> is set to 0 and features bits 32 to 63 if <i>DeviceFeaturesSel</i> is set to 1. Also see 2.2 Feature Bits .
<i>DeviceFeaturesSel</i> 0x014 W	Device (host) features word selection. Writing to this register selects a set of 32 device feature bits accessible by reading from <i>DeviceFeatures</i> .
<i>DriverFeatures</i> 0x020 W	Flags representing device features understood and activated by the driver Writing to this register sets 32 consecutive flag bits, the least significant bit depending on the last value written to <i>DriverFeaturesSel</i> . Access to this register sets bits $DriverFeaturesSel * 32$ to $(DriverFeaturesSel * 32) + 31$, eg. feature bits 0 to 31 if <i>DriverFeaturesSel</i> is set to 0 and features bits 32 to 63 if <i>DriverFeaturesSel</i> is set to 1. Also see 2.2 Feature Bits .
<i>DriverFeaturesSel</i> 0x024 W	Activated (guest) features word selection Writing to this register selects a set of 32 activated feature bits accessible by writing to <i>DriverFeatures</i> .
<i>QueueSel</i> 0x030 W	Virtual queue index Writing to this register selects the virtual queue that the following operations on <i>QueueNumMax</i> , <i>QueueNum</i> , <i>QueueReady</i> , <i>QueueDescLow</i> , <i>QueueDescHigh</i> , <i>QueueAvailLow</i> , <i>QueueAvailHigh</i> , <i>QueueUsedLow</i> and <i>QueueUsedHigh</i> apply to. The index number of the first queue is zero (0x0).
<i>QueueNumMax</i> 0x034 R	Maximum virtual queue size Reading from the register returns the maximum size (number of elements) of the queue the device is ready to process or zero (0x0) if the queue is not available. This applies to the queue selected by writing to <i>QueueSel</i> .
<i>QueueNum</i> 0x038 W	Virtual queue size Queue size is the number of elements in the queue. Writing to this register notifies the device what size of the queue the driver will use. This applies to the queue selected by writing to <i>QueueSel</i> .
<i>QueueReady</i> 0x044 RW	Virtual queue ready bit Writing one (0x1) to this register notifies the device that it can execute requests from this virtual queue. Reading from this register returns the last value written to it. Both read and write accesses apply to the queue selected by writing to <i>QueueSel</i> .

Name Offset from the base Direction	Function Description
<i>QueueNotify</i> 0x050 W	Queue notifier Writing a value queue index to this register notifies the device that there are new buffers to process in a queue . When VIRTIO_F_NOTIFICATION_DATA has not been negotiated, the value written is the queue index. When VIRTIO_F_NOTIFICATION_DATA has been negotiated, the value has the following format:- See ?? for the definition of the components-
<i>InterruptStatus</i> 0x60 R	Interrupt status Reading from this register returns a bit mask of events that caused the device interrupt to be asserted. The following events are possible: Used Buffer Update - bit 0 - the interrupt was asserted because the device has used a buffer in at least one of the active virtual queues. Configuration Change - bit 1 - the interrupt was asserted because the configuration of the device has changed.
<i>InterruptACK</i> 0x064 W	Interrupt acknowledge Writing a value with bits set as defined in <i>InterruptStatus</i> to this register notifies the device that events causing the interrupt have been handled.
<i>Status</i> 0x070 RW	Device status Reading from this register returns the current device status flags. Writing non-zero values to this register sets the status flags, indicating the driver progress. Writing zero (0x0) to this register triggers a device reset. See also p. 4.2.3.1 Device Initialization .
<i>QueueDescLow</i> 0x080 <i>QueueDescHigh</i> 0x084 W	Virtual queue's Descriptor Area 64 bit long physical address Writing to these two registers (lower 32 bits of the address to <i>QueueDescLow</i> , higher 32 bits to <i>QueueDescHigh</i>) notifies the device about location of the Descriptor Area of the queue selected by writing to <i>QueueSel</i> register.
<i>QueueDriverLow</i> 0x090 <i>QueueDriverHigh</i> 0x094 W	Virtual queue's Driver Area 64 bit long physical address Writing to these two registers (lower 32 bits of the address to <i>QueueAvailLow</i> , higher 32 bits to <i>QueueAvailHigh</i>) notifies the device about location of the Driver Area of the queue selected by writing to <i>QueueSel</i> .
<i>QueueDeviceLow</i> 0x0a0 <i>QueueDeviceHigh</i> 0x0a4 W	Virtual queue's Device Area 64 bit long physical address Writing to these two registers (lower 32 bits of the address to <i>QueueUsedLow</i> , higher 32 bits to <i>QueueUsedHigh</i>) notifies the device about location of the Device Area of the queue selected by writing to <i>QueueSel</i> .
<i>ConfigGeneration</i> 0x0fc R	Configuration atomicity value Reading from this register returns a value describing a version of the device-specific configuration space (see <i>Config</i>). The driver can then access the configuration space and, when finished, read <i>ConfigGeneration</i> again. If no part of the configuration space has changed between these two <i>ConfigGeneration</i> reads, the returned values are identical. If the values are different, the configuration space accesses were not atomic and the driver has to perform the operations again. See also 2.3.

<i>Name</i>	Function
Offset from the base Direction	Description
<i>Config</i> 0x100+ RW	Configuration space Device-specific configuration space starts at the offset 0x100 and is accessed with byte alignment. Its meaning and size depend on the device and the driver.

4.2.2.1 Device Requirements: MMIO Device Register Layout

The device MUST return 0x74726976 in *MagicValue*.

The device MUST return value 0x2 in *Version*.

The device MUST present each event by setting the corresponding bit in *InterruptStatus* from the moment it takes place, until the driver acknowledges the interrupt by writing a corresponding bit mask to the *InterruptACK* register. Bits which do not represent events which took place MUST be zero.

Upon reset, the device MUST clear all bits in *InterruptStatus* and ready bits in the *QueueReady* register for all queues in the device.

The device MUST change value returned in *ConfigGeneration* if there is any risk of a driver seeing an inconsistent configuration state.

The device MUST NOT access virtual queue contents when *QueueReady* is zero (0x0).

4.2.2.2 Driver Requirements: MMIO Device Register Layout

The driver MUST NOT access memory locations not described in the table 4.1 (or, in case of the configuration space, described in the device specification), MUST NOT write to the read-only registers (direction R) and MUST NOT read from the write-only registers (direction W).

The driver MUST only use 32 bit wide and aligned reads and writes to access the control registers described in table 4.1. For the device-specific configuration space, the driver MUST use 8 bit wide accesses for 8 bit wide fields, 16 bit wide and aligned accesses for 16 bit wide fields and 32 bit wide and aligned accesses for 32 and 64 bit wide fields.

The driver MUST ignore a device with *MagicValue* which is not 0x74726976, although it MAY report an error.

The driver MUST ignore a device with *Version* which is not 0x2, although it MAY report an error.

The driver MUST ignore a device with *DeviceID* 0x0, but MUST NOT report any error.

Before reading from *DeviceFeatures*, the driver MUST write a value to *DeviceFeaturesSel*.

Before writing to the *DriverFeatures* register, the driver MUST write a value to the *DriverFeaturesSel* register.

The driver MUST write a value to *QueueNum* which is less than or equal to the value presented by the device in *QueueNumMax*.

When *QueueReady* is not zero, the driver MUST NOT access *QueueNum*, *QueueDescLow*, *QueueDescHigh*, *QueueAvailLow*, *QueueAvailHigh*, *QueueUsedLow*, *QueueUsedHigh*.

To stop using the queue the driver MUST write zero (0x0) to this *QueueReady* and MUST read the value back to ensure synchronization.

The driver MUST ignore undefined bits in *InterruptStatus*.

The driver MUST write a value with a bit mask describing events it handled into *InterruptACK* when it finishes handling an interrupt and MUST NOT set any of the undefined bits in the value.

4.2.3 MMIO-specific Initialization And Device Operation

4.2.3.1 Device Initialization

4.2.3.1.1 Driver Requirements: Device Initialization

The driver MUST start the device initialization by reading and checking values from *MagicValue* and *Version*. If both values are valid, it MUST read *DeviceID* and if its value is zero (0x0) MUST abort initialization and MUST NOT access any other register.

Further initialization MUST follow the procedure described in [3.1 Device Initialization](#).

4.2.3.2 Virtqueue Configuration

The driver will typically initialize the virtual queue in the following way:

1. Select the queue writing its index (first queue is 0) to *QueueSel*.
2. Check if the queue is not already in use: read *QueueReady*, and expect a returned value of zero (0x0).
3. Read maximum queue size (number of elements) from *QueueNumMax*. If the returned value is zero (0x0) the queue is not available.
4. Allocate and zero the queue memory, making sure the memory is physically contiguous.
5. Notify the device about the queue size by writing the size to *QueueNum*.
6. Write physical addresses of the queue's Descriptor Area, Driver Area and Device Area to (respectively) the *QueueDescLow/QueueDescHigh*, *QueueDriverLow/QueueDriverHigh* and *QueueDeviceLow/QueueDeviceHigh* register pairs.
7. Write 0x1 to *QueueReady*.

4.2.3.3 Notifying The Device

The driver notifies the device about new buffers being available in a queue by writing the index of the updated queue to *QueueNotify*.

4.2.3.4 Notifications From The Device

The memory mapped virtio device is using a single, dedicated interrupt signal, which is asserted when at least one of the bits described in the description of *InterruptStatus* is set. This is how the device notifies the driver about a new used buffer being available in the queue or about a change in the device configuration.

4.2.3.4.1 Driver Requirements: Notifications From The Device

After receiving an interrupt, the driver MUST read *InterruptStatus* to check what caused the interrupt (see the register description). After the interrupt is handled, the driver MUST acknowledge it by writing a bit mask corresponding to the handled events to the InterruptACK register.

4.2.4 Legacy interface

The legacy MMIO transport used page-based addressing, resulting in a slightly different control register layout, the device initialization and the virtual queue configuration procedure.

Table [4.2](#) presents control registers layout, omitting descriptions of registers which did not change their function nor behaviour:

Table 4.2: MMIO Device Legacy Register Layout

<i>Name</i> Offset from base Direction	Function Description
<i>MagicValue</i> 0x000 R	Magic value
<i>Version</i> 0x004 R	Device version number Legacy device returns value 0x1.
<i>DeviceID</i> 0x008 R	Virtio Subsystem Device ID
<i>VendorID</i> 0x00c R	Virtio Subsystem Vendor ID
<i>HostFeatures</i> 0x010 R	Flags representing features the device supports
<i>HostFeaturesSel</i> 0x014 W	Device (host) features word selection.
<i>GuestFeatures</i> 0x020 W	Flags representing device features understood and activated by the driver
<i>GuestFeaturesSel</i> 0x024 W	Activated (guest) features word selection
<i>GuestPageSize</i> 0x028 W	Guest page size The driver writes the guest page size in bytes to the register during initialization, before any queues are used. This value should be a power of 2 and is used by the device to calculate the Guest address of the first queue page (see <i>QueuePFN</i>).
<i>QueueSel</i> 0x030 W	Virtual queue index Writing to this register selects the virtual queue that the following operations on the <i>QueueNumMax</i> , <i>QueueNum</i> , <i>QueueAlign</i> and <i>QueuePFN</i> registers apply to. The index number of the first queue is zero (0x0).
<i>QueueNumMax</i> 0x034 R	Maximum virtual queue size Reading from the register returns the maximum size of the queue the device is ready to process or zero (0x0) if the queue is not available. This applies to the queue selected by writing to <i>QueueSel</i> and is allowed only when <i>QueuePFN</i> is set to zero (0x0), so when the queue is not actively used.
<i>QueueNum</i> 0x038 W	Virtual queue size Queue size is the number of elements in the queue. Writing to this register notifies the device what size of the queue the driver will use. This applies to the queue selected by writing to <i>QueueSel</i> .
<i>QueueAlign</i> 0x03c W	Used Ring alignment in the virtual queue Writing to this register notifies the device about alignment boundary of the Used Ring in bytes. This value should be a power of 2 and applies to the queue selected by writing to <i>QueueSel</i> .

<i>Name</i> Offset from the base Direction	Function Description
<i>QueuePFN</i> 0x040 RW	Guest physical page number of the virtual queue Writing to this register notifies the device about location of the virtual queue in the Guest's physical address space. This value is the index number of a page starting with the queue Descriptor Table. Value zero (0x0) means physical address zero (0x00000000) and is illegal. When the driver stops using the queue it writes zero (0x0) to this register. Reading from this register returns the currently used page number of the queue, therefore a value other than zero (0x0) means that the queue is in use. Both read and write accesses apply to the queue selected by writing to <i>QueueSel</i> .
<i>QueueNotify</i> 0x050 W	Queue notifier
<i>InterruptStatus</i> 0x60 R	Interrupt status
<i>InterruptACK</i> 0x064 W	Interrupt acknowledge
<i>Status</i> 0x070 RW	Device status Reading from this register returns the current device status flags. Writing non-zero values to this register sets the status flags, indicating the OS/driver progress. Writing zero (0x0) to this register triggers a device reset. The device sets <i>QueuePFN</i> to zero (0x0) for all queues in the device. Also see 3.1 Device Initialization .
<i>Config</i> 0x100+ RW	Configuration space

The virtual queue page size is defined by writing to *GuestPageSize*, as written by the guest. The driver does this before the virtual queues are configured.

The virtual queue layout follows p. [2.5.2 Legacy Interfaces: A Note on Virtqueue Layout](#), with the alignment defined in *QueueAlign*.

The virtual queue is configured as follows:

1. Select the queue writing its index (first queue is 0) to *QueueSel*.
2. Check if the queue is not already in use: read *QueuePFN*, expecting a returned value of zero (0x0).
3. Read maximum queue size (number of elements) from *QueueNumMax*. If the returned value is zero (0x0) the queue is not available.
4. Allocate and zero the queue pages in contiguous virtual memory, aligning the Used Ring to an optimal boundary (usually page size). The driver should choose a queue size smaller than or equal to *QueueNumMax*.
5. Notify the device about the queue size by writing the size to *QueueNum*.
6. Notify the device about the used alignment by writing its value in bytes to *QueueAlign*.
7. Write the physical number of the first page of the queue to the *QueuePFN* register.

Notification mechanisms did not change.

4.3 Virtio Over Channel I/O

S/390 based virtual machines support neither PCI nor MMIO, so a different transport is needed there.

virtio-ccw uses the standard channel I/O based mechanism used for the majority of devices on S/390. A virtual channel device with a special control unit type acts as proxy to the virtio device (similar to the way virtio-pci uses a PCI device) and configuration and operation of the virtio device is accomplished (mostly) via channel commands. This means virtio devices are discoverable via standard operating system algorithms, and adding virtio support is mainly a question of supporting a new control unit type.

As the S/390 is a big endian machine, the data structures transmitted via channel commands are big-endian: this is made clear by use of the types be16, be32 and be64.

4.3.1 Basic Concepts

As a proxy device, virtio-ccw uses a channel-attached I/O control unit with a special control unit type (0x3832) and a control unit model corresponding to the attached virtio device's subsystem device ID, accessed via a virtual I/O subchannel and a virtual channel path of type 0x32. This proxy device is discoverable via normal channel subsystem device discovery (usually a STORE SUBCHANNEL loop) and answers to the basic channel commands:

- NO-OPERATION (0x03)
- BASIC SENSE (0x04)
- TRANSFER IN CHANNEL (0x08)
- SENSE ID (0xe4)

For a virtio-ccw proxy device, SENSE ID will return the following information:

Bytes	Description	Contents
0	reserved	0xff
1-2	control unit type	0x3832
3	control unit model	<virtio device id>
4-5	device type	zeroes (unset)
6	device model	zeroes (unset)
7-255	extended Senseld data	zeroes (unset)

In addition to the basic channel commands, virtio-ccw defines a set of channel commands related to configuration and operation of virtio:

```
#define CCW_CMD_SET_VQ 0x13
#define CCW_CMD_VDEV_RESET 0x33
#define CCW_CMD_SET_IND 0x43
#define CCW_CMD_SET_CONF_IND 0x53
#define CCW_CMD_SET_IND_ADAPTER 0x73
#define CCW_CMD_READ_FEAT 0x12
#define CCW_CMD_WRITE_FEAT 0x11
#define CCW_CMD_READ_CONF 0x22
#define CCW_CMD_WRITE_CONF 0x21
#define CCW_CMD_WRITE_STATUS 0x31
#define CCW_CMD_READ_VQ_CONF 0x32
#define CCW_CMD_SET_VIRTIO_REV 0x83
#define CCW_CMD_READ_STATUS 0x72
```

4.3.1.1 Device Requirements: Basic Concepts

The virtio-ccw device acts like a normal channel device, as specified in [S390 PoP] and [S390 Common I/O]. In particular:

- A device **MUST** post a unit check with command reject for any command it does not support.
- If a driver did not suppress length checks for a channel command, the device **MUST** present a sub-channel status as detailed in the architecture when the actual length did not match the expected length.
- If a driver did suppress length checks for a channel command, the device **MUST** present a check condition if the transmitted data does not contain enough data to process the command. If the driver submitted a buffer that was too long, the device **SHOULD** accept the command.

4.3.1.2 Driver Requirements: Basic Concepts

A driver for virtio-ccw devices **MUST** check for a control unit type of 0x3832 and **MUST** ignore the device type and model.

A driver **SHOULD** attempt to provide the correct length in a channel command even if it suppresses length checks for that command.

4.3.2 Device Initialization

virtio-ccw uses several channel commands to set up a device.

4.3.2.1 Setting the Virtio Revision

CCW_CMD_SET_VIRTIO_REV is issued by the driver to set the revision of the virtio-ccw transport it intends to drive the device with. It uses the following communication structure:

```
struct virtio_rev_info {  
    be16 revision;  
    be16 length;  
    u8 data[];  
};
```

revision contains the desired revision id, *length* the length of the data portion and *data* revision-dependent additional desired options.

The following values are supported:

<i>revision</i>	<i>length</i>	<i>data</i>	remarks
0	0	<empty>	legacy interface; transitional devices only
1	0	<empty>	Virtio 1.0
2	0	<empty>	CCW_CMD_READ_STATUS support
3-n			reserved for later revisions

Note that a change in the virtio standard does not necessarily correspond to a change in the virtio-ccw revision.

4.3.2.1.1 Device Requirements: Setting the Virtio Revision

A device **MUST** post a unit check with command reject for any *revision* it does not support. For any invalid combination of *revision*, *length* and *data*, it **MUST** post a unit check with command reject as well. A non-transitional device **MUST** reject revision id 0.

A device **MUST** answer with command reject to any virtio-ccw specific channel command that is not contained in the revision selected by the driver.

A device **MUST** answer with command reject to any attempt to select a different revision after a revision has been successfully selected by the driver.

A device **MUST** treat the revision as unset from the time the associated subchannel has been enabled until a revision has been successfully set by the driver. This implies that revisions are not persistent across disabling and enabling of the associated subchannel.

4.3.2.1.2 Driver Requirements: Setting the Virtio Revision

A driver **SHOULD** start with trying to set the highest revision it supports and continue with lower revisions if it gets a command reject.

A driver **MUST NOT** issue any other virtio-ccw specific channel commands prior to setting the revision.

After a revision has been successfully selected by the driver, it **MUST NOT** attempt to select a different revision.

4.3.2.1.3 Legacy Interfaces: A Note on Setting the Virtio Revision

A legacy device will not support the CCW_CMD_SET_VIRTIO_REV and answer with a command reject. A non-transitional driver **MUST** stop trying to operate this device in that case. A transitional driver **MUST** operate the device as if it had been able to set revision 0.

A legacy driver will not issue the CCW_CMD_SET_VIRTIO_REV prior to issuing other virtio-ccw specific channel commands. A non-transitional device therefore **MUST** answer any such attempts with a command reject. A transitional device **MUST** assume in this case that the driver is a legacy driver and continue as if the driver selected revision 0. This implies that the device **MUST** reject any command not valid for revision 0, including a subsequent CCW_CMD_SET_VIRTIO_REV.

4.3.2.2 Configuring a Virtqueue

CCW_CMD_READ_VQ_CONF is issued by the driver to obtain information about a queue. It uses the following structure for communicating:

```
struct vq_config_block {
    be16 index;
    be16 max_num;
};
```

The requested number of buffers for queue *index* is returned in *max_num*.

Afterwards, CCW_CMD_SET_VQ is issued by the driver to inform the device about the location used for its queue. The transmitted structure is

```
struct vq_info_block {
    be64 desc;
    be32 res0;
    be16 index;
    be16 num;
    be64 driver;
    be64 device;
};
```

desc, *driver* and *device* contain the guest addresses for the descriptor area, available area and used area for queue *index*, respectively. The actual virtqueue size (number of allocated buffers) is transmitted in *num*.

4.3.2.2.1 Device Requirements: Configuring a Virtqueue

res0 is reserved and MUST be ignored by the device.

4.3.2.2.2 Legacy Interface: A Note on Configuring a Virtqueue

For a legacy driver or for a driver that selected revision 0, CCW_CMD_SET_VQ uses the following communication block:

```
struct vq_info_block_legacy {
    be64 queue;
    be32 align;
    be16 index;
    be16 num;
};
```

queue contains the guest address for queue *index*, *num* the number of buffers and *align* the alignment. The queue layout follows [2.5.2 Legacy Interfaces: A Note on Virtqueue Layout](#).

4.3.2.3 Communicating Status Information

The driver changes the status of a device via the CCW_CMD_WRITE_STATUS command, which transmits an 8 bit status value.

As described in [2.2.2](#), a device sometimes fails to set the *status* field: For example, it might fail to accept the FEATURES_OK status bit during device initialization.

With revision 2, CCW_CMD_READ_STATUS is defined: It reads an 8 bit status value from the device and acts as a reverse operation to CCW_CMD_WRITE_STATUS.

4.3.2.3.1 Driver Requirements: Communicating Status Information

If the device posts a unit check with command reject in response to the CCW_CMD_WRITE_STATUS command, the driver MUST assume that the device failed to set the status and the *status* field retained its previous value.

If at least revision 2 has been negotiated, the driver SHOULD use the CCW_CMD_READ_STATUS command to retrieve the *status* field after a configuration change has been detected.

If not at least revision 2 has been negotiated, the driver MUST NOT attempt to issue the CCW_CMD_READ_STATUS command.

4.3.2.3.2 Device Requirements: Communicating Status Information

If the device fails to set the *status* field to the value written by the driver, the device MUST assure that the *status* field is left unchanged and MUST post a unit check with command reject.

If at least revision 2 has been negotiated, the device MUST return the current *status* field if the CCW_CMD_READ_STATUS command is issued.

4.3.2.4 Handling Device Features

Feature bits are arranged in an array of 32 bit values, making for a total of 8192 feature bits. Feature bits are in little-endian byte order.

The CCW commands dealing with features use the following communication block:

```
struct virtio_feature_desc {
    le32 features;
    u8 index;
};
```

features are the 32 bits of features currently accessed, while *index* describes which of the feature bit values is to be accessed. No padding is added at the end of the structure, it is exactly 5 bytes in length.

The guest obtains the device's device feature set via the CCW_CMD_READ_FEAT command. The device stores the features at *index* to *features*.

For communicating its supported features to the device, the driver uses the CCW_CMD_WRITE_FEAT command, denoting a *features/index* combination.

4.3.2.5 Device Configuration

The device's configuration space is located in host memory.

To obtain information from the configuration space, the driver uses CCW_CMD_READ_CONF, specifying the guest memory for the device to write to.

For changing configuration information, the driver uses CCW_CMD_WRITE_CONF, specifying the guest memory for the device to read from.

In both cases, the complete configuration space is transmitted. This allows the driver to compare the new configuration space with the old version, and keep a generation count internally whenever it changes.

4.3.2.6 Setting Up Indicators

In order to set up the indicator bits for host->guest notification, the driver uses different channel commands depending on whether it wishes to use traditional I/O interrupts tied to a subchannel or adapter I/O interrupts for virtqueue notifications. For any given device, the two mechanisms are mutually exclusive.

For the configuration change indicators, only a mechanism using traditional I/O interrupts is provided, regardless of whether traditional or adapter I/O interrupts are used for virtqueue notifications.

4.3.2.6.1 Setting Up Classic Queue Indicators

Indicators for notification via classic I/O interrupts are contained in a 64 bit value per virtio-ccw proxy device.

To communicate the location of the indicator bits for host->guest notification, the driver uses the CCW_CMD_SET_IND command, pointing to a location containing the guest address of the indicators in a 64 bit value.

If the driver has already set up two-staged queue indicators via the CCW_CMD_SET_IND_ADAPTER command, the device MUST post a unit check with command reject to any subsequent CCW_CMD_SET_IND command.

4.3.2.6.2 Setting Up Configuration Change Indicators

Indicators for configuration change host->guest notification are contained in a 64 bit value per virtio-ccw proxy device.

To communicate the location of the indicator bits used in the configuration change host->guest notification, the driver issues the CCW_CMD_SET_CONF_IND command, pointing to a location containing the guest address of the indicators in a 64 bit value.

4.3.2.6.3 Setting Up Two-Stage Queue Indicators

Indicators for notification via adapter I/O interrupts consist of two stages:

- a summary indicator byte covering the virtqueues for one or more virtio-ccw proxy devices
- a set of contiguous indicator bits for the virtqueues for a virtio-ccw proxy device

To communicate the location of the summary and queue indicator bits, the driver uses the `CCW_CMD_SET_IND_ADAPTER` command with the following payload:

```
struct virtio_thinint_area {  
    be64 summary_indicator;  
    be64 indicator;  
    be64 bit_nr;  
    u8 isc;  
} __attribute__((packed));
```

summary_indicator contains the guest address of the 8 bit summary indicator. *indicator* contains the guest address of an area wherein the indicators for the devices are contained, starting at *bit_nr*, one bit per virtqueue of the device. Bit numbers start at the left, i.e. the most significant bit in the first byte is assigned the bit number 0. *isc* contains the I/O interruption subclass to be used for the adapter I/O interrupt. It MAY be different from the *isc* used by the proxy virtio-ccw device's subchannel. No padding is added at the end of the structure, it is exactly 25 bytes in length.

4.3.2.6.3.1 Device Requirements: Setting Up Two-Stage Queue Indicators

If the driver has already set up classic queue indicators via the `CCW_CMD_SET_IND` command, the device MUST post a unit check with command reject to any subsequent `CCW_CMD_SET_IND_ADAPTER` command.

4.3.2.6.4 Legacy Interfaces: A Note on Setting Up Indicators

In some cases, legacy devices will only support classic queue indicators; in that case, they will reject `CCW_CMD_SET_IND_ADAPTER` as they don't know that command. Some legacy devices will support two-stage queue indicators, though, and a driver will be able to successfully use `CCW_CMD_SET_IND_ADAPTER` to set them up.

4.3.3 Device Operation

4.3.3.1 Host->Guest Notification

There are two modes of operation regarding host->guest notification, classic I/O interrupts and adapter I/O interrupts. The mode to be used is determined by the driver by using `CCW_CMD_SET_IND` respectively `CCW_CMD_SET_IND_ADAPTER` to set up queue indicators.

For configuration changes, the driver always uses classic I/O interrupts.

4.3.3.1.1 Notification via Classic I/O Interrupts

If the driver used the `CCW_CMD_SET_IND` command to set up queue indicators, the device will use classic I/O interrupts for host->guest notification about virtqueue activity.

For notifying the driver of virtqueue buffers, the device sets the corresponding bit in the guest-provided indicators. If an interrupt is not already pending for the subchannel, the device generates an unsolicited I/O interrupt.

If the device wants to notify the driver about configuration changes, it sets bit 0 in the configuration indicators and generates an unsolicited I/O interrupt, if needed. This also applies if adapter I/O interrupts are used for queue notifications.

4.3.3.1.2 Notification via Adapter I/O Interrupts

If the driver used the CCW_CMD_SET_IND_ADAPTER command to set up queue indicators, the device will use adapter I/O interrupts for host->guest notification about virtqueue activity.

For notifying the driver of virtqueue buffers, the device sets the bit in the guest-provided indicator area at the corresponding offset. The guest-provided summary indicator is set to 0x01. An adapter I/O interrupt for the corresponding interruption subclass is generated.

The recommended way to process an adapter I/O interrupt by the driver is as follows:

- Process all queue indicator bits associated with the summary indicator.
- Clear the summary indicator, performing a synchronization (memory barrier) afterwards.
- Process all queue indicator bits associated with the summary indicator again.

4.3.3.1.2.1 Device Requirements: Notification via Adapter I/O Interrupts

The device SHOULD only generate an adapter I/O interrupt if the summary indicator had not been set prior to notification.

4.3.3.1.2.2 Driver Requirements: Notification via Adapter I/O Interrupts

The driver MUST clear the summary indicator after receiving an adapter I/O interrupt before it processes the queue indicators.

4.3.3.1.3 Legacy Interfaces: A Note on Host->Guest Notification

As legacy devices and drivers support only classic queue indicators, host->guest notification will always be done via classic I/O interrupts.

4.3.3.2 Guest->Host Notification

For notifying the device of virtqueue buffers, the driver unfortunately can't use a channel command (the asynchronous characteristics of channel I/O interact badly with the host block I/O backend). Instead, it uses a diagnose 0x500 call with subcode 3 specifying the queue, as follows:

GPR	Input Value	Output Value
1	0x3	
2	Subchannel ID	Host Cookie
3	Notification data -Virtqueue number	
4	Host Cookie	

~~When VIRTIO_F_NOTIFICATION_DATA has not been negotiated, Notification data includes the Virtqueue number.~~

~~When VIRTIO_F_NOTIFICATION_DATA has been negotiated, Notification data has the following format:~~

~~See ?? for the definition of the components.~~

4.3.3.2.1 Device Requirements: Guest->Host Notification

The device MUST ignore bits 0-31 (counting from the left) of GPR2. This aligns passing the subchannel ID with the way it is passed for the existing I/O instructions.

The device MAY return a 64-bit host cookie in GPR2 to speed up the notification execution.

4.3.3.2.2 Driver Requirements: Guest->Host Notification

For each notification, the driver SHOULD use GPR4 to pass the host cookie received in GPR2 from the previous notification.

Note: For example:

```
info->cookie = do_notify(schid,
                        virtqueue_get_queue_index(vq),
                        info->cookie);
```

4.3.3.3 Resetting Devices

In order to reset a device, a driver sends the CCW_CMD_VDEV_RESET command.

5 Device Types

On top of the queues, config space and feature negotiation facilities built into virtio, several devices are defined.

The following device IDs are used to identify different types of virtio devices. Some device IDs are reserved for devices which are not currently defined in this standard.

Discovering what devices are available and their type is bus-dependent.

Device ID	Virtio Device
0	reserved (invalid)
1	network card
2	block device
3	console
4	entropy source
5	memory ballooning (traditional)
6	ioMemory
7	rpmsg
8	SCSI host
9	9P transport
10	mac80211 wlan
11	rproc serial
12	virtio CAIF
13	memory balloon
16	GPU device
17	Timer/Clock device
18	Input device
19	Socket device
20	Crypto device
21	Signal Distribution Module
22	pstore device
23	IOMMU device

Some of the devices above are unspecified by this document, because they are seen as immature or especially niche. Be warned that some are only specified by the sole existing implementation; they could become part of a future specification, be abandoned entirely, or live on outside this standard. We shall speak of them no further.

5.1 Network Device

The virtio network device is a virtual ethernet card, and is the most complex of the devices supported so far by virtio. It has enhanced rapidly and demonstrates clearly how support for new features are added to an

existing device. Empty buffers are placed in one virtqueue for receiving packets, and outgoing packets are enqueued into another for transmission in that order. A third command queue is used to control advanced filtering features.

5.1.1 Device ID

1

5.1.2 Virtqueues

0 receiveq1

1 transmitq1

...

2(N-1) receiveqN

2(N-1)+1 transmitqN

2N controlq

N=1 if VIRTIO_NET_F_MQ is not negotiated, otherwise N is set by *max_virtqueue_pairs*.

controlq only exists if VIRTIO_NET_F_CTRL_VQ set.

5.1.3 Feature bits

VIRTIO_NET_F_CSUM (0) Device handles packets with partial checksum. This “checksum offload” is a common feature on modern network cards.

VIRTIO_NET_F_GUEST_CSUM (1) Driver handles packets with partial checksum.

VIRTIO_NET_F_CTRL_GUEST_OFFLOADS (2) Control channel offloads reconfiguration support.

VIRTIO_NET_F_MTU(3) Device maximum MTU reporting is supported. If offered by the device, device advises driver about the value of its maximum MTU. If negotiated, the driver uses *mtu* as the maximum MTU value.

VIRTIO_NET_F_MAC (5) Device has given MAC address.

VIRTIO_NET_F_GUEST_TSO4 (7) Driver can receive TSOv4.

VIRTIO_NET_F_GUEST_TSO6 (8) Driver can receive TSOv6.

VIRTIO_NET_F_GUEST_ECN (9) Driver can receive TSO with ECN.

VIRTIO_NET_F_GUEST_UFO (10) Driver can receive UFO.

VIRTIO_NET_F_HOST_TSO4 (11) Device can receive TSOv4.

VIRTIO_NET_F_HOST_TSO6 (12) Device can receive TSOv6.

VIRTIO_NET_F_HOST_ECN (13) Device can receive TSO with ECN.

VIRTIO_NET_F_HOST_UFO (14) Device can receive UFO.

VIRTIO_NET_F_MRG_RXBUF (15) Driver can merge receive buffers.

VIRTIO_NET_F_STATUS (16) Configuration status field is available.

VIRTIO_NET_F_CTRL_VQ (17) Control channel is available.

VIRTIO_NET_F_CTRL_RX (18) Control channel RX mode support.

VIRTIO_NET_F_CTRL_VLAN (19) Control channel VLAN filtering.

VIRTIO_NET_F_GUEST_ANNOUNCE(21) Driver can send gratuitous packets.

VIRTIO_NET_F_MQ(22) Device supports multiqueue with automatic receive steering.

VIRTIO_NET_F_CTRL_MAC_ADDR(23) Set MAC address through control channel.

5.1.3.1 Feature bit requirements

Some networking feature bits require other networking feature bits (see 2.2.1):

VIRTIO_NET_F_GUEST_TSO4 Requires **VIRTIO_NET_F_GUEST_CSUM**.

VIRTIO_NET_F_GUEST_TSO6 Requires **VIRTIO_NET_F_GUEST_CSUM**.

VIRTIO_NET_F_GUEST_ECN Requires **VIRTIO_NET_F_GUEST_TSO4** or **VIRTIO_NET_F_GUEST_TSO6**.

VIRTIO_NET_F_GUEST_UFO Requires **VIRTIO_NET_F_GUEST_CSUM**.

VIRTIO_NET_F_HOST_TSO4 Requires **VIRTIO_NET_F_CSUM**.

VIRTIO_NET_F_HOST_TSO6 Requires **VIRTIO_NET_F_CSUM**.

VIRTIO_NET_F_HOST_ECN Requires **VIRTIO_NET_F_HOST_TSO4** or **VIRTIO_NET_F_HOST_TSO6**.

VIRTIO_NET_F_HOST_UFO Requires **VIRTIO_NET_F_CSUM**.

VIRTIO_NET_F_CTRL_RX Requires **VIRTIO_NET_F_CTRL_VQ**.

VIRTIO_NET_F_CTRL_VLAN Requires **VIRTIO_NET_F_CTRL_VQ**.

VIRTIO_NET_F_GUEST_ANNOUNCE Requires **VIRTIO_NET_F_CTRL_VQ**.

VIRTIO_NET_F_MQ Requires **VIRTIO_NET_F_CTRL_VQ**.

VIRTIO_NET_F_CTRL_MAC_ADDR Requires **VIRTIO_NET_F_CTRL_VQ**.

5.1.3.2 Legacy Interface: Feature bits

VIRTIO_NET_F_GSO (6) Device handles packets with any GSO type.

This was supposed to indicate segmentation offload support, but upon further investigation it became clear that multiple bits were needed.

5.1.4 Device configuration layout

Three driver-read-only configuration fields are currently defined. The *mac* address field always exists (though is only valid if **VIRTIO_NET_F_MAC** is set), and *status* only exists if **VIRTIO_NET_F_STATUS** is set. Two read-only bits (for the driver) are currently defined for the status field: **VIRTIO_NET_S_LINK_UP** and **VIRTIO_NET_S_ANNOUNCE**.

```
#define VIRTIO_NET_S_LINK_UP      1
#define VIRTIO_NET_S_ANNOUNCE    2
```

The following driver-read-only field, *max_virtqueue_pairs* only exists if **VIRTIO_NET_F_MQ** is set. This field specifies the maximum number of each of transmit and receive virtqueues (receiveq1...receiveqN and transmitq1...transmitqN respectively) that can be configured once **VIRTIO_NET_F_MQ** is negotiated.

The following driver-read-only field, *mtu* only exists if **VIRTIO_NET_F_MTU** is set. This field specifies the maximum MTU for the driver to use.

```
struct virtio_net_config {
    u8 mac[6];
    le16 status;
    le16 max_virtqueue_pairs;
    le16 mtu;
};
```


5.1.4.1 Device Requirements: Device configuration layout

The device MUST set *max_virtqueue_pairs* to between 1 and 0x8000 inclusive, if it offers VIRTIO_NET_F_MQ.

The device MUST set *mtu* to between 68 and 65535 inclusive, if it offers VIRTIO_NET_F_MTU.

The device SHOULD set *mtu* to at least 1280, if it offers VIRTIO_NET_F_MTU.

The device MUST NOT modify *mtu* once it has been set.

The device MUST NOT pass received packets that exceed *mtu* (plus low level ethernet header length) size with *gso_type* NONE or ECN after VIRTIO_NET_F_MTU has been successfully negotiated.

The device MUST forward transmitted packets of up to *mtu* (plus low level ethernet header length) size with *gso_type* NONE or ECN, and do so without fragmentation, after VIRTIO_NET_F_MTU has been successfully negotiated.

5.1.4.2 Driver Requirements: Device configuration layout

A driver SHOULD negotiate VIRTIO_NET_F_MAC if the device offers it. If the driver negotiates the VIRTIO_NET_F_MAC feature, the driver MUST set the physical address of the NIC to *mac*. Otherwise, it SHOULD use a locally-administered MAC address (see [IEEE 802](#), “9.2 48-bit universal LAN MAC addresses”).

If the driver does not negotiate the VIRTIO_NET_F_STATUS feature, it SHOULD assume the link is active, otherwise it SHOULD read the link status from the bottom bit of *status*.

A driver SHOULD negotiate VIRTIO_NET_F_MTU if the device offers it.

If the driver negotiates VIRTIO_NET_F_MTU, it MUST supply enough receive buffers to receive at least one receive packet of size *mtu* (plus low level ethernet header length) with *gso_type* NONE or ECN.

If the driver negotiates VIRTIO_NET_F_MTU, it MUST NOT transmit packets of size exceeding the value of *mtu* (plus low level ethernet header length) with *gso_type* NONE or ECN.

5.1.4.3 Legacy Interface: Device configuration layout

When using the legacy interface, transitional devices and drivers MUST format *status* and *max_virtqueue_pairs* in struct *virtio_net_config* according to the native endian of the guest rather than (necessarily when not using the legacy interface) little-endian.

When using the legacy interface, *mac* is driver-writable which provided a way for drivers to update the MAC without negotiating VIRTIO_NET_F_CTRL_MAC_ADDR.

5.1.5 Device Initialization

A driver would perform a typical initialization routine like so:

1. Identify and initialize the receive and transmission virtqueues, up to N of each kind. If VIRTIO_NET_F_MQ feature bit is negotiated, $N = \text{max_virtqueue_pairs}$, otherwise identify $N=1$.
2. If the VIRTIO_NET_F_CTRL_VQ feature bit is negotiated, identify the control virtqueue.
3. Fill the receive queues with buffers: see [5.1.6.3](#).
4. Even with VIRTIO_NET_F_MQ, only *receiveq1*, *transmitq1* and *controlq* are used by default. The driver would send the VIRTIO_NET_CTRL_MQ_VQ_PAIRS_SET command specifying the number of the transmit and receive queues to use.
5. If the VIRTIO_NET_F_MAC feature bit is set, the configuration space *mac* entry indicates the “physical” address of the network card, otherwise the driver would typically generate a random local MAC address.

6. If the VIRTIO_NET_F_STATUS feature bit is negotiated, the link status comes from the bottom bit of *status*. Otherwise, the driver assumes it's active.
7. A performant driver would indicate that it will generate checksumless packets by negotiating the VIRTIO_NET_F_CSUM feature.
8. If that feature is negotiated, a driver can use TCP or UDP segmentation offload by negotiating the VIRTIO_NET_F_HOST_TSO4 (IPv4 TCP), VIRTIO_NET_F_HOST_TSO6 (IPv6 TCP) and VIRTIO_NET_F_HOST_UFO (UDP fragmentation) features.
9. The converse features are also available: a driver can save the virtual device some work by negotiating these features.

Note: For example, a network packet transported between two guests on the same system might not need checksumming at all, nor segmentation, if both guests are amenable. The VIRTIO_NET_F_GUEST_CSUM feature indicates that partially checksummed packets can be received, and if it can do that then the VIRTIO_NET_F_GUEST_TSO4, VIRTIO_NET_F_GUEST_TSO6, VIRTIO_NET_F_GUEST_UFO and VIRTIO_NET_F_GUEST_ECN are the input equivalents of the features described above. See [5.1.6.3 Setting Up Receive Buffers](#) and [5.1.6.4 Processing of Incoming Packets](#) below.

A truly minimal driver would only accept VIRTIO_NET_F_MAC and ignore everything else.

5.1.6 Device Operation

Packets are transmitted by placing them in the transmitq1...transmitqN, and buffers for incoming packets are placed in the receiveq1...receiveqN. In each case, the packet itself is preceded by a header:

```
struct virtio_net_hdr {
#define VIRTIO_NET_HDR_F_NEEDS_CSUM    1
    u8 flags;
#define VIRTIO_NET_HDR_GSO_NONE        0
#define VIRTIO_NET_HDR_GSO_TCPV4      1
#define VIRTIO_NET_HDR_GSO_UDP        3
#define VIRTIO_NET_HDR_GSO_TCPV6      4
#define VIRTIO_NET_HDR_GSO_ECN        0x80
    u8 gso_type;
    le16 hdr_len;
    le16 gso_size;
    le16 csum_start;
    le16 csum_offset;
    le16 num_buffers;
};
```

The controlq is used to control device features such as filtering.

5.1.6.1 Legacy Interface: Device Operation

When using the legacy interface, transitional devices and drivers MUST format the fields in struct virtio_net_hdr according to the native endian of the guest rather than (necessarily when not using the legacy interface) little-endian.

The legacy driver only presented *num_buffers* in the struct virtio_net_hdr when VIRTIO_NET_F_MRG_RXBUF was negotiated; without that feature the structure was 2 bytes shorter.

When using the legacy interface, the driver SHOULD ignore the used length for the transmit queues and the controlq queue.

Note: Historically, some devices put the total descriptor length there, even though no data was actually written.

5.1.6.2 Packet Transmission

Transmitting a single packet is simple, but varies depending on the different features the driver negotiated.

1. The driver can send a completely checksummed packet. In this case, *flags* will be zero, and *gso_type* will be `VIRTIO_NET_HDR_GSO_NONE`.
2. If the driver negotiated `VIRTIO_NET_F_CSUM`, it can skip checksumming the packet:
 - *flags* has the `VIRTIO_NET_HDR_F_NEEDS_CSUM` set,
 - *csum_start* is set to the offset within the packet to begin checksumming, and
 - *csum_offset* indicates how many bytes after the *csum_start* the new (16 bit ones' complement) checksum is placed by the device.
 - The TCP checksum field in the packet is set to the sum of the TCP pseudo header, so that replacing it by the ones' complement checksum of the TCP header and body will give the correct result.

Note: For example, consider a partially checksummed TCP (IPv4) packet. It will have a 14 byte ethernet header and 20 byte IP header followed by the TCP header (with the TCP checksum field 16 bytes into that header). *csum_start* will be $14+20 = 34$ (the TCP checksum includes the header), and *csum_offset* will be 16.

3. If the driver negotiated `VIRTIO_NET_F_HOST_TSO4`, `TSO6` or `UFO`, and the packet requires TCP segmentation or UDP fragmentation, then *gso_type* is set to `VIRTIO_NET_HDR_GSO_TCPV4`, `TCPV6` or `UDP`. (Otherwise, it is set to `VIRTIO_NET_HDR_GSO_NONE`). In this case, packets larger than 1514 bytes can be transmitted: the metadata indicates how to replicate the packet header to cut it into smaller packets. The other *gso* fields are set:
 - *hdr_len* is a hint to the device as to how much of the header needs to be kept to copy into each packet, usually set to the length of the headers, including the transport header¹.
 - *gso_size* is the maximum size of each packet beyond that header (ie. MSS).
 - If the driver negotiated the `VIRTIO_NET_F_HOST_ECN` feature, the `VIRTIO_NET_HDR_GSO_-ECN` bit in *gso_type* indicates that the TCP packet has the ECN bit set².
4. *num_buffers* is set to zero. This field is unused on transmitted packets.
5. The header and packet are added as one output descriptor to the `transmitq`, and the device is notified of the new entry (see [5.1.5 Device Initialization](#)).

5.1.6.2.1 Driver Requirements: Packet Transmission

The driver **MUST** set *num_buffers* to zero.

If `VIRTIO_NET_F_CSUM` is not negotiated, the driver **MUST** set *flags* to zero and **SHOULD** supply a fully checksummed packet to the device.

If `VIRTIO_NET_F_HOST_TSO4` is negotiated, the driver **MAY** set *gso_type* to `VIRTIO_NET_HDR_GSO_-TCPV4` to request TCPv4 segmentation, otherwise the driver **MUST NOT** set *gso_type* to `VIRTIO_NET_-HDR_GSO_TCPV4`.

If `VIRTIO_NET_F_HOST_TSO6` is negotiated, the driver **MAY** set *gso_type* to `VIRTIO_NET_HDR_GSO_-TCPV6` to request TCPv6 segmentation, otherwise the driver **MUST NOT** set *gso_type* to `VIRTIO_NET_-HDR_GSO_TCPV6`.

If `VIRTIO_NET_F_HOST_UFO` is negotiated, the driver **MAY** set *gso_type* to `VIRTIO_NET_HDR_GSO_-UDP` to request UDP segmentation, otherwise the driver **MUST NOT** set *gso_type* to `VIRTIO_NET_HDR_-GSO_UDP`.

¹Due to various bugs in implementations, this field is not useful as a guarantee of the transport header size.

²This case is not handled by some older hardware, so is called out specifically in the protocol.

The driver SHOULD NOT send to the device TCP packets requiring segmentation offload which have the Explicit Congestion Notification bit set, unless the VIRTIO_NET_F_HOST_ECN feature is negotiated, in which case the driver MUST set the VIRTIO_NET_HDR_GSO_ECN bit in *gso_type*.

If the VIRTIO_NET_F_CSUM feature has been negotiated, the driver MAY set the VIRTIO_NET_HDR_F_NEEDS_CSUM bit in *flags*, if so:

1. the driver MUST validate the packet checksum at offset *csum_offset* from *csum_start* as well as all preceding offsets;
2. the driver MUST set the packet checksum stored in the buffer to the TCP/UDP pseudo header;
3. the driver MUST set *csum_start* and *csum_offset* such that calculating a ones' complement checksum from *csum_start* up until the end of the packet and storing the result at offset *csum_offset* from *csum_start* will result in a fully checksummed packet;

If none of the VIRTIO_NET_F_HOST_TSO4, TSO6 or UFO options have been negotiated, the driver MUST set *gso_type* to VIRTIO_NET_HDR_GSO_NONE.

If *gso_type* differs from VIRTIO_NET_HDR_GSO_NONE, then the driver MUST also set the VIRTIO_NET_HDR_F_NEEDS_CSUM bit in *flags* and MUST set *gso_size* to indicate the desired MSS.

If one of the VIRTIO_NET_F_HOST_TSO4, TSO6 or UFO options have been negotiated, the driver SHOULD set *hdr_len* to a value not less than the length of the headers, including the transport header.

The driver MUST NOT set the VIRTIO_NET_HDR_F_DATA_VALID bit in *flags*.

5.1.6.2.2 Device Requirements: Packet Transmission

The device MUST ignore *flag* bits that it does not recognize.

If VIRTIO_NET_HDR_F_NEEDS_CSUM bit in *flags* is not set, the device MUST NOT use the *csum_start* and *csum_offset*.

If one of the VIRTIO_NET_F_HOST_TSO4, TSO6 or UFO options have been negotiated, the device MAY use *hdr_len* only as a hint about the transport header size. The device MUST NOT rely on *hdr_len* to be correct.

Note: This is due to various bugs in implementations.

If VIRTIO_NET_HDR_F_NEEDS_CSUM is not set, the device MUST NOT rely on the packet checksum being correct.

5.1.6.2.3 Packet Transmission Interrupt

Often a driver will suppress transmission virtqueue interrupts and check for used packets in the transmit path of following packets.

The normal behavior in this interrupt handler is to retrieve used buffers from the virtqueue and free the corresponding headers and packets.

5.1.6.3 Setting Up Receive Buffers

It is generally a good idea to keep the receive virtqueue as fully populated as possible: if it runs out, network performance will suffer.

If the VIRTIO_NET_F_GUEST_TSO4, VIRTIO_NET_F_GUEST_TSO6 or VIRTIO_NET_F_GUEST_UFO features are used, the maximum incoming packet will be to 65550 bytes long (the maximum size of a TCP or UDP packet, plus the 14 byte ethernet header), otherwise 1514 bytes. The 12-byte struct *virtio_net_hdr* is prepended to this, making for 65562 or 1526 bytes.

5.1.6.3.1 Driver Requirements: Setting Up Receive Buffers

- If VIRTIO_NET_F_MRG_RXBUF is not negotiated:
 - If VIRTIO_NET_F_GUEST_TSO4, VIRTIO_NET_F_GUEST_TSO6 or VIRTIO_NET_F_GUEST_UFO are negotiated, the driver SHOULD populate the receive queue(s) with buffers of at least 65562 bytes.
 - Otherwise, the driver SHOULD populate the receive queue(s) with buffers of at least 1526 bytes.
- If VIRTIO_NET_F_MRG_RXBUF is negotiated, each buffer MUST be at least the size of the struct `virtio_net_hdr`.

Note: Obviously each buffer can be split across multiple descriptor elements.

If VIRTIO_NET_F_MQ is negotiated, each of `receiveq1...receiveqN` that will be used SHOULD be populated with receive buffers.

5.1.6.3.2 Device Requirements: Setting Up Receive Buffers

The device MUST set `num_buffers` to the number of descriptors used to hold the incoming packet.

The device MUST use only a single descriptor if VIRTIO_NET_F_MRG_RXBUF was not negotiated.

Note: This means that `num_buffers` will always be 1 if VIRTIO_NET_F_MRG_RXBUF is not negotiated.

5.1.6.4 Processing of Incoming Packets

When a packet is copied into a buffer in the `receiveq`, the optimal path is to disable further interrupts for the `receiveq` and process packets until no more are found, then re-enable them.

Processing incoming packets involves:

1. `num_buffers` indicates how many descriptors this packet is spread over (including this one): this will always be 1 if VIRTIO_NET_F_MRG_RXBUF was not negotiated. This allows receipt of large packets without having to allocate large buffers: a packet that does not fit in a single buffer can flow over to the next buffer, and so on. In this case, there will be at least `num_buffers` used buffers in the `virtqueue`, and the device chains them together to form a single packet in a way similar to how it would store it in a single buffer spread over multiple descriptors. The other buffers will not begin with a struct `virtio_net_hdr`.
2. If `num_buffers` is one, then the entire packet will be contained within this buffer, immediately following the struct `virtio_net_hdr`.
3. If the VIRTIO_NET_F_GUEST_CSUM feature was negotiated, the VIRTIO_NET_HDR_F_DATA_VALID bit in `flags` can be set: if so, device has validated the packet checksum. In case of multiple encapsulated protocols, one level of checksums has been validated.

Additionally, VIRTIO_NET_F_GUEST_CSUM, TSO4, TSO6, UDP and ECN features enable receive checksum, large receive offload and ECN support which are the input equivalents of the transmit checksum, transmit segmentation offloading and ECN features, as described in [5.1.6.2](#):

1. If the VIRTIO_NET_F_GUEST_CSUM feature was negotiated, the VIRTIO_NET_HDR_F_NEEDS_CSUM bit in `flags` can be set: if so, the packet checksum at offset `csum_offset` from `csum_start` and any preceding checksums have been validated. The checksum on the packet is incomplete and `csum_start` and `csum_offset` indicate how to calculate it (see Packet Transmission point 1).
2. If the VIRTIO_NET_F_GUEST_TSO4, TSO6 or UFO options were negotiated, then `gso_type` MAY be something other than VIRTIO_NET_HDR_GSO_NONE, and `gso_size` field indicates the desired MSS (see Packet Transmission point 2).

5.1.6.4.1 Device Requirements: Processing of Incoming Packets

If VIRTIO_NET_F_MRG_RXBUF has not been negotiated, the device MUST set *num_buffers* to 1.

If VIRTIO_NET_F_MRG_RXBUF has been negotiated, the device MUST set *num_buffers* to indicate the number of buffers the packet (including the header) is spread over.

If a receive packet is spread over multiple buffers, the device MUST use all buffers but the last (i.e. the first *num_buffers* - 1 buffers) completely up to the full length of each buffer supplied by the driver.

The device MUST use all buffers used by a single receive packet together, such that at least *num_buffers* are observed by driver as used.

If VIRTIO_NET_F_GUEST_CSUM is not negotiated, the device MUST set *flags* to zero and SHOULD supply a fully checksummed packet to the driver.

If VIRTIO_NET_F_GUEST_TSO4 is not negotiated, the device MUST NOT set *gso_type* to VIRTIO_NET_HDR_GSO_TCPV4.

If VIRTIO_NET_F_GUEST_UDP is not negotiated, the device MUST NOT set *gso_type* to VIRTIO_NET_HDR_GSO_UDP.

If VIRTIO_NET_F_GUEST_TSO6 is not negotiated, the device MUST NOT set *gso_type* to VIRTIO_NET_HDR_GSO_TCPV6.

The device SHOULD NOT send to the driver TCP packets requiring segmentation offload which have the Explicit Congestion Notification bit set, unless the VIRTIO_NET_F_GUEST_ECN feature is negotiated, in which case the device MUST set the VIRTIO_NET_HDR_GSO_ECN bit in *gso_type*.

If the VIRTIO_NET_F_GUEST_CSUM feature has been negotiated, the device MAY set the VIRTIO_NET_HDR_F_NEEDS_CSUM bit in *flags*, if so:

1. the device MUST validate the packet checksum at offset *csum_offset* from *csum_start* as well as all preceding offsets;
2. the device MUST set the packet checksum stored in the receive buffer to the TCP/UDP pseudo header;
3. the device MUST set *csum_start* and *csum_offset* such that calculating a ones' complement checksum from *csum_start* up until the end of the packet and storing the result at offset *csum_offset* from *csum_start* will result in a fully checksummed packet;

If none of the VIRTIO_NET_F_GUEST_TSO4, TSO6 or UFO options have been negotiated, the device MUST set *gso_type* to VIRTIO_NET_HDR_GSO_NONE.

If *gso_type* differs from VIRTIO_NET_HDR_GSO_NONE, then the device MUST also set the VIRTIO_NET_HDR_F_NEEDS_CSUM bit in *flags* MUST set *gso_size* to indicate the desired MSS.

If one of the VIRTIO_NET_F_GUEST_TSO4, TSO6 or UFO options have been negotiated, the device SHOULD set *hdr_len* to a value not less than the length of the headers, including the transport header.

If the VIRTIO_NET_F_GUEST_CSUM feature has been negotiated, the device MAY set the VIRTIO_NET_HDR_F_DATA_VALID bit in *flags*, if so, the device MUST validate the packet checksum (in case of multiple encapsulated protocols, one level of checksums is validated).

5.1.6.4.2 Driver Requirements: Processing of Incoming Packets

The driver MUST ignore *flag* bits that it does not recognize.

If VIRTIO_NET_HDR_F_NEEDS_CSUM bit in *flags* is not set, the driver MUST NOT use the *csum_start* and *csum_offset*.

If one of the VIRTIO_NET_F_GUEST_TSO4, TSO6 or UFO options have been negotiated, the driver MAY use *hdr_len* only as a hint about the transport header size. The driver MUST NOT rely on *hdr_len* to be correct.

Note: This is due to various bugs in implementations.

If neither `VIRTIO_NET_HDR_F_NEEDS_CSUM` nor `VIRTIO_NET_HDR_F_DATA_VALID` is set, the driver **MUST NOT** rely on the packet checksum being correct.

5.1.6.5 Control Virtqueue

The driver uses the control virtqueue (if `VIRTIO_NET_F_CTRL_VQ` is negotiated) to send commands to manipulate various features of the device which would not easily map into the configuration space.

All commands are of the following form:

```
struct virtio_net_ctrl {
    u8 class;
    u8 command;
    u8 command-specific-data[];
    u8 ack;
};

/* ack values */
#define VIRTIO_NET_OK      0
#define VIRTIO_NET_ERR    1
```

The *class*, *command* and *command-specific-data* are set by the driver, and the device sets the *ack* byte. There is little it can do except issue a diagnostic if *ack* is not `VIRTIO_NET_OK`.

5.1.6.5.1 Packet Receive Filtering

If the `VIRTIO_NET_F_CTRL_RX` and `VIRTIO_NET_F_CTRL_RX_EXTRA` features are negotiated, the driver can send control commands for promiscuous mode, multicast, unicast and broadcast receiving.

Note: In general, these commands are best-effort: unwanted packets could still arrive.

```
#define VIRTIO_NET_CTRL_RX      0
#define VIRTIO_NET_CTRL_RX_PROMISC    0
#define VIRTIO_NET_CTRL_RX_ALLMULTI  1
#define VIRTIO_NET_CTRL_RX_ALLUNI    2
#define VIRTIO_NET_CTRL_RX_NOMULTI   3
#define VIRTIO_NET_CTRL_RX_NOUNI     4
#define VIRTIO_NET_CTRL_RX_NOBCAST    5
```

5.1.6.5.1.1 Device Requirements: Packet Receive Filtering

If the `VIRTIO_NET_F_CTRL_RX` feature has been negotiated, the device **MUST** support the following `VIRTIO_NET_CTRL_RX` class commands:

- `VIRTIO_NET_CTRL_RX_PROMISC` turns promiscuous mode on and off. The command-specific-data is one byte containing 0 (off) or 1 (on). If promiscuous mode is on, the device **SHOULD** receive all incoming packets. This **SHOULD** take effect even if one of the other modes set by a `VIRTIO_NET_CTRL_RX` class command is on.
- `VIRTIO_NET_CTRL_RX_ALLMULTI` turns all-multicast receive on and off. The command-specific-data is one byte containing 0 (off) or 1 (on). When all-multicast receive is on the device **SHOULD** allow all incoming multicast packets.

If the `VIRTIO_NET_F_CTRL_RX_EXTRA` feature has been negotiated, the device **MUST** support the following `VIRTIO_NET_CTRL_RX` class commands:

- `VIRTIO_NET_CTRL_RX_ALLUNI` turns all-unicast receive on and off. The command-specific-data is one byte containing 0 (off) or 1 (on). When all-unicast receive is on the device **SHOULD** allow all incoming unicast packets.

- VIRTIO_NET_CTRL_RX_NOMULTI suppresses multicast receive. The command-specific-data is one byte containing 0 (multicast receive allowed) or 1 (multicast receive suppressed). When multicast receive is suppressed, the device SHOULD NOT send multicast packets to the driver. This SHOULD take effect even if VIRTIO_NET_CTRL_RX_ALLMULTI is on. This filter SHOULD NOT apply to broadcast packets.
- VIRTIO_NET_CTRL_RX_NOUNI suppresses unicast receive. The command-specific-data is one byte containing 0 (unicast receive allowed) or 1 (unicast receive suppressed). When unicast receive is suppressed, the device SHOULD NOT send unicast packets to the driver. This SHOULD take effect even if VIRTIO_NET_CTRL_RX_ALLUNI is on.
- VIRTIO_NET_CTRL_RX_NOBCAST suppresses broadcast receive. The command-specific-data is one byte containing 0 (broadcast receive allowed) or 1 (broadcast receive suppressed). When broadcast receive is suppressed, the device SHOULD NOT send broadcast packets to the driver. This SHOULD take effect even if VIRTIO_NET_CTRL_RX_ALLMULTI is on.

5.1.6.5.1.2 Driver Requirements: Packet Receive Filtering

If the VIRTIO_NET_F_CTRL_RX feature has not been negotiated, the driver MUST NOT issue commands VIRTIO_NET_CTRL_RX_PROMISC or VIRTIO_NET_CTRL_RX_ALLMULTI.

If the VIRTIO_NET_F_CTRL_RX_EXTRA feature has not been negotiated, the driver MUST NOT issue commands VIRTIO_NET_CTRL_RX_ALLUNI, VIRTIO_NET_CTRL_RX_NOMULTI, VIRTIO_NET_CTRL_RX_NOUNI or VIRTIO_NET_CTRL_RX_NOBCAST.

5.1.6.5.2 Setting MAC Address Filtering

If the VIRTIO_NET_F_CTRL_RX feature is negotiated, the driver can send control commands for MAC address filtering.

```
struct virtio_net_ctrl_mac {
    le32 entries;
    u8 macs[entries][6];
};

#define VIRTIO_NET_CTRL_MAC      1
#define VIRTIO_NET_CTRL_MAC_TABLE_SET    0
#define VIRTIO_NET_CTRL_MAC_ADDR_SET    1
```

The device can filter incoming packets by any number of destination MAC addresses³. This table is set using the class VIRTIO_NET_CTRL_MAC and the command VIRTIO_NET_CTRL_MAC_TABLE_SET. The command-specific-data is two variable length tables of 6-byte MAC addresses (as described in struct virtio_net_ctrl_mac). The first table contains unicast addresses, and the second contains multicast addresses.

The VIRTIO_NET_CTRL_MAC_ADDR_SET command is used to set the default MAC address which rx filtering accepts (and if VIRTIO_NET_F_MAC_ADDR has been negotiated, this will be reflected in *mac* in config space).

The command-specific-data for VIRTIO_NET_CTRL_MAC_ADDR_SET is the 6-byte MAC address.

5.1.6.5.2.1 Device Requirements: Setting MAC Address Filtering

The device MUST have an empty MAC filtering table on reset.

The device MUST update the MAC filtering table before it consumes the VIRTIO_NET_CTRL_MAC_TABLE_SET command.

³Since there are no guarantees, it can use a hash filter or silently switch to allmulti or promiscuous mode if it is given too many addresses.

The device MUST update *mac* in config space before it consumes the VIRTIO_NET_CTRL_MAC_ADDR_SET command, if VIRTIO_NET_F_MAC_ADDR has been negotiated.

The device SHOULD drop incoming packets which have a destination MAC which matches neither the *mac* (or that set with VIRTIO_NET_CTRL_MAC_ADDR_SET) nor the MAC filtering table.

5.1.6.5.2.2 Driver Requirements: Setting MAC Address Filtering

If VIRTIO_NET_F_CTRL_RX has not been negotiated, the driver MUST NOT issue VIRTIO_NET_CTRL_MAC class commands.

If VIRTIO_NET_F_CTRL_RX has been negotiated, the driver SHOULD issue VIRTIO_NET_CTRL_MAC_ADDR_SET to set the default mac if it is different from *mac*.

The driver MUST follow the VIRTIO_NET_CTRL_MAC_TABLE_SET command by a le32 number, followed by that number of non-multicast MAC addresses, followed by another le32 number, followed by that number of multicast addresses. Either number MAY be 0.

5.1.6.5.2.3 Legacy Interface: Setting MAC Address Filtering

When using the legacy interface, transitional devices and drivers MUST format *entries* in struct `virtio_net_ctrl_mac` according to the native endian of the guest rather than (necessarily when not using the legacy interface) little-endian.

Legacy drivers that didn't negotiate VIRTIO_NET_F_CTRL_MAC_ADDR changed *mac* in config space when NIC is accepting incoming packets. These drivers always wrote the mac value from first to last byte, therefore after detecting such drivers, a transitional device MAY defer MAC update, or MAY defer processing incoming packets until driver writes the last byte of *mac* in the config space.

5.1.6.5.3 VLAN Filtering

If the driver negotiates the VIRTIO_NET_F_CTRL_VLAN feature, it can control a VLAN filter table in the device.

```
#define VIRTIO_NET_CTRL_VLAN      2
#define VIRTIO_NET_CTRL_VLAN_ADD  0
#define VIRTIO_NET_CTRL_VLAN_DEL  1
```

Both the VIRTIO_NET_CTRL_VLAN_ADD and VIRTIO_NET_CTRL_VLAN_DEL command take a little-endian 16-bit VLAN id as the command-specific-data.

5.1.6.5.3.1 Legacy Interface: VLAN Filtering

When using the legacy interface, transitional devices and drivers MUST format the VLAN id according to the native endian of the guest rather than (necessarily when not using the legacy interface) little-endian.

5.1.6.5.4 Gratuitous Packet Sending

If the driver negotiates the VIRTIO_NET_F_GUEST_ANNOUNCE (depends on VIRTIO_NET_F_CTRL_VQ), the device can ask the driver to send gratuitous packets; this is usually done after the guest has been physically migrated, and needs to announce its presence on the new network links. (As hypervisor does not have the knowledge of guest network configuration (eg. tagged vlan) it is simplest to prod the guest in this way).

```
#define VIRTIO_NET_CTRL_ANNOUNCE  3
#define VIRTIO_NET_CTRL_ANNOUNCE_ACK 0
```

The driver checks VIRTIO_NET_S_ANNOUNCE bit in the device configuration *status* field when it notices the changes of device configuration. The command VIRTIO_NET_CTRL_ANNOUNCE_ACK is used to indicate that driver has received the notification and device clears the VIRTIO_NET_S_ANNOUNCE bit in *status*.

Processing this notification involves:

1. Sending the gratuitous packets (eg. ARP) or marking there are pending gratuitous packets to be sent and letting deferred routine to send them.
2. Sending VIRTIO_NET_CTRL_ANNOUNCE_ACK command through control vq.

5.1.6.5.4.1 Driver Requirements: Gratuitous Packet Sending

If the driver negotiates VIRTIO_NET_F_GUEST_ANNOUNCE, it SHOULD notify network peers of its new location after it sees the VIRTIO_NET_S_ANNOUNCE bit in *status*. The driver MUST send a command on the command queue with class VIRTIO_NET_CTRL_ANNOUNCE and command VIRTIO_NET_CTRL_ANNOUNCE_ACK.

5.1.6.5.4.2 Device Requirements: Gratuitous Packet Sending

If VIRTIO_NET_F_GUEST_ANNOUNCE is negotiated, the device MUST clear the VIRTIO_NET_S_ANNOUNCE bit in *status* upon receipt of a command buffer with class VIRTIO_NET_CTRL_ANNOUNCE and command VIRTIO_NET_CTRL_ANNOUNCE_ACK before marking the buffer as used.

5.1.6.5.5 Automatic receive steering in multiqueue mode

If the driver negotiates the VIRTIO_NET_F_MQ feature bit (depends on VIRTIO_NET_F_CTRL_VQ), it MAY transmit outgoing packets on one of the multiple transmitq1...transmitqN and ask the device to queue incoming packets into one of the multiple receiveq1...receiveqN depending on the packet flow.

```
struct virtio_net_ctrl_mq {
    le16 virtqueue_pairs;
};

#define VIRTIO_NET_CTRL_MQ      4
#define VIRTIO_NET_CTRL_MQ_VQ_PAIRS_SET    0
#define VIRTIO_NET_CTRL_MQ_VQ_PAIRS_MIN    1
#define VIRTIO_NET_CTRL_MQ_VQ_PAIRS_MAX    0x8000
```

Multiqueue is disabled by default. The driver enables multiqueue by executing the VIRTIO_NET_CTRL_MQ_VQ_PAIRS_SET command, specifying the number of the transmit and receive queues to be used up to *max_virtqueue_pairs*; subsequently, transmitq1...transmitqn and receiveq1...receiveqn where n=*virtqueue_pairs* MAY be used.

When multiqueue is enabled, the device MUST use automatic receive steering based on packet flow. Programming of the receive steering classifier is implicit. After the driver transmitted a packet of a flow on transmitqX, the device SHOULD cause incoming packets for that flow to be steered to receiveqX. For uni-directional protocols, or where no packets have been transmitted yet, the device MAY steer a packet to a random queue out of the specified receiveq1...receiveqn.

Multiqueue is disabled by setting *virtqueue_pairs* to 1 (this is the default) and waiting for the device to use the command buffer.

5.1.6.5.5.1 Driver Requirements: Automatic receive steering in multiqueue mode

The driver MUST configure the virtqueues before enabling them with the VIRTIO_NET_CTRL_MQ_VQ_PAIRS_SET command.

The driver MUST NOT request a *virtqueue_pairs* of 0 or greater than *max_virtqueue_pairs* in the device configuration space.

The driver MUST queue packets only on any transmitq1 before the VIRTIO_NET_CTRL_MQ_VQ_PAIRS_SET command.

The driver MUST NOT queue packets on transmit queues greater than *virtqueue_pairs* once it has placed the VIRTIO_NET_CTRL_MQ_VQ_PAIRS_SET command in the available ring.

5.1.6.5.5.2 Device Requirements: Automatic receive steering in multiqueue mode

The device MUST queue packets only on any receiveq1 before the VIRTIO_NET_CTRL_MQ_VQ_PAIRS_SET command.

The device MUST NOT queue packets on receive queues greater than *virtqueue_pairs* once it has placed the VIRTIO_NET_CTRL_MQ_VQ_PAIRS_SET command in a used buffer.

5.1.6.5.5.3 Legacy Interface: Automatic receive steering in multiqueue mode

When using the legacy interface, transitional devices and drivers MUST format *virtqueue_pairs* according to the native endian of the guest rather than (necessarily when not using the legacy interface) little-endian.

5.1.6.5.6 Offloads State Configuration

If the VIRTIO_NET_F_CTRL_GUEST_OFFLOADS feature is negotiated, the driver can send control commands for dynamic offloads state configuration.

5.1.6.5.6.1 Setting Offloads State

```
le64 offloads;

#define VIRTIO_NET_F_GUEST_CSUM      1
#define VIRTIO_NET_F_GUEST_TSO4     7
#define VIRTIO_NET_F_GUEST_TSO6     8
#define VIRTIO_NET_F_GUEST_ECN      9
#define VIRTIO_NET_F_GUEST_UFO     10

#define VIRTIO_NET_CTRL_GUEST_OFFLOADS 5
#define VIRTIO_NET_CTRL_GUEST_OFFLOADS_SET 0
```

The class VIRTIO_NET_CTRL_GUEST_OFFLOADS has one command: VIRTIO_NET_CTRL_GUEST_OFFLOADS_SET applies the new offloads configuration.

le64 value passed as command data is a bitmask, bits set define offloads to be enabled, bits cleared - offloads to be disabled.

There is a corresponding device feature for each offload. Upon feature negotiation corresponding offload gets enabled to preserve backward compatibility.

5.1.6.5.6.2 Driver Requirements: Setting Offloads State

A driver MUST NOT enable an offload for which the appropriate feature has not been negotiated.

5.1.6.5.6.3 Legacy Interface: Setting Offloads State

When using the legacy interface, transitional devices and drivers MUST format *offloads* according to the native endian of the guest rather than (necessarily when not using the legacy interface) little-endian.

5.1.6.6 Legacy Interface: Framing Requirements

When using legacy interfaces, transitional drivers which have not negotiated VIRTIO_F_ANY_LAYOUT MUST use a single descriptor for the struct `virtio_net_hdr` on both transmit and receive, with the network data in the following descriptors.

Additionally, when using the control virtqueue (see 5.1.6.5) , transitional drivers which have not negotiated VIRTIO_F_ANY_LAYOUT MUST:

- for all commands, use a single 2-byte descriptor including the first two fields: *class* and *command*
- for all commands except VIRTIO_NET_CTRL_MAC_TABLE_SET use a single descriptor including command-specific-data with no padding.
- for the VIRTIO_NET_CTRL_MAC_TABLE_SET command use exactly two descriptors including command-specific-data with no padding: the first of these descriptors MUST include the `virtio_net_ctrl_mac` table structure for the unicast addresses with no padding, the second of these descriptors MUST include the `virtio_net_ctrl_mac` table structure for the multicast addresses with no padding.
- for all commands, use a single 1-byte descriptor for the *ack* field

See 2.5.4.

5.2 Block Device

The virtio block device is a simple virtual block device (ie. disk). Read and write requests (and other exotic requests) are placed in the queue, and serviced (probably out of order) by the device except where noted.

5.2.1 Device ID

2

5.2.2 Virtqueues

0 requestq

5.2.3 Feature bits

VIRTIO_BLK_F_SIZE_MAX (1) Maximum size of any single segment is in *size_max*.

VIRTIO_BLK_F_SEG_MAX (2) Maximum number of segments in a request is in *seg_max*.

VIRTIO_BLK_F_GEOMETRY (4) Disk-style geometry specified in *geometry*.

VIRTIO_BLK_F_RO (5) Device is read-only.

VIRTIO_BLK_F_BLK_SIZE (6) Block size of disk is in *blk_size*.

VIRTIO_BLK_F_FLUSH (9) Cache flush command support.

VIRTIO_BLK_F_TOPOLOGY (10) Device exports information on optimal I/O alignment.

VIRTIO_BLK_F_CONFIG_WCE (11) Device can toggle its cache between writeback and writethrough modes.

5.2.3.1 Legacy Interface: Feature bits

VIRTIO_BLK_F_BARRIER (0) Device supports request barriers.

VIRTIO_BLK_F SCSI (7) Device supports scsi packet commands.

Note: In the legacy interface, VIRTIO_BLK_F_FLUSH was also called VIRTIO_BLK_F_WCE.

5.2.4 Device configuration layout

The *capacity* of the device (expressed in 512-byte sectors) is always present. The availability of the others all depend on various feature bits as indicated above.

```
struct virtio_blk_config {
    le64 capacity;
    le32 size_max;
    le32 seg_max;
    struct virtio_blk_geometry {
        le16 cylinders;
        u8 heads;
        u8 sectors;
    } geometry;
    le32 blk_size;
    struct virtio_blk_topology {
        // # of logical blocks per physical block (log2)
        u8 physical_block_exp;
        // offset of first aligned logical block
        u8 alignment_offset;
        // suggested minimum I/O size in blocks
        le16 min_io_size;
        // optimal (suggested maximum) I/O size in blocks
        le32 opt_io_size;
    } topology;
    u8 writeback;
};
```

5.2.4.1 Legacy Interface: Device configuration layout

When using the legacy interface, transitional devices and drivers **MUST** format the fields in struct `virtio_blk_config` according to the native endian of the guest rather than (necessarily when not using the legacy interface) little-endian.

5.2.5 Device Initialization

1. The device size can be read from *capacity*.
2. If the VIRTIO_BLK_F_BLK_SIZE feature is negotiated, *blk_size* can be read to determine the optimal sector size for the driver to use. This does not affect the units used in the protocol (always 512 bytes), but awareness of the correct value can affect performance.
3. If the VIRTIO_BLK_F_RO feature is set by the device, any write requests will fail.
4. If the VIRTIO_BLK_F_TOPOLOGY feature is negotiated, the fields in the *topology* struct can be read to determine the physical block size and optimal I/O lengths for the driver to use. This also does not affect the units in the protocol, only performance.
5. If the VIRTIO_BLK_F_CONFIG_WCE feature is negotiated, the cache mode can be read or set through the *writeback* field. 0 corresponds to a writethrough cache, 1 to a writeback cache⁴. The cache mode

⁴Consistent with 5.2.6.2, a writethrough cache can be defined broadly as a cache that commits writes to persistent device backend storage before reporting their completion. For example, a battery-backed writeback cache actually counts as writethrough according to this definition.

after reset can be either writeback or writethrough. The actual mode can be determined by reading *writeback* after feature negotiation.

5.2.5.1 Driver Requirements: Device Initialization

Drivers SHOULD NOT negotiate VIRTIO_BLK_F_FLUSH if they are incapable of sending VIRTIO_BLK_T_FLUSH commands.

If neither VIRTIO_BLK_F_CONFIG_WCE nor VIRTIO_BLK_F_FLUSH are negotiated, the driver MAY deduce the presence of a writethrough cache. If VIRTIO_BLK_F_CONFIG_WCE was not negotiated but VIRTIO_BLK_F_FLUSH was, the driver SHOULD assume presence of a writeback cache.

The driver MUST NOT read *writeback* before setting the FEATURES_OK *status* bit.

5.2.5.2 Device Requirements: Device Initialization

Devices SHOULD always offer VIRTIO_BLK_F_FLUSH, and MUST offer it if they offer VIRTIO_BLK_F_CONFIG_WCE.

If VIRTIO_BLK_F_CONFIG_WCE is negotiated but VIRTIO_BLK_F_FLUSH is not, the device MUST initialize *writeback* to 0.

5.2.5.3 Legacy Interface: Device Initialization

Because legacy devices do not have FEATURES_OK, transitional devices MUST implement slightly different behavior around feature negotiation when used through the legacy interface. In particular, when using the legacy interface:

- the driver MAY read or write *writeback* before setting the DRIVER or DRIVER_OK *status* bit
- the device MUST NOT modify the cache mode (and *writeback*) as a result of a driver setting a status bit, unless the DRIVER_OK bit is being set and the driver has not set the VIRTIO_BLK_F_CONFIG_WCE driver feature bit.
- the device MUST NOT modify the cache mode (and *writeback*) as a result of a driver modifying the driver feature bits, for example if the driver sets the VIRTIO_BLK_F_CONFIG_WCE driver feature bit but does not set the VIRTIO_BLK_F_FLUSH bit.

5.2.6 Device Operation

The driver queues requests to the virtqueue, and they are used by the device (not necessarily in order). Each request is of form:

```
struct virtio_blk_req {
    le32 type;
    le32 reserved;
    le64 sector;
    u8 data[][512];
    u8 status;
};
```

The type of the request is either a read (VIRTIO_BLK_T_IN), a write (VIRTIO_BLK_T_OUT), or a flush (VIRTIO_BLK_T_FLUSH).

```
#define VIRTIO_BLK_T_IN      0
#define VIRTIO_BLK_T_OUT    1
#define VIRTIO_BLK_T_FLUSH  4
```

The *sector* number indicates the offset (multiplied by 512) where the read or write is to occur. This field is unused and set to 0 for scsi packet commands and for flush commands.

The final *status* byte is written by the device: either VIRTIO_BLK_S_OK for success, VIRTIO_BLK_S_IOERR for device or driver error or VIRTIO_BLK_S_UNSUPP for a request unsupported by device:

```
#define VIRTIO_BLK_S_OK      0
#define VIRTIO_BLK_S_IOERR  1
#define VIRTIO_BLK_S_UNSUPP 2
```

5.2.6.1 Driver Requirements: Device Operation

A driver **MUST NOT** submit a request which would cause a read or write beyond *capacity*.

A driver **SHOULD** accept the VIRTIO_BLK_F_RO feature if offered.

A driver **MUST** set *sector* to 0 for a VIRTIO_BLK_T_FLUSH request. A driver **SHOULD NOT** include any data in a VIRTIO_BLK_T_FLUSH request.

If the VIRTIO_BLK_F_CONFIG_WCE feature is negotiated, the driver **MAY** switch to writethrough or write-back mode by writing respectively 0 and 1 to the *writeback* field. After writing a 0 to *writeback*, the driver **MUST NOT** assume that any volatile writes have been committed to persistent device backend storage.

5.2.6.2 Device Requirements: Device Operation

A device **MUST** set the *status* byte to VIRTIO_BLK_S_IOERR for a write request if the VIRTIO_BLK_F_RO feature is offered, and **MUST NOT** write any data.

A write is considered volatile when it is submitted; the contents of sectors covered by a volatile write are undefined in persistent device backend storage until the write becomes stable. A write becomes stable once it is completed and one or more of the following conditions is true:

1. neither VIRTIO_BLK_F_CONFIG_WCE nor VIRTIO_BLK_F_FLUSH feature were negotiated, but VIRTIO_BLK_F_FLUSH was offered by the device;
2. the VIRTIO_BLK_F_CONFIG_WCE feature was negotiated and the *writeback* field in configuration space was 0 **all the time between the submission of the write and its completion**;
3. a VIRTIO_BLK_T_FLUSH request is sent **after the write is completed** and is completed itself.

If the device is backed by persistent storage, the device **MUST** ensure that stable writes are committed to it, before reporting completion of the write (cases 1 and 2) or the flush (case 3). Failure to do so can cause data loss in case of a crash.

If the driver changes *writeback* between the submission of the write and its completion, the write could be either volatile or stable when its completion is reported; in other words, the exact behavior is undefined.

If VIRTIO_BLK_F_FLUSH was not offered by the device⁵, the device **MAY** also commit writes to persistent device backend storage before reporting their completion. Unlike case 1, however, this is not an absolute requirement of the specification.

Note: An implementation that does not offer VIRTIO_BLK_F_FLUSH and does not commit completed writes will not be resilient to data loss in case of crashes. Not offering VIRTIO_BLK_F_FLUSH is an absolute requirement for implementations that do not wish to be safe against such data losses.

5.2.6.3 Legacy Interface: Device Operation

When using the legacy interface, transitional devices and drivers **MUST** format the fields in struct virtio_blk_req according to the native endian of the guest rather than (necessarily when not using the legacy interface) little-endian.

⁵Note that in this case, according to 5.2.5.2, the device will not have offered VIRTIO_BLK_F_CONFIG_WCE either.

When using the legacy interface, transitional drivers SHOULD ignore the used length values.

Note: Historically, some devices put the total descriptor length, or the total length of device-writable buffers there, even when only the status byte was actually written.

The *reserved* field was previously called *ioprio*. *ioprio* is a hint about the relative priorities of requests to the device: higher numbers indicate more important requests.

```
#define VIRTIO_BLK_T_FLUSH_OUT    5
```

The command `VIRTIO_BLK_T_FLUSH_OUT` was a synonym for `VIRTIO_BLK_T_FLUSH`; a driver MUST treat it as a `VIRTIO_BLK_T_FLUSH` command.

```
#define VIRTIO_BLK_T_BARRIER    0x80000000
```

If the device has `VIRTIO_BLK_F_BARRIER` feature the high bit (`VIRTIO_BLK_T_BARRIER`) indicates that this request acts as a barrier and that all preceding requests SHOULD be complete before this one, and all following requests SHOULD NOT be started until this is complete.

Note: A barrier does not flush caches in the underlying backend device in host, and thus does not serve as data consistency guarantee. Only a `VIRTIO_BLK_T_FLUSH` request does that.

Some older legacy devices did not commit completed writes to persistent device backend storage when `VIRTIO_BLK_F_FLUSH` was offered but not negotiated. In order to work around this, the driver MAY set the *writeback* to 0 (if available) or it MAY send an explicit flush request after every completed write.

If the device has `VIRTIO_BLK_F_SCSI` feature, it can also support scsi packet command requests, each of these requests is of form:

```
/* All fields are in guest's native endian. */
struct virtio_scsi_pc_req {
    u32 type;
    u32 ioprio;
    u64 sector;
    u8 cmd[];
    u8 data[][512];
#define SCSI_SENSE_BUFFERSIZE    96
    u8 sense[SCSI_SENSE_BUFFERSIZE];
    u32 errors;
    u32 data_len;
    u32 sense_len;
    u32 residual;
    u8 status;
};
```

A request type can also be a scsi packet command (`VIRTIO_BLK_T_SCSI_CMD` or `VIRTIO_BLK_T_SCSI_CMD_OUT`). The two types are equivalent, the device does not distinguish between them:

```
#define VIRTIO_BLK_T_SCSI_CMD      2
#define VIRTIO_BLK_T_SCSI_CMD_OUT 3
```

The *cmd* field is only present for scsi packet command requests, and indicates the command to perform. This field MUST reside in a single, separate device-readable buffer; command length can be derived from the length of this buffer.

Note that these first three (four for scsi packet commands) fields are always device-readable: *data* is either device-readable or device-writable, depending on the request. The size of the read or write can be derived from the total size of the request buffers.

sense is only present for scsi packet command requests, and indicates the buffer for scsi sense data.

data_len is only present for scsi packet command requests, this field is deprecated, and SHOULD be ignored by the driver. Historically, devices copied data length there.

sense_len is only present for scsi packet command requests and indicates the number of bytes actually written to the *sense* buffer.

residual field is only present for scsi packet command requests and indicates the residual size, calculated as data length - number of bytes actually transferred.

5.2.6.4 Legacy Interface: Framing Requirements

When using legacy interfaces, transitional drivers which have not negotiated VIRTIO_F_ANY_LAYOUT:

- MUST use a single 8-byte descriptor containing *type*, *reserved* and *sector*, followed by descriptors for *data*, then finally a separate 1-byte descriptor for *status*.
- For SCSI commands there are additional constraints. *errors*, *data_len*, *sense_len* and *residual* MUST reside in a single, separate device-writable descriptor, *sense* MUST reside in a single separate device-writable descriptor of size 96 bytes, and *errors*, *data_len*, *sense_len* and *residual* MUST reside a single separate device-writable descriptor.

See 2.5.4.

5.3 Console Device

The virtio console device is a simple device for data input and output. A device MAY have one or more ports. Each port has a pair of input and output virtqueues. Moreover, a device has a pair of control IO virtqueues. The control virtqueues are used to communicate information between the device and the driver about ports being opened and closed on either side of the connection, indication from the device about whether a particular port is a console port, adding new ports, port hot-plug/unplug, etc., and indication from the driver about whether a port or a device was successfully added, port open/close, etc. For data IO, one or more empty buffers are placed in the receive queue for incoming data and outgoing characters are placed in the transmit queue.

5.3.1 Device ID

3

5.3.2 Virtqueues

- 0 receiveq(port0)
- 1 transmitq(port0)
- 2 control receiveq
- 3 control transmitq
- 4 receiveq(port1)
- 5 transmitq(port1)

...

The port 0 receive and transmit queues always exist: other queues only exist if VIRTIO_CONSOLE_F_MULTIPORT is set.

5.3.3 Feature bits

VIRTIO_CONSOLE_F_SIZE (0) Configuration *cols* and *rows* are valid.

VIRTIO_CONSOLE_F_MULTIPORT (1) Device has support for multiple ports; *max_nr_ports* is valid and control virtqueues will be used.

VIRTIO_CONSOLE_F_EMERG_WRITE (2) Device has support for emergency write. Configuration field `emerg_wr` is valid.

5.3.4 Device configuration layout

The size of the console is supplied in the configuration space if the `VIRTIO_CONSOLE_F_SIZE` feature is set. Furthermore, if the `VIRTIO_CONSOLE_F_MULTIPORT` feature is set, the maximum number of ports supported by the device can be fetched.

If `VIRTIO_CONSOLE_F_EMERG_WRITE` is set then the driver can use emergency write to output a single character without initializing virtio queues, or even acknowledging the feature.

```
struct virtio_console_config {
    le16 cols;
    le16 rows;
    le32 max_nr_ports;
    le32 emerg_wr;
};
```

5.3.4.1 Legacy Interface: Device configuration layout

When using the legacy interface, transitional devices and drivers **MUST** format the fields in struct `virtio_console_config` according to the native endian of the guest rather than (necessarily when not using the legacy interface) little-endian.

5.3.5 Device Initialization

1. If the `VIRTIO_CONSOLE_F_EMERG_WRITE` feature is offered, *emerg_wr* field of the configuration can be written at any time. Thus it works for very early boot debugging output as well as catastrophic OS failures (eg. virtio ring corruption).
2. If the `VIRTIO_CONSOLE_F_SIZE` feature is negotiated, the driver can read the console dimensions from *cols* and *rows*.
3. If the `VIRTIO_CONSOLE_F_MULTIPORT` feature is negotiated, the driver can spawn multiple ports, not all of which are necessarily attached to a console. Some could be generic ports. In this case, the control virtqueues are enabled and according to *max_nr_ports*, the appropriate number of virtqueues are created. A control message indicating the driver is ready is sent to the device. The device can then send control messages for adding new ports to the device. After creating and initializing each port, a `VIRTIO_CONSOLE_PORT_READY` control message is sent to the device for that port so the device can let the driver know of any additional configuration options set for that port.
4. The receiveq for each port is populated with one or more receive buffers.

5.3.5.1 Device Requirements: Device Initialization

The device **MUST** allow a write to *emerg_wr*, even on an unconfigured device.

The device **SHOULD** transmit the lower byte written to *emerg_wr* to an appropriate log or output method.

5.3.6 Device Operation

1. For output, a buffer containing the characters is placed in the port's transmitq⁶.

⁶Because this is high importance and low bandwidth, the current Linux implementation polls for the buffer to be used, rather than waiting for an interrupt, simplifying the implementation significantly. However, for generic serial ports with the `O_NONBLOCK` flag set, the polling limitation is relaxed and the consumed buffers are freed upon the next write or poll call or when a port is closed or hot-unplugged.

2. When a buffer is used in the receiveq (signalled by an interrupt), the contents is the input to the port associated with the virtqueue for which the notification was received.
3. If the driver negotiated the VIRTIO_CONSOLE_F_SIZE feature, a configuration change interrupt indicates that the updated size can be read from the configuration fields. This size applies to port 0 only.
4. If the driver negotiated the VIRTIO_CONSOLE_F_MULTIPORT feature, active ports are announced by the device using the VIRTIO_CONSOLE_PORT_ADD control message. The same message is used for port hot-plug as well.

5.3.6.1 Driver Requirements: Device Operation

The driver **MUST NOT** put a device-readable in a receiveq. The driver **MUST NOT** put a device-writable buffer in a transmitq.

5.3.6.2 Multiport Device Operation

If the driver negotiated the VIRTIO_CONSOLE_F_MULTIPORT, the two control queues are used to manipulate the different console ports: the control receiveq for messages from the device to the driver, and the control sendq for driver-to-device messages. The layout of the control messages is:

```
struct virtio_console_control {
    le32 id; /* Port number */
    le16 event; /* The kind of control event */
    le16 value; /* Extra information for the event */
};
```

The values for *event* are:

VIRTIO_CONSOLE_DEVICE_READY (0) Sent by the driver at initialization to indicate that it is ready to receive control messages. A value of 1 indicates success, and 0 indicates failure. The port number *id* is unused.

VIRTIO_CONSOLE_DEVICE_ADD (1) Sent by the device, to create a new port. *value* is unused.

VIRTIO_CONSOLE_DEVICE_REMOVE (2) Sent by the device, to remove an existing port. *value* is unused.

VIRTIO_CONSOLE_PORT_READY (3) Sent by the driver in response to the device's VIRTIO_CONSOLE_PORT_ADD message, to indicate that the port is ready to be used. A *value* of 1 indicates success, and 0 indicates failure.

VIRTIO_CONSOLE_CONSOLE_PORT (4) Sent by the device to nominate a port as a console port. There MAY be more than one console port.

VIRTIO_CONSOLE_RESIZE (5) Sent by the device to indicate a console size change. *value* is unused. The buffer is followed by the number of columns and rows:

```
struct virtio_console_resize {
    le16 cols;
    le16 rows;
};
```

VIRTIO_CONSOLE_PORT_OPEN (6) This message is sent by both the device and the driver. *value* indicates the state: 0 (port closed) or 1 (port open). This allows for ports to be used directly by guest and host processes to communicate in an application-defined manner.

VIRTIO_CONSOLE_PORT_NAME (7) Sent by the device to give a tag to the port. This control command is immediately followed by the UTF-8 name of the port for identification within the guest (without a NUL terminator).

5.3.6.2.1 Device Requirements: Multiport Device Operation

The device MUST NOT specify a port which exists in a VIRTIO_CONSOLE_DEVICE_ADD message, nor a port which is equal or greater than *max_nr_ports*.

The device MUST NOT specify a port in VIRTIO_CONSOLE_DEVICE_REMOVE which has not been created with a previous VIRTIO_CONSOLE_DEVICE_ADD.

5.3.6.2.2 Driver Requirements: Multiport Device Operation

The driver MUST send a VIRTIO_CONSOLE_DEVICE_READY message if VIRTIO_CONSOLE_F_MULTIPORT is negotiated.

Upon receipt of a VIRTIO_CONSOLE_CONSOLE_PORT message, the driver SHOULD treat the port in a manner suitable for text console access and MUST respond with a VIRTIO_CONSOLE_PORT_OPEN message, which MUST have *value* set to 1.

5.3.6.3 Legacy Interface: Device Operation

When using the legacy interface, transitional devices and drivers MUST format the fields in struct *virtio_console_control* according to the native endian of the guest rather than (necessarily when not using the legacy interface) little-endian.

When using the legacy interface, the driver SHOULD ignore the used length values for the transmit queues and the control transmitq.

Note: Historically, some devices put the total descriptor length there, even though no data was actually written.

5.3.6.4 Legacy Interface: Framing Requirements

When using legacy interfaces, transitional drivers which have not negotiated VIRTIO_F_ANY_LAYOUT MUST use only a single descriptor for all buffers in the control receiveq and control transmitq.

5.4 Entropy Device

The virtio entropy device supplies high-quality randomness for guest use.

5.4.1 Device ID

4

5.4.2 Virtqueues

0 requestq

5.4.3 Feature bits

None currently defined

5.4.4 Device configuration layout

None currently defined.

5.4.5 Device Initialization

1. The virtqueue is initialized

5.4.6 Device Operation

When the driver requires random bytes, it places the descriptor of one or more buffers in the queue. It will be completely filled by random data by the device.

5.4.6.1 Driver Requirements: Device Operation

The driver **MUST NOT** place driver-readable buffers into the queue.

The driver **MUST** examine the length written by the device to determine how many random bytes were received.

5.4.6.2 Device Requirements: Device Operation

The device **MUST** place one or more random bytes into the buffer, but it **MAY** use less than the entire buffer length.

5.5 Traditional Memory Balloon Device

This is the traditional balloon device. The device number 13 is reserved for a new memory balloon interface, with different semantics, which is expected in a future version of the standard.

The traditional virtio memory balloon device is a primitive device for managing guest memory: the device asks for a certain amount of memory, and the driver supplies it (or withdraws it, if the device has more than it asks for). This allows the guest to adapt to changes in allowance of underlying physical memory. If the feature is negotiated, the device can also be used to communicate guest memory statistics to the host.

5.5.1 Device ID

5

5.5.2 Virtqueues

0 inflateq

1 deflateq

2 statsq.

Virtqueue 2 only exists if VIRTIO_BALLOON_F_STATS_VQ set.

5.5.3 Feature bits

VIRTIO_BALLOON_F_MUST_TELL_HOST (0) Host has to be told before pages from the balloon are used.

VIRTIO_BALLOON_F_STATS_VQ (1) A virtqueue for reporting guest memory statistics is present.

VIRTIO_BALLOON_F_DEFLATE_ON_OOM (2) Deflate balloon on guest out of memory condition.

5.5.3.1 Driver Requirements: Feature bits

The driver SHOULD accept the VIRTIO_BALLOON_F_MUST_TELL_HOST feature if offered by the device.

5.5.3.2 Device Requirements: Feature bits

If the device offers the VIRTIO_BALLOON_F_MUST_TELL_HOST feature bit, and if the driver did not accept this feature bit, the device MAY signal failure by failing to set FEATURES_OK *device status* bit when the driver writes it.

5.5.3.2.0.1 Legacy Interface: Feature bits

As the legacy interface does not have a way to gracefully report feature negotiation failure, when using the legacy interface, transitional devices MUST support guests which do not negotiate VIRTIO_BALLOON_F_MUST_TELL_HOST feature, and SHOULD allow guest to use memory before notifying host if VIRTIO_BALLOON_F_MUST_TELL_HOST is not negotiated.

5.5.4 Device configuration layout

Both fields of this configuration are always available.

```
struct virtio_balloon_config {
    le32 num_pages;
    le32 actual;
};
```

5.5.4.0.0.1 Legacy Interface: Device configuration layout

When using the legacy interface, transitional devices and drivers MUST format the fields in struct virtio_balloon_config according to the little-endian format.

Note: This is unlike the usual convention that legacy device fields are guest endian.

5.5.5 Device Initialization

The device initialization process is outlined below:

1. The inflate and deflate virtqueues are identified.
2. If the VIRTIO_BALLOON_F_STATS_VQ feature bit is negotiated:
 - (a) Identify the stats virtqueue.
 - (b) Add one empty buffer to the stats virtqueue.
 - (c) DRIVER_OK is set: device operation begins.
 - (d) Notify the device about the stats virtqueue buffer.

5.5.6 Device Operation

The device is driven either by the receipt of a configuration change interrupt, or by changing guest memory needs, such as performing memory compaction or responding to out of memory conditions.

1. *num_pages* configuration field is examined. If this is greater than the *actual* number of pages, the balloon wants more memory from the guest. If it is less than *actual*, the balloon doesn't need it all.
2. To supply memory to the balloon (aka. inflate):
 - (a) The driver constructs an array of addresses of unused memory pages. These addresses are divided by 4096⁷ and the descriptor describing the resulting 32-bit array is added to the inflateq.
3. To remove memory from the balloon (aka. deflate):
 - (a) The driver constructs an array of addresses of memory pages it has previously given to the balloon, as described above. This descriptor is added to the deflateq.
 - (b) If the VIRTIO_BALLOON_F_MUST_TELL_HOST feature is negotiated, the guest informs the device of pages before it uses them.
 - (c) Otherwise, the guest is allowed to re-use pages previously given to the balloon before the device has acknowledged their withdrawal⁸.
4. In either case, the device acknowledges inflate and deflate requests by using the descriptor.
5. Once the device has acknowledged the inflation or deflation, the driver updates *actual* to reflect the new number of pages in the balloon.

5.5.6.1 Driver Requirements: Device Operation

The driver SHOULD supply pages to the balloon when *num_pages* is greater than the actual number of pages in the balloon.

The driver MAY use pages from the balloon when *num_pages* is less than the actual number of pages in the balloon.

The driver MAY supply pages to the balloon when *num_pages* is greater than or equal to the actual number of pages in the balloon.

If VIRTIO_BALLOON_F_DEFLATE_ON_OOM has not been negotiated, the driver MUST NOT use pages from the balloon when *num_pages* is less than or equal to the actual number of pages in the balloon.

If VIRTIO_BALLOON_F_DEFLATE_ON_OOM has been negotiated, the driver MAY use pages from the balloon when *num_pages* is less than or equal to the actual number of pages in the balloon if this is required for system stability (e.g. if memory is required by applications running within the guest).

The driver MUST use the deflateq to inform the device of pages that it wants to use from the balloon.

If the VIRTIO_BALLOON_F_MUST_TELL_HOST feature is negotiated, the driver MUST NOT use pages from the balloon until the device has acknowledged the deflate request.

Otherwise, if the VIRTIO_BALLOON_F_MUST_TELL_HOST feature is not negotiated, the driver MAY begin to re-use pages previously given to the balloon before the device has acknowledged the deflate request.

In any case, the driver MUST NOT use pages from the balloon after adding the pages to the balloon, but before the device has acknowledged the inflate request.

The driver MUST NOT request deflation of pages in the balloon before the device has acknowledged the inflate request.

The driver MUST update *actual* after changing the number of pages in the balloon.

The driver MAY update *actual* once after multiple inflate and deflate operations.

⁷This is historical, and independent of the guest page size.

⁸In this case, deflation advice is merely a courtesy.

5.5.6.2 Device Requirements: Device Operation

The device MAY modify the contents of a page in the balloon after detecting its physical number in an inflate request and before acknowledging the inflate request by using the inflateq descriptor.

If the VIRTIO_BALLOON_F_MUST_TELL_HOST feature is negotiated, the device MAY modify the contents of a page in the balloon after detecting its physical number in an inflate request and before detecting its physical number in a deflate request and acknowledging the deflate request.

5.5.6.2.1 Legacy Interface: Device Operation

When using the legacy interface, the driver SHOULD ignore the used length values.

Note: Historically, some devices put the total descriptor length there, even though no data was actually written.

When using the legacy interface, the driver MUST write out all 4 bytes each time it updates the *actual* value in the configuration space, using a single atomic operation.

When using the legacy interface, the device SHOULD NOT use the *actual* value written by the driver in the configuration space, until the last, most-significant byte of the value has been written.

Note: Historically, devices used the *actual* value, even though when using Virtio Over PCI Bus the device-specific configuration space was not guaranteed to be atomic. Using intermediate values during update by driver is best avoided, except for debugging.

Historically, drivers using Virtio Over PCI Bus wrote the *actual* value by using multiple single-byte writes in order, from the least-significant to the most-significant value.

5.5.6.3 Memory Statistics

The stats virtqueue is atypical because communication is driven by the device (not the driver). The channel becomes active at driver initialization time when the driver adds an empty buffer and notifies the device. A request for memory statistics proceeds as follows:

1. The device uses the buffer and sends an interrupt.
2. The driver pops the used buffer and discards it.
3. The driver collects memory statistics and writes them into a new buffer.
4. The driver adds the buffer to the virtqueue and notifies the device.
5. The device pops the buffer (retaining it to initiate a subsequent request) and consumes the statistics.

Within the buffer, statistics are an array of 6-byte entries. Each statistic consists of a 16 bit tag and a 64 bit value. All statistics are optional and the driver chooses which ones to supply. To guarantee backwards compatibility, devices omit unsupported statistics.

```
struct virtio_balloon_stat {
#define VIRTIO_BALLOON_S_SWAP_IN 0
#define VIRTIO_BALLOON_S_SWAP_OUT 1
#define VIRTIO_BALLOON_S_MAJFLT 2
#define VIRTIO_BALLOON_S_MINFLT 3
#define VIRTIO_BALLOON_S_MEMFREE 4
#define VIRTIO_BALLOON_S_MEMTOT 5
    le16 tag;
    le64 val;
} __attribute__((packed));
```


5.5.6.3.1 Driver Requirements: Memory Statistics

Normative statements in this section apply if and only if the `VIRTIO_BALLOON_F_STATS_VQ` feature has been negotiated.

The driver **MUST** make at most one buffer available to the device in the statsq, at all times.

After initializing the device, the driver **MUST** make an output buffer available in the statsq.

Upon detecting that device has used a buffer in the statsq, the driver **MUST** make an output buffer available in the statsq.

Before making an output buffer available in the statsq, the driver **MUST** initialize it, including one struct `virtio_balloon_stat` entry for each statistic that it supports.

Driver **MUST** use an output buffer size which is a multiple of 6 bytes for all buffers submitted to the statsq.

Driver **MAY** supply struct `virtio_balloon_stat` entries in the output buffer submitted to the statsq in any order, without regard to *tag* values.

Driver **MAY** supply a subset of all statistics in the output buffer submitted to the statsq.

Driver **MUST** supply the same subset of statistics in all buffers submitted to the statsq.

5.5.6.3.2 Device Requirements: Memory Statistics

Normative statements in this section apply if and only if the `VIRTIO_BALLOON_F_STATS_VQ` feature has been negotiated.

Within an output buffer submitted to the statsq, the device **MUST** ignore entries with *tag* values that it does not recognize.

Within an output buffer submitted to the statsq, the device **MUST** accept struct `virtio_balloon_stat` entries in any order without regard to *tag* values.

5.5.6.3.3 Legacy Interface: Memory Statistics

When using the legacy interface, transitional devices and drivers **MUST** format the fields in struct `virtio_balloon_stat` according to the native endian of the guest rather than (necessarily when not using the legacy interface) little-endian.

When using the legacy interface, the device **SHOULD** ignore all values in the first buffer in the statsq supplied by the driver after device initialization.

Note: Historically, drivers supplied an uninitialized buffer in the first buffer.

5.5.6.4 Memory Statistics Tags

VIRTIO_BALLOON_S_SWAP_IN (0) The amount of memory that has been swapped in (in bytes).

VIRTIO_BALLOON_S_SWAP_OUT (1) The amount of memory that has been swapped out to disk (in bytes).

VIRTIO_BALLOON_S_MAJFLT (2) The number of major page faults that have occurred.

VIRTIO_BALLOON_S_MINFLT (3) The number of minor page faults that have occurred.

VIRTIO_BALLOON_S_MEMFREE (4) The amount of memory not being used for any purpose (in bytes).

VIRTIO_BALLOON_S_MEMTOT (5) The total amount of memory available (in bytes).

5.6 SCSI Host Device

The virtio SCSI host device groups together one or more virtual logical units (such as disks), and allows communicating to them using the SCSI protocol. An instance of the device represents a SCSI host to which many targets and LUNs are attached.

The virtio SCSI device services two kinds of requests:

- command requests for a logical unit;
- task management functions related to a logical unit, target or command.

The device is also able to send out notifications about added and removed logical units. Together, these capabilities provide a SCSI transport protocol that uses virtqueues as the transfer medium. In the transport protocol, the virtio driver acts as the initiator, while the virtio SCSI host provides one or more targets that receive and process the requests.

This section relies on definitions from [SAM](#).

5.6.1 Device ID

8

5.6.2 Virtqueues

0 controlq

1 eventq

2...n request queues

5.6.3 Feature bits

VIRTIO SCSI_F_INOUT (0) A single request can include both device-readable and device-writable data buffers.

VIRTIO SCSI_F_HOTPLUG (1) The host SHOULD enable reporting of hot-plug and hot-unplug events for LUNs and targets on the SCSI bus. The guest SHOULD handle hot-plug and hot-unplug events.

VIRTIO SCSI_F_CHANGE (2) The host will report changes to LUN parameters via a VIRTIO SCSI_T_-PARAM_CHANGE event; the guest SHOULD handle them.

VIRTIO SCSI_F_T10_PI (3) The extended fields for T10 protection information (DIF/DIX) are included in the SCSI request header.

5.6.4 Device configuration layout

All fields of this configuration are always available.

```
struct virtio_scsi_config {
    le32 num_queues;
    le32 seg_max;
    le32 max_sectors;
    le32 cmd_per_lun;
    le32 event_info_size;
    le32 sense_size;
    le32 cdb_size;
    le16 max_channel;
    le16 max_target;
    le32 max_lun;
};
```

num_queues is the total number of request virtqueues exposed by the device. The driver MAY use only one request queue, or it can use more to achieve better performance.

seg_max is the maximum number of segments that can be in a command. A bidirectional command can include **seg_max** input segments and **seg_max** output segments.

max_sectors is a hint to the driver about the maximum transfer size to use.

cmd_per_lun is tells the driver the maximum number of linked commands it can send to one LUN.

event_info_size is the maximum size that the device will fill for buffers that the driver places in the eventq. It is written by the device depending on the set of negotiated features.

sense_size is the maximum size of the sense data that the device will write. The default value is written by the device and MUST be 96, but the driver can modify it. It is restored to the default when the device is reset.

cdb_size is the maximum size of the CDB that the driver will write. The default value is written by the device and MUST be 32, but the driver can likewise modify it. It is restored to the default when the device is reset.

max_channel, **max_target** and **max_lun** can be used by the driver as hints to constrain scanning the logical units on the host to channel/target/logical unit numbers that are less than or equal to the value of the fields. **max_channel** SHOULD be zero. **max_target** SHOULD be less than or equal to 255. **max_lun** SHOULD be less than or equal to 16383.

5.6.4.1 Driver Requirements: Device configuration layout

The driver MUST NOT write to device configuration fields other than **sense_size** and **cdb_size**.

The driver MUST NOT send more than **cmd_per_lun** linked commands to one LUN, and MUST NOT send more than the virtqueue size number of linked commands to one LUN.

5.6.4.2 Device Requirements: Device configuration layout

On reset, the device MUST set **sense_size** to 96 and **cdb_size** to 32.

5.6.4.3 Legacy Interface: Device configuration layout

When using the legacy interface, transitional devices and drivers MUST format the fields in struct **virtio_scsi_config** according to the native endian of the guest rather than (necessarily when not using the legacy interface) little-endian.

5.6.5 Device Requirements: Device Initialization

On initialization the driver SHOULD first discover the device's virtqueues.

If the driver uses the eventq, the driver SHOULD place at least one buffer in the eventq.

The driver MAY immediately issue requests⁹ or task management functions¹⁰.

5.6.6 Device Operation

Device operation consists of operating request queues, the control queue and the event queue.

⁹For example, INQUIRY or REPORT LUNS.

¹⁰For example, I_T RESET.

5.6.6.0.1 Legacy Interface: Device Operation

When using the legacy interface, the driver SHOULD ignore the used length values.

Note: Historically, devices put the total descriptor length, or the total length of device-writable buffers there, even when only part of the buffers were actually written.

5.6.6.1 Device Operation: Request Queues

The driver queues requests to an arbitrary request queue, and they are used by the device on that same queue. It is the responsibility of the driver to ensure strict request ordering for commands placed on different queues, because they will be consumed with no order constraints.

Requests have the following format:

```
struct virtio_scsi_req_cmd {
    // Device-readable part
    u8 lun[8];
    le64 id;
    u8 task_attr;
    u8 prio;
    u8 crn;
    u8 cdb[cdb_size];
    // The next two fields are only present if VIRTIO SCSI_F_T10_PI
    // is negotiated.
    le32 pi_bytesout;
    le32 pi_bytesin;
    u8 pi_out[pi_bytesout];
    u8 dataout[];

    // Device-writable part
    le32 sense_len;
    le32 residual;
    le16 status_qualifier;
    u8 status;
    u8 response;
    u8 sense[sense_size];
    // The next two fields are only present if VIRTIO SCSI_F_T10_PI
    // is negotiated
    u8 pi_in[pi_bytesin];
    u8 datain[];
};

/* command-specific response values */
#define VIRTIO_SCSI_S_OK 0
#define VIRTIO_SCSI_S_OVERRUN 1
#define VIRTIO_SCSI_S_ABORTED 2
#define VIRTIO_SCSI_S_BAD_TARGET 3
#define VIRTIO_SCSI_S_RESET 4
#define VIRTIO_SCSI_S_BUSY 5
#define VIRTIO_SCSI_S_TRANSPORT_FAILURE 6
#define VIRTIO_SCSI_S_TARGET_FAILURE 7
#define VIRTIO_SCSI_S_NEXUS_FAILURE 8
#define VIRTIO_SCSI_S_FAILURE 9

/* task_attr */
#define VIRTIO_SCSI_S_SIMPLE 0
#define VIRTIO_SCSI_S_ORDERED 1
#define VIRTIO_SCSI_S_HEAD 2
#define VIRTIO_SCSI_S_ACA 3
```

lun addresses the REPORT LUNS well-known logical unit, or a target and logical unit in the virtio-scsi device's SCSI domain. When used to address the REPORT LUNS logical unit, *lun* is 0xC1, 0x01 and six zero bytes. The virtio-scsi device SHOULD implement the REPORT LUNS well-known logical unit.

When used to address a target and logical unit, the only supported format for *lun* is: first byte set to 1,

second byte set to target, third and fourth byte representing a single level LUN structure, followed by four zero bytes. With this representation, a virtio-scsi device can serve up to 256 targets and 16384 LUNs per target. The device MAY also support having a well-known logical units in the third and fourth byte.

id is the command identifier (“tag”).

task_attr defines the task attribute as in the table above, but all task attributes MAY be mapped to SIMPLE by the device. Some commands are defined by SCSI standards as “implicit head of queue”; for such commands, all task attributes MAY also be mapped to HEAD OF QUEUE. Drivers and applications SHOULD NOT send a command with the ORDERED task attribute if the command has an implicit HEAD OF QUEUE attribute, because whether the ORDERED task attribute is honored is vendor-specific.

crn may also be provided by clients, but is generally expected to be 0. The maximum CRN value defined by the protocol is 255, since CRN is stored in an 8-bit integer.

The CDB is included in *cdb* and its size, *cdb_size*, is taken from the configuration space.

All of these fields are defined in [SAM](#) and are always device-readable.

pi_bytesout determines the size of the *pi_out* field in bytes. If it is nonzero, the *pi_out* field contains outgoing protection information for write operations. *pi_bytesin* determines the size of the *pi_in* field in the device-writable section, in bytes. All three fields are only present if VIRTIO_SCSI_F_T10_PI has been negotiated.

The remainder of the device-readable part is the data output buffer, *dataout*.

sense and subsequent fields are always device-writable. *sense_len* indicates the number of bytes actually written to the sense buffer.

residual indicates the residual size, calculated as “data_length - number_of_transferred_bytes”, for read or write operations. For bidirectional commands, the number_of_transferred_bytes includes both read and written bytes. A *residual* that is less than the size of *datain* means that *dataout* was processed entirely. A *residual* that exceeds the size of *datain* means that *dataout* was processed partially and *datain* was not processed at all.

If the *pi_bytesin* is nonzero, the *pi_in* field contains incoming protection information for read operations. *pi_in* is only present if VIRTIO_SCSI_F_T10_PI has been negotiated¹¹.

The remainder of the device-writable part is the data input buffer, *datain*.

5.6.6.1.1 Device Requirements: Device Operation: Request Queues

The device MUST write the *status* byte as the status code as defined in [SAM](#).

The device MUST write the *response* byte as one of the following:

VIRTIO_SCSI_S_OK when the request was completed and the *status* byte is filled with a SCSI status code (not necessarily “GOOD”).

VIRTIO_SCSI_S_OVERRUN if the content of the CDB (such as the allocation length, parameter length or transfer size) requires more data than is available in the *datain* and *dataout* buffers.

VIRTIO_SCSI_S_ABORTED if the request was cancelled due to an ABORT TASK or ABORT TASK SET task management function.

VIRTIO_SCSI_S_BAD_TARGET if the request was never processed because the target indicated by *lun* does not exist.

VIRTIO_SCSI_S_RESET if the request was cancelled due to a bus or device reset (including a task management function).

VIRTIO_SCSI_S_TRANSPORT_FAILURE if the request failed due to a problem in the connection between the host and the target (severed link).

¹¹There is no separate residual size for *pi_bytesout* and *pi_bytesin*. It can be computed from the *residual* field, the size of the data integrity information per sector, and the sizes of *pi_out*, *pi_in*, *dataout* and *datain*.

VIRTIO_SCSI_S_TARGET_FAILURE if the target is suffering a failure and to tell the driver not to retry on other paths.

VIRTIO_SCSI_S_NEXUS_FAILURE if the nexus is suffering a failure but retrying on other paths might yield a different result.

VIRTIO_SCSI_S_BUSY if the request failed but retrying on the same path is likely to work.

VIRTIO_SCSI_S_FAILURE for other host or driver error. In particular, if neither *dataout* nor *datain* is empty, and the VIRTIO_SCSI_F_INOUT feature has not been negotiated, the request will be immediately returned with a response equal to VIRTIO_SCSI_S_FAILURE.

All commands must be completed before the virtio-scsi device is reset or unplugged. The device MAY choose to abort them, or if it does not do so MUST pick the VIRTIO_SCSI_S_FAILURE response.

5.6.6.1.2 Driver Requirements: Device Operation: Request Queues

task_attr, *prio* and *crn* SHOULD be zero.

Upon receiving a VIRTIO_SCSI_S_TARGET_FAILURE response, the driver SHOULD NOT retry the request on other paths.

5.6.6.1.3 Legacy Interface: Device Operation: Request Queues

When using the legacy interface, transitional devices and drivers MUST format the fields in struct *virtio_scsi_req_cmd* according to the native endian of the guest rather than (necessarily when not using the legacy interface) little-endian.

5.6.6.2 Device Operation: controlq

The controlq is used for other SCSI transport operations. Requests have the following format:

```
struct virtio_scsi_ctrl {
    le32 type;
    ...
    u8 response;
};

/* response values valid for all commands */
#define VIRTIO_SCSI_S_OK 0
#define VIRTIO_SCSI_S_BAD_TARGET 3
#define VIRTIO_SCSI_S_BUSY 5
#define VIRTIO_SCSI_S_TRANSPORT_FAILURE 6
#define VIRTIO_SCSI_S_TARGET_FAILURE 7
#define VIRTIO_SCSI_S_NEXUS_FAILURE 8
#define VIRTIO_SCSI_S_FAILURE 9
#define VIRTIO_SCSI_S_INCORRECT_LUN 12
```

The *type* identifies the remaining fields.

The following commands are defined:

- Task management function.

```
#define VIRTIO_SCSI_T_TMF 0

#define VIRTIO_SCSI_T_TMF_ABORT_TASK 0
#define VIRTIO_SCSI_T_TMF_ABORT_TASK_SET 1
#define VIRTIO_SCSI_T_TMF_CLEAR_ACA 2
#define VIRTIO_SCSI_T_TMF_CLEAR_TASK_SET 3
#define VIRTIO_SCSI_T_TMF_I_T_NEXUS_RESET 4
#define VIRTIO_SCSI_T_TMF_LOGICAL_UNIT_RESET 5
#define VIRTIO_SCSI_T_TMF_QUERY_TASK 6
#define VIRTIO_SCSI_T_TMF_QUERY_TASK_SET 7
```

```

struct virtio_scsi_ctrl_tmf
{
    // Device-readable part
    le32 type;
    le32 subtype;
    u8   lun[8];
    le64 id;
    // Device-writable part
    u8   response;
}

/* command-specific response values */
#define VIRTIO_SCSI_S_FUNCTION_COMPLETE      0
#define VIRTIO_SCSI_S_FUNCTION_SUCCEEDED    10
#define VIRTIO_SCSI_S_FUNCTION_REJECTED     11

```

The *type* is `VIRTIO_SCSI_T_TMF`; *subtype* defines which task management function. All fields except *response* are filled by the driver.

Other fields which are irrelevant for the requested TMF are ignored but they are still present. *lun* is in the same format specified for request queues; the single level LUN is ignored when the task management function addresses a whole I_T nexus. When relevant, the value of *id* is matched against the id values passed on the requestq.

The outcome of the task management function is written by the device in *response*. The command-specific response values map 1-to-1 with those defined in [SAM](#).

Task management function can affect the response value for commands that are in the request queue and have not been completed yet. For example, the device MUST complete all active commands on a logical unit or target (possibly with a `VIRTIO_SCSI_S_RESET` response code) upon receiving a "logical unit reset" or "I_T nexus reset" TMF. Similarly, the device MUST complete the selected commands (possibly with a `VIRTIO_SCSI_S_ABORTED` response code) upon receiving an "abort task" or "abort task set" TMF. Such effects MUST take place before the TMF itself is successfully completed, and the device MUST use memory barriers appropriately in order to ensure that the driver sees these writes in the correct order.

- Asynchronous notification query.

```

#define VIRTIO_SCSI_T_AN_QUERY      1

struct virtio_scsi_ctrl_an {
    // Device-readable part
    le32 type;
    u8   lun[8];
    le32 event_requested;
    // Device-writable part
    le32 event_actual;
    u8   response;
}

#define VIRTIO_SCSI_EVT_ASYNC_OPERATIONAL_CHANGE  2
#define VIRTIO_SCSI_EVT_ASYNC_POWER_MGMT         4
#define VIRTIO_SCSI_EVT_ASYNC_EXTERNAL_REQUEST   8
#define VIRTIO_SCSI_EVT_ASYNC_MEDIA_CHANGE       16
#define VIRTIO_SCSI_EVT_ASYNC_MULTI_HOST         32
#define VIRTIO_SCSI_EVT_ASYNC_DEVICE_BUSY        64

```

By sending this command, the driver asks the device which events the given LUN can report, as described in paragraphs 6.6 and A.6 of [SCSI MMC](#). The driver writes the events it is interested in into *event_requested*; the device responds by writing the events that it supports into *event_actual*.

The *type* is `VIRTIO_SCSI_T_AN_QUERY`. *lun* and *event_requested* are written by the driver. *event_actual* and *response* fields are written by the device.

No command-specific values are defined for the *response* byte.

- Asynchronous notification subscription.

```
#define VIRTIO SCSI_T_AN_SUBSCRIBE 2

struct virtio_scsi_ctrl_an {
    // Device-readable part
    le32 type;
    u8 lun[8];
    le32 event_requested;
    // Device-writable part
    le32 event_actual;
    u8 response;
}
```

By sending this command, the driver asks the specified LUN to report events for its physical interface, again as described in [SCSI MMC](#). The driver writes the events it is interested in into *event_requested*; the device responds by writing the events that it supports into *event_actual*.

Event types are the same as for the asynchronous notification query message.

The *type* is VIRTIO SCSI_T_AN_SUBSCRIBE. *lun* and *event_requested* are written by the driver. *event_actual* and *response* are written by the device.

No command-specific values are defined for the response byte.

5.6.6.2.1 Legacy Interface: Device Operation: controlq

When using the legacy interface, transitional devices and drivers MUST format the fields in struct *virtio_scsi_ctrl*, struct *virtio_scsi_ctrl_tmf*, struct *virtio_scsi_ctrl_an* and struct *virtio_scsi_ctrl_an* according to the native endian of the guest rather than (necessarily when not using the legacy interface) little-endian.

5.6.6.3 Device Operation: eventq

The eventq is populated by the driver for the device to report information on logical units that are attached to it. In general, the device will not queue events to cope with an empty eventq, and will end up dropping events if it finds no buffer ready. However, when reporting events for many LUNs (e.g. when a whole target disappears), the device can throttle events to avoid dropping them. For this reason, placing 10-15 buffers on the event queue is sufficient.

Buffers returned by the device on the eventq will be referred to as “events” in the rest of this section. Events have the following format:

```
#define VIRTIO SCSI_T_EVENTS_MISSED 0x80000000

struct virtio_scsi_event {
    // Device-writable part
    le32 event;
    u8 lun[8];
    le32 reason;
}
```

The device sets bit 31 in *event* to report lost events due to missing buffers.

The meaning of *reason* depends on the contents of *event*. The following events are defined:

- No event.

```
#define VIRTIO SCSI_T_NO_EVENT 0
```

This event is fired in the following cases:

- When the device detects in the eventq a buffer that is shorter than what is indicated in the configuration field, it MAY use it immediately and put this dummy value in *event*. A well-written driver will never observe this situation.

- When events are dropped, the device MAY signal this event as soon as the drivers makes a buffer available, in order to request action from the driver. In this case, of course, this event will be reported with the `VIRTIO_SCSI_T_EVENTS_MISSED` flag.

- Transport reset

```
#define VIRTIO_SCSI_T_TRANSPORT_RESET 1

#define VIRTIO_SCSI_EVT_RESET_HARD      0
#define VIRTIO_SCSI_EVT_RESET_RESCAN   1
#define VIRTIO_SCSI_EVT_RESET_REMOVED   2
```

By sending this event, the device signals that a logical unit on a target has been reset, including the case of a new device appearing or disappearing on the bus. The device fills in all fields. *event* is set to `VIRTIO_SCSI_T_TRANSPORT_RESET`. *lun* addresses a logical unit in the SCSI host.

The *reason* value is one of the three `#define` values appearing above:

VIRTIO_SCSI_EVT_RESET_REMOVED (“LUN/target removed”) is used if the target or logical unit is no longer able to receive commands.

VIRTIO_SCSI_EVT_RESET_HARD (“LUN hard reset”) is used if the logical unit has been reset, but is still present.

VIRTIO_SCSI_EVT_RESET_RESCAN (“rescan LUN/target”) is used if a target or logical unit has just appeared on the device.

The “removed” and “rescan” events can happen when `VIRTIO_SCSI_F_HOTPLUG` feature was negotiated; when sent for LUN 0, they MAY apply to the entire target so the driver can ask the initiator to rescan the target to detect this.

Events will also be reported via sense codes (this obviously does not apply to newly appeared buses or targets, since the application has never discovered them):

- “LUN/target removed” maps to sense key `ILLEGAL REQUEST`, `asc 0x25`, `ascq 0x00` (`LOGICAL UNIT NOT SUPPORTED`)
- “LUN hard reset” maps to sense key `UNIT ATTENTION`, `asc 0x29` (`POWER ON, RESET OR BUS DEVICE RESET OCCURRED`)
- “rescan LUN/target” maps to sense key `UNIT ATTENTION`, `asc 0x3f`, `ascq 0x0e` (`REPORTED LUNS DATA HAS CHANGED`)

The preferred way to detect transport reset is always to use events, because sense codes are only seen by the driver when it sends a SCSI command to the logical unit or target. However, in case events are dropped, the initiator will still be able to synchronize with the actual state of the controller if the driver asks the initiator to rescan of the SCSI bus. During the rescan, the initiator will be able to observe the above sense codes, and it will process them as if it the driver had received the equivalent event.

- Asynchronous notification

```
#define VIRTIO_SCSI_T_ASYNC_NOTIFY 2
```

By sending this event, the device signals that an asynchronous event was fired from a physical interface.

All fields are written by the device. *event* is set to `VIRTIO_SCSI_T_ASYNC_NOTIFY`. *lun* addresses a logical unit in the SCSI host. *reason* is a subset of the events that the driver has subscribed to via the “Asynchronous notification subscription” command.

- LUN parameter change

```
#define VIRTIO_SCSI_T_PARAM_CHANGE 3
```

By sending this event, the device signals a change in the configuration parameters of a logical unit, for example the capacity or cache mode. *event* is set to VIRTIO_SCSI_T_PARAM_CHANGE. *lun* addresses a logical unit in the SCSI host.

The same event SHOULD also be reported as a unit attention condition. *reason* contains the additional sense code and additional sense code qualifier, respectively in bits 0...7 and 8...15.

Note: For example, a change in capacity will be reported as asc 0x2a, ascq 0x09 (CAPACITY DATA HAS CHANGED).

For MMC devices (inquiry type 5) there would be some overlap between this event and the asynchronous notification event, so for simplicity the host never reports this event for MMC devices.

5.6.6.3.1 Driver Requirements: Device Operation: eventq

The driver SHOULD keep the eventq populated with buffers. These buffers MUST be device-writable, and SHOULD be at least *event_info_size* bytes long, and MUST be at least the size of struct virtio_scsi_event.

If *event* has bit 31 set, the driver SHOULD poll the logical units for unit attention conditions, and/or do whatever form of bus scan is appropriate for the guest operating system and SHOULD poll for asynchronous events manually using SCSI commands.

When receiving a VIRTIO_SCSI_T_TRANSPORT_RESET message with *reason* set to VIRTIO_SCSI_EVT_RESET_REMOVED or VIRTIO_SCSI_EVT_RESET_RESCAN for LUN 0, the driver SHOULD ask the initiator to rescan the target, in order to detect the case when an entire target has appeared or disappeared.

5.6.6.3.2 Device Requirements: Device Operation: eventq

The device MUST set bit 31 in *event* if events were lost due to missing buffers, and it MAY use a VIRTIO_SCSI_T_NO_EVENT event to report this.

The device MUST NOT send VIRTIO_SCSI_T_TRANSPORT_RESET messages with *reason* set to VIRTIO_SCSI_EVT_RESET_REMOVED or VIRTIO_SCSI_EVT_RESET_RESCAN unless VIRTIO_SCSI_F_HOTPLUG was negotiated.

The device MUST NOT report VIRTIO_SCSI_T_PARAM_CHANGE for MMC devices.

5.6.6.3.3 Legacy Interface: Device Operation: eventq

When using the legacy interface, transitional devices and drivers MUST format the fields in struct virtio_scsi_event according to the native endian of the guest rather than (necessarily when not using the legacy interface) little-endian.

5.6.6.4 Legacy Interface: Framing Requirements

When using legacy interfaces, transitional drivers which have not negotiated VIRTIO_F_ANY_LAYOUT MUST use a single descriptor for the *lun*, *id*, *task_attr*, *prio*, *crn* and *cdb* fields, and MUST only use a single descriptor for the *sense_len*, *residual*, *status_qualifier*, *status*, *response* and *sense* fields.

6 Reserved Feature Bits

Currently these device-independent feature bits defined:

VIRTIO_F_RING_INDIRECT_DESC (28) Negotiating this feature indicates that the driver can use descriptors with the `VIRTQ_DESC_F_INDIRECT` flag set, as described in [2.5.5.3 Indirect Descriptors](#) and [2.6.7 Indirect Flag: Scatter-Gather Support](#).

VIRTIO_F_RING_EVENT_IDX(29) This feature enables the *used_event* and the *avail_event* fields as described in [2.5.7](#), [2.5.8](#) and [2.6.10](#).

VIRTIO_F_VERSION_1(32) This indicates compliance with this specification, giving a simple way to detect legacy devices or drivers.

VIRTIO_F_IOMMU_PLATFORM(33) This feature indicates that the device is behind an IOMMU that translates bus addresses from the device into physical addresses in memory. If this feature bit is set to 0, then the device emits physical addresses which are not translated further, even though an IOMMU may be present.

VIRTIO_F_RING_PACKED(34) This feature indicates support for the packed virtqueue layout as described in [2.6 Packed Virtqueues](#).

VIRTIO_F_IN_ORDER(35) This feature indicates that all buffers are used by the device in the same order in which they have been made available. ~~This feature indicates that drivers pass extra data (besides identifying the Virtqueue) in their device notifications. See ??.~~

6.1 Driver Requirements: Reserved Feature Bits

A driver **MUST** accept `VIRTIO_F_VERSION_1` if it is offered. A driver **MAY** fail to operate further if `VIRTIO_F_VERSION_1` is not offered.

A driver **SHOULD** accept `VIRTIO_F_IOMMU_PLATFORM` if it is offered, and it **MUST** then either disable the IOMMU or configure the IOMMU to translate bus addresses passed to the device into physical addresses in memory. If `VIRTIO_F_IOMMU_PLATFORM` is not offered, then a driver **MUST** pass only physical addresses to the device.

A driver **SHOULD** accept `VIRTIO_F_RING_PACKED` if it is offered.

6.2 Device Requirements: Reserved Feature Bits

A device **MUST** offer `VIRTIO_F_VERSION_1`. A device **MAY** fail to operate further if `VIRTIO_F_VERSION_1` is not accepted.

A device **SHOULD** offer `VIRTIO_F_IOMMU_PLATFORM` if it is behind an IOMMU that translates bus addresses from the device into physical addresses in memory. A device **MAY** fail to operate further if `VIRTIO_F_IOMMU_PLATFORM` is not accepted.

If `VIRTIO_F_IN_ORDER` has been negotiated, a device **MUST** use buffers in the same order in which they have been available.

6.3 Legacy Interface: Reserved Feature Bits

Transitional devices MAY offer the following:

VIRTIO_F_NOTIFY_ON_EMPTY (24) If this feature has been negotiated by driver, the device MUST issue an interrupt if the device runs out of available descriptors on a virtqueue, even though interrupts are suppressed using the `VIRTQ_AVAIL_F_NO_INTERRUPT` flag or the *used_event* field.

Note: An example of a driver using this feature is the legacy networking driver: it doesn't need to know every time a packet is transmitted, but it does need to free the transmitted packets a finite time after they are transmitted. It can avoid using a timer if the device interrupts it when all the packets are transmitted.

Transitional devices MUST offer, and if offered by the device transitional drivers MUST accept the following:

VIRTIO_F_ANY_LAYOUT (27) This feature indicates that the device accepts arbitrary descriptor layouts, as described in Section [2.5.4.3 Legacy Interface: Message Framing](#).

UNUSED (30) Bit 30 is used by qemu's implementation to check for experimental early versions of virtio which did not perform correct feature negotiation, and SHOULD NOT be negotiated.

7 Conformance

This chapter lists the conformance targets and clauses for each; this also forms a useful checklist which authors are asked to consult for their implementations!

7.1 Conformance Targets

Conformance targets:

Driver A driver MUST conform to three conformance clauses:

- Clause [7.2](#),
- One of clauses [7.2.1](#), [7.2.2](#) or [7.2.3](#).
- One of clauses [7.2.4](#), [7.2.5](#), [7.2.6](#), [7.2.7](#), [7.2.8](#) or [7.2.9](#).

Device A device MUST conform to three conformance clauses:

- Clause [7.3](#),
- One of clauses [7.3.1](#), [7.3.2](#) or [7.3.3](#).
- One of clauses [7.3.4](#), [7.3.5](#), [7.3.6](#), [7.3.7](#), [7.3.8](#) or [7.3.9](#).

7.2 Driver Conformance

A driver MUST conform to the following normative statements:

- [2.1.1](#)
- [2.2.1](#)
- [2.3.1](#)
- [2.5.1](#)
- [2.5.4.2](#)
- [2.5.5.2](#)
- [2.5.5.3.1](#)
- [2.5.7.1](#)
- [2.5.6.1](#)
- [2.5.8.3](#)
- [2.5.9.1](#)
- [2.5.12.3.1](#)
- [2.5.12.4.1](#)
- [3.1.1](#)
- [3.3.1](#)

- [6.1](#)

7.2.1 PCI Driver Conformance

A PCI driver MUST conform to the following normative statements:

- [4.1.2.2](#)
- [4.1.3.1](#)
- [4.1.4.1](#)
- [4.1.4.3.2](#)
- [4.1.4.5.2](#)
- [4.1.4.7.2](#)
- [4.1.5.1.2.2](#)
- [4.1.5.4.2](#)

7.2.2 MMIO Driver Conformance

An MMIO driver MUST conform to the following normative statements:

- [4.2.2.2](#)
- [4.2.3.1.1](#)
- [4.2.3.4.1](#)

7.2.3 Channel I/O Driver Conformance

A Channel I/O driver MUST conform to the following normative statements:

- [4.3.1.2](#)
- [4.3.2.1.2](#)
- [4.3.2.3.1](#)
- [4.3.3.1.2.2](#)
- [4.3.3.2.2](#)

7.2.4 Network Driver Conformance

A network driver MUST conform to the following normative statements:

- [5.1.4.2](#)
- [5.1.6.2.1](#)
- [5.1.6.3.1](#)
- [5.1.6.4.2](#)
- [5.1.6.5.1.2](#)
- [5.1.6.5.2.2](#)
- [5.1.6.5.4.1](#)
- [5.1.6.5.5.1](#)

- [5.1.6.5.6.2](#)

7.2.5 Block Driver Conformance

A block driver MUST conform to the following normative statements:

- [5.2.5.1](#)
- [5.2.6.1](#)

7.2.6 Console Driver Conformance

A console driver MUST conform to the following normative statements:

- [5.3.6.1](#)
- [5.3.6.2.2](#)

7.2.7 Entropy Driver Conformance

An entropy driver MUST conform to the following normative statements:

- [5.4.6.1](#)

7.2.8 Traditional Memory Balloon Driver Conformance

A traditional memory balloon driver MUST conform to the following normative statements:

- [5.5.3.1](#)
- [5.5.6.1](#)
- [5.5.6.3.1](#)

7.2.9 SCSI Host Driver Conformance

An SCSI host driver MUST conform to the following normative statements:

- [5.6.4.1](#)
- [5.6.6.1.2](#)
- [5.6.6.3.1](#)

7.3 Device Conformance

A device MUST conform to the following normative statements:

- [2.1.2](#)
- [2.2.2](#)
- [2.3.2](#)
- [2.5.4.1](#)
- [2.5.5.1](#)
- [2.5.5.3.2](#)

- [2.5.7.2](#)
- [2.5.8.2](#)
- [2.5.9.2](#)
- [6.2](#)

7.3.1 PCI Device Conformance

A PCI device MUST conform to the following normative statements:

- [4.1.1](#)
- [4.1.2.1](#)
- [4.1.3.2](#)
- [4.1.4.2](#)
- [4.1.4.3.1](#)
- [4.1.4.4.1](#)
- [4.1.4.5.1](#)
- [4.1.4.6.1](#)
- [4.1.4.7.1](#)
- [4.1.4.9.0.1](#)
- [4.1.5.1.2.1](#)
- [4.1.5.3.1](#)
- [4.1.5.4.1](#)

7.3.2 MMIO Device Conformance

An MMIO device MUST conform to the following normative statements:

- [4.2.2.1](#)

7.3.3 Channel I/O Device Conformance

A Channel I/O device MUST conform to the following normative statements:

- [4.3.1.1](#)
- [4.3.2.1.1](#)
- [4.3.2.2.1](#)
- [4.3.2.3.2](#)
- [4.3.2.6.3.1](#)
- [4.3.3.1.2.1](#)
- [4.3.3.2.1](#)

7.3.4 Network Device Conformance

A network device MUST conform to the following normative statements:

- [5.1.4.1](#)
- [5.1.6.2.2](#)
- [5.1.6.3.2](#)
- [5.1.6.4.1](#)
- [5.1.6.5.1.1](#)
- [5.1.6.5.2.1](#)
- [5.1.6.5.4.2](#)
- [5.1.6.5.5.2](#)

7.3.5 Block Device Conformance

A block device MUST conform to the following normative statements:

- [5.2.5.2](#)
- [5.2.6.2](#)

7.3.6 Console Device Conformance

A console device MUST conform to the following normative statements:

- [5.3.5.1](#)
- [5.3.6.2.1](#)

7.3.7 Entropy Device Conformance

An entropy device MUST conform to the following normative statements:

- [5.4.6.2](#)

7.3.8 Traditional Memory Balloon Device Conformance

A traditional memory balloon device MUST conform to the following normative statements:

- [5.5.3.2](#)
- [5.5.6.2](#)
- [5.5.6.3.2](#)

7.3.9 SCSI Host Device Conformance

An SCSI host device MUST conform to the following normative statements:

- [5.6.4.2](#)
- [5.6.5](#)
- [5.6.6.1.1](#)
- [5.6.6.3.2](#)

7.4 Legacy Interface: Transitional Device and Transitional Driver Conformance

A conformant implementation MUST be either transitional or non-transitional, see [1.3.1](#).

A non-transitional implementation conforms to this specification if it satisfies all of the MUST or REQUIRED level requirements defined above.

An implementation MAY choose to implement OPTIONAL support for the legacy interface, including support for legacy drivers or devices, by additionally conforming to all of the MUST or REQUIRED level requirements for the legacy interface for the transitional devices and drivers.

The requirements for the legacy interface for transitional implementations are located in sections named “Legacy Interface” listed below:

- Section [2.2.3](#)
- Section [2.3.3](#)
- Section [2.3.4](#)
- Section [2.5.2](#)
- Section [2.5.3](#)
- Section [2.5.4.3](#)
- Section [3.1.2](#)
- Section [4.1.2.3](#)
- Section [4.1.4.8](#)
- Section [4.1.5.1.1.1](#)
- Section [4.1.5.1.3.1](#)
- Section [4.2.4](#)
- Section [4.3.2.1.3](#)
- Section [4.3.2.2.2](#)
- Section [4.3.3.1.3](#)
- Section [4.3.2.6.4](#)
- Section [5.1.3.2](#)
- Section [5.1.4.3](#)
- Section [5.1.6.1](#)
- Section [5.1.6.5.2.3](#)
- Section [5.1.6.5.3.1](#)
- Section [5.1.6.5.5.3](#)
- Section [5.1.6.5.6.3](#)
- Section [5.2.3.1](#)
- Section [5.2.4.1](#)
- Section [5.2.5.3](#)
- Section [5.2.6.3](#)
- Section [5.3.4.1](#)
- Section [5.3.6.3](#)

- Section [5.5.3.2.0.1](#)
- Section [5.5.6.2.1](#)
- Section [5.5.6.3.3](#)
- Section [5.6.4.3](#)
- Section [5.6.6.0.1](#)
- Section [5.6.6.1.3](#)
- Section [5.6.6.2.1](#)
- Section [5.6.6.3.3](#)
- Section [6.3](#)

Appendix A. virtio_queue.h

This file is also available at the link http://docs.oasis-open.org/virtio/virtio/v1.0/wd10/listings/virtio_queue.h. All definitions in this section are for non-normative reference only.

```
#ifndef VIRTQUEUE_H
#define VIRTQUEUE_H
/* An interface for efficient virtio implementation.
 *
 * This header is BSD licensed so anyone can use the definitions
 * to implement compatible drivers/servers.
 *
 * Copyright 2007, 2009, IBM Corporation
 * Copyright 2011, Red Hat, Inc
 * All rights reserved.
 *
 * Redistribution and use in source and binary forms, with or without
 * modification, are permitted provided that the following conditions
 * are met:
 * 1. Redistributions of source code must retain the above copyright
 *    notice, this list of conditions and the following disclaimer.
 * 2. Redistributions in binary form must reproduce the above copyright
 *    notice, this list of conditions and the following disclaimer in the
 *    documentation and/or other materials provided with the distribution.
 * 3. Neither the name of IBM nor the names of its contributors
 *    may be used to endorse or promote products derived from this software
 *    without specific prior written permission.
 * THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS ``AS IS'' AND
 * ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE
 * IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE
 * ARE DISCLAIMED. IN NO EVENT SHALL IBM OR CONTRIBUTORS BE LIABLE
 * FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL
 * DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS
 * OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION)
 * HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT
 * LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY
 * OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF
 * SUCH DAMAGE.
 */
#include <stdint.h>

/* This marks a buffer as continuing via the next field. */
#define VIRTQ_DESC_F_NEXT 1
/* This marks a buffer as write-only (otherwise read-only). */
#define VIRTQ_DESC_F_WRITE 2
/* This means the buffer contains a list of buffer descriptors. */
#define VIRTQ_DESC_F_INDIRECT 4

/* The device uses this in used->flags to advise the driver: don't kick me
 * when you add a buffer. It's unreliable, so it's simply an
 * optimization. */
#define VIRTQ_USED_F_NO_NOTIFY 1
/* The driver uses this in avail->flags to advise the device: don't
 * interrupt me when you consume a buffer. It's unreliable, so it's
 * simply an optimization. */
#define VIRTQ_AVAIL_F_NO_INTERRUPT 1

/* Support for indirect descriptors */
#define VIRTIO_F_INDIRECT_DESC 28

/* Support for avail_event and used_event fields */
#define VIRTIO_F_EVENT_IDX 29
```

```

/* Arbitrary descriptor layouts. */
#define VIRTIO_F_ANY_LAYOUT 27

/* Virtqueue descriptors: 16 bytes.
 * These can chain together via "next". */
struct virtq_desc {
    /* Address (guest-physical). */
    le64 addr;
    /* Length. */
    le32 len;
    /* The flags as indicated above. */
    le16 flags;
    /* We chain unused descriptors via this, too */
    le16 next;
};

struct virtq_avail {
    le16 flags;
    le16 idx;
    le16 ring[];
    /* Only if VIRTIO_F_EVENT_IDX: le16 used_event; */
};

/* le32 is used here for ids for padding reasons. */
struct virtq_used_elem {
    /* Index of start of used descriptor chain. */
    le32 id;
    /* Total length of the descriptor chain which was written to. */
    le32 len;
};

struct virtq_used {
    le16 flags;
    le16 idx;
    struct virtq_used_elem ring[];
    /* Only if VIRTIO_F_EVENT_IDX: le16 avail_event; */
};

struct virtq {
    unsigned int num;

    struct virtq_desc *desc;
    struct virtq_avail *avail;
    struct virtq_used *used;
};

static inline int virtq_need_event(uint16_t event_idx, uint16_t new_idx, uint16_t old_idx)
{
    return (uint16_t)(new_idx - event_idx - 1) < (uint16_t)(new_idx - old_idx);
}

/* Get location of event indices (only with VIRTIO_F_EVENT_IDX) */
static inline le16 *virtq_used_event(struct virtq *vq)
{
    /* For backwards compat, used event index is at *end* of avail ring. */
    return &vq->avail->ring[vq->num];
}

static inline le16 *virtq_avail_event(struct virtq *vq)
{
    /* For backwards compat, avail event index is at *end* of used ring. */
    return (le16 *)&vq->used->ring[vq->num];
}
#endif /* VIRTQUEUE_H */

```

Appendix B. Creating New Device Types

Various considerations are necessary when creating a new device type.

B.1 How Many Virtqueues?

It is possible that a very simple device will operate entirely through its device configuration space, but most will need at least one virtqueue in which it will place requests. A device with both input and output (eg. console and network devices described here) need two queues: one which the driver fills with buffers to receive input, and one which the driver places buffers to transmit output.

B.2 What Device Configuration Space Layout?

Device configuration space should only be used for initialization-time parameters. It is a limited resource with no synchronization between field written by the driver, so for most uses it is better to use a virtqueue to update configuration information (the network device does this for filtering, otherwise the table in the config space could potentially be very large).

Remember that configuration fields over 32 bits wide might not be atomically writable by the driver. Therefore, no writeable field which triggers an action ought to be wider than 32 bits.

B.3 What Device Number?

Device numbers can be reserved by the OASIS committee: email virtio-dev@lists.oasis-open.org to secure a unique one.

Meanwhile for experimental drivers, use 65535 and work backwards.

B.4 How many MSI-X vectors? (for PCI)

Using the optional MSI-X capability devices can speed up interrupt processing by removing the need to read ISR Status register by guest driver (which might be an expensive operation), reducing interrupt sharing between devices and queues within the device, and handling interrupts from multiple CPUs. However, some systems impose a limit (which might be as low as 256) on the total number of MSI-X vectors that can be allocated to all devices. Devices and/or drivers should take this into account, limiting the number of vectors used unless the device is expected to cause a high volume of interrupts. Devices can control the number of vectors used by limiting the MSI-X Table Size or not presenting MSI-X capability in PCI configuration space. Drivers can control this by mapping events to as small number of vectors as possible, or disabling MSI-X capability altogether.

B.5 Device Improvements

Any change to device configuration space, or new virtqueues, or behavioural changes, should be indicated by negotiation of a new feature bit. This establishes clarity¹ and avoids future expansion problems.

Clusters of functionality which are always implemented together can use a single bit, but if one feature makes sense without the others they should not be gratuitously grouped together to conserve feature bits.

¹Even if it does mean documenting design or implementation mistakes!

Appendix C. Acknowledgements

The following individuals have participated in the creation of this specification and are gratefully acknowledged:

Participants:

Amit Shah, Red Hat
Amos Kong, Red Hat
Anthony Liguori, IBM
Bruce Rogers, Novell
Bryan Venteicher, NetApp
Cornelia Huck, Red Hat
Daniel Kiper, Oracle
Geoff Brown, Machine-to-Machine Intelligence (M2MI) Corporation
Gershon Janssen, Individual Member
James Bottomley, Parallels IP Holdings GmbH
Luiz Capitulino, Red Hat
Michael S. Tsirkin, Red Hat
Paolo Bonzini, Red Hat
Pawel Moll, ARM
Richard Sohn, Alcatel-Lucent
Rusty Russell, IBM
Sasha Levin, Oracle
Sergey Tverdyshev, Thales e-Security
Stefan Hajnoczi, Red Hat
Tom Lyon, Samya Systems, Inc.

The following non-members have provided valuable feedback on this specification and are gratefully acknowledged:

Reviewers:

Andrew Thornton, Google
Arun Subbarao, LynuxWorks
Brian Foley, ARM
David Alan Gilbert, Red Hat
Fam Zheng, Red Hat
Gerd Hoffmann, Red Hat
Jason Wang, Red Hat
Laura Novich, Red Hat
Patrick Durusau, Technical Advisory Board, OASIS
Thomas Huth, Red Hat
Yan Vugenfirer, Red Hat / Daynix
Kevin Lo, MSI

Appendix D. Revision History

The following changes have been made since the previous version of this specification:

Revision	Date	Editor	Changes Made
540	11 Oct 2015	Greg Kurz	virtqueues: fix trivial typo See 2.5.7 .
541	11 Oct 2015	Paolo Bonzini	virtio-blk: fix typo in legacy framing requirements section See 5.2.6.4 .

Revision	Date	Editor	Changes Made
545	18 Oct 2015	Paolo Bonzini	<p>virtio-blk: restore VIRTIO_BLK_F_FLUSH and VIRTIO_BLK_F_CONFIG_WCE</p> <p>VIRTIO_BLK_F_CONFIG_WCE is important in order to achieve good performance (up to 2x, though more realistically +30-40%) in latency-bound workloads. However, it was removed by mistake together with VIRTIO_BLK_F_FLUSH.</p> <p>In addition, even removing VIRTIO_BLK_F_FLUSH was probably not a great idea, because it simplifies simple drivers (e.g. firmware) that are okay with a writethrough cache but still need data to persist after power loss. What really should have been removed is just the possibility that devices not propose VIRTIO_BLK_F_FLUSH, but even that only deserves a "SHOULD" in the new world of conformance statements.</p> <p>Restore these, with the following changes:</p> <ul style="list-style-type: none"> * clarify and use conformance statements in order to define writeback and writethrough caching according to what is commonly done by high-end storage. * clarify (with conformance statements) the influence of the VIRTIO_BLK_F_FLUSH feature on caching and how to proceed if only one of VIRTIO_BLK_F_FLUSH and VIRTIO_BLK_F_CONFIG_WCE is negotiated. * strengthen the requirement for persisting writes to MUST after a VIRTIO_BLK_T_FLUSH request (and in other cases too involving the new features). <p>The suggested behavior upon feature negotiation is okay for the Linux implementation of virtio1, even after the implementation is modified to support the two new features.</p> <p>This fixes VIRTIO-144. See 5.2, 7.2.5 and 7.3.5.</p>
546	18 Oct 2015	Michael S. Tsirkin	<p>pci: clarify configuration access capability rules</p> <p>The point of the configuration access capability is to enable access to other capabilities. The intent never was to allow writes to a random place within device BARs. Limiting drivers simplifies devices - and devices can always add another capability if drivers ever want to access some other range.</p> <p>This resolves VIRTIO-145. See 4.1.4.7.2.</p>

Revision	Date	Editor	Changes Made
547	18 Oct 2015	Michael S. Tsirkin	add advice on transition from legacy interfaces Reading legacy chapters gives a hint about what changed, let's help readers discover this useful shortcut. This resolves VIRTIO-146. See 1.3.2 .
554	16 Feb 2016	Thomas Huth	virtio-net: fix inconsistent legacy header size Current text says: The legacy driver only presented num_buffers in the struct virtio_net_hdr when VIRTIO_NET_F_MRG_RXBUF was not negotiated; Should be: "...was negotiated ..." instead of "...was not negotiated ..." To be consistent with the following: without that feature the structure was 2 bytes shorter. See 5.1.6.1 .
555	16 Feb 2016	Michael S. Tsirkin	virtio header: tweak change motivation The changes are not just to remove Linux assumptions, we have also renamed ring->queue. Tweak the header description accordingly. See 2.5.10 .
558	16 Feb 2016	Michael S. Tsirkin	rename virtio_ring.h to virtio_queue.h Since vring* and VRING* have been replaced with virtq* and VIRTQ* respectively, rename the header virtio_ring.h to virtio_queue.h. See A .
559	16 Feb 2016	Michael S. Tsirkin	init: sort status bits Status bit order is inconsistent: they are neither in increasing order nor in the order they are likely to be used. The second approach seems more useful since there aren't that many bits, so the numerical order does not help much. A typical order of use would be: <ul style="list-style-type: none"> • ACKNOWLEDGE • DRIVER • then either FAILED or FEATURES_OK • then either FAILED or DRIVER_OK • then DEVICE_NEEDS_RESET (if device detects an error) Sort the bits accordingly. See 2.1 .